

Hadoop Basics

Nagabhushan



Agenda – Day 1

- Introduction to Big Data and Hadoop
- History of Hadoop & Use Cases
- HDFS Storage Architecture
- Hadoop Setup

What is Big Data?

- *3 Vs of Big Data*

- *Velocity → Speed*
- *Variety → Different forms of data*
- *Volume → Size of data*

- *Hadoop's 4th V → VALUE*

- *How to store Big Data? → HDFS –*

- *How to process Big Data? → MapReduce (1.x) / YARN (2.x)*




Hadoop Core

Data Measurement Scale

▣ 1 KB	Kilobyte	1000
▣ 1 MB	Megabyte	1000000
▣ 1 GB	Gigabyte	1000000000
▣ 1 TB	Terabyte	1000000000000
▣ 1 PB	Petabyte	1000000000000000
▣ 1 EB	Exabyte	1000000000000000000
▣ <u>1 ZB</u>	<u>Zetabyte</u>	<u>1000000000000000000000000 X 5</u>
▣ 1 YB	Yotabyte	1000000000000000000000000000000

Traditional System


Moving Data to Processing



Big Data!

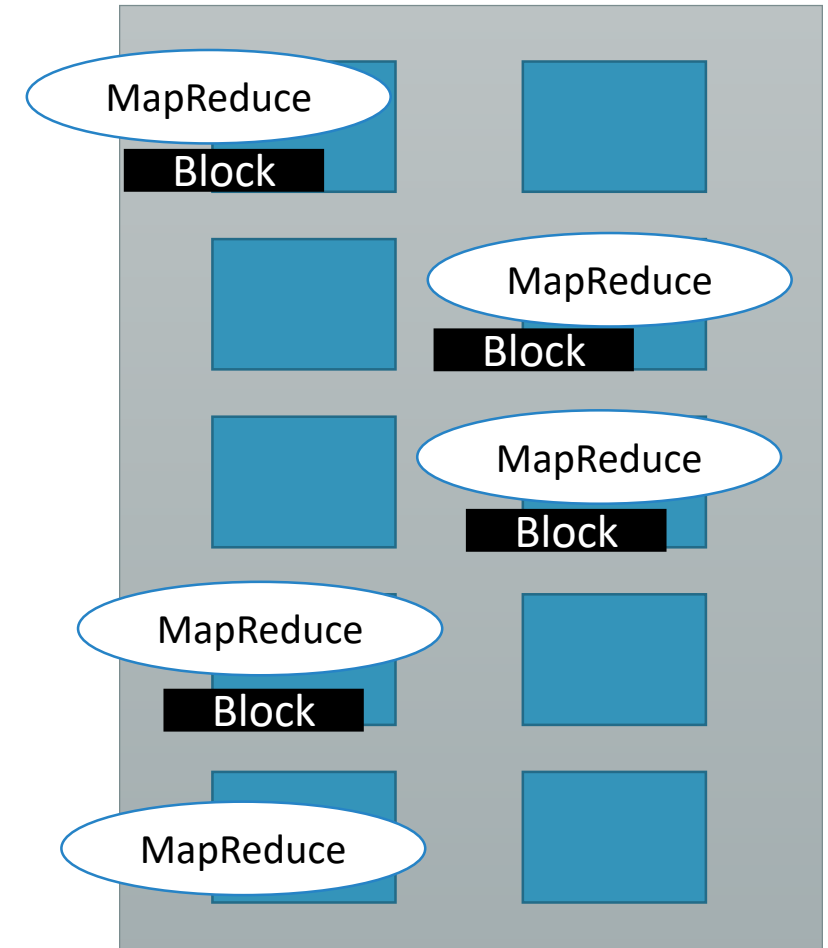


8 cores
64 GB RAM
240 TB Storage



Hadoop System

Moving Processing to Data



Features of Hadoop

- Commodity Hardware
- Open Source
- Fault Tolerance
- Distributed Storage
- Read Only File System → Parallel Processing (free)
- Horizontal Scaling
- Data Locality → Move Processing to Data
- Simplified Programming


Limitations of Hadoop

- WORM → Write Once Read Many
- Sequential Access to Data – No Random Reads

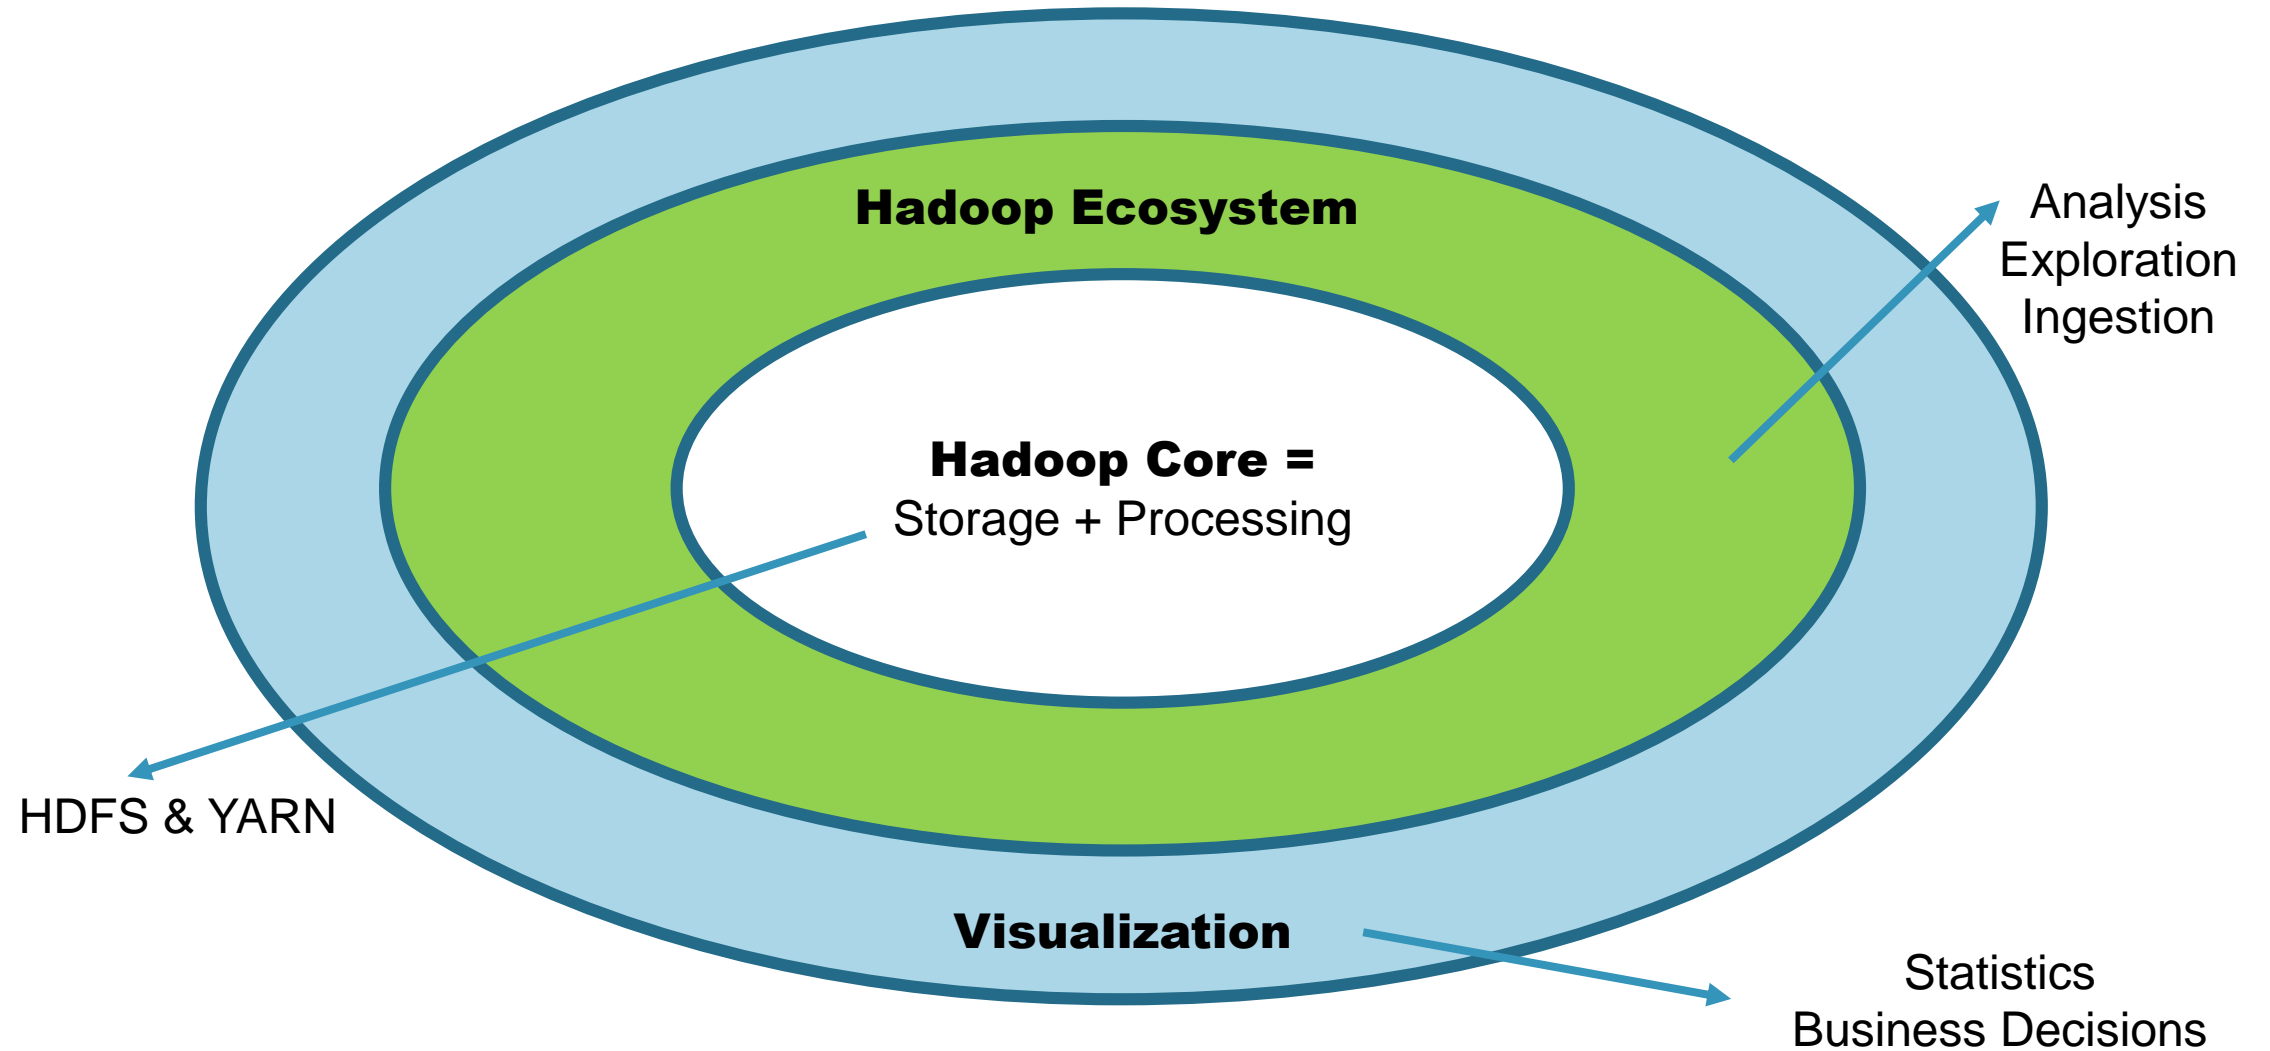
A simple use case!

Business Owner

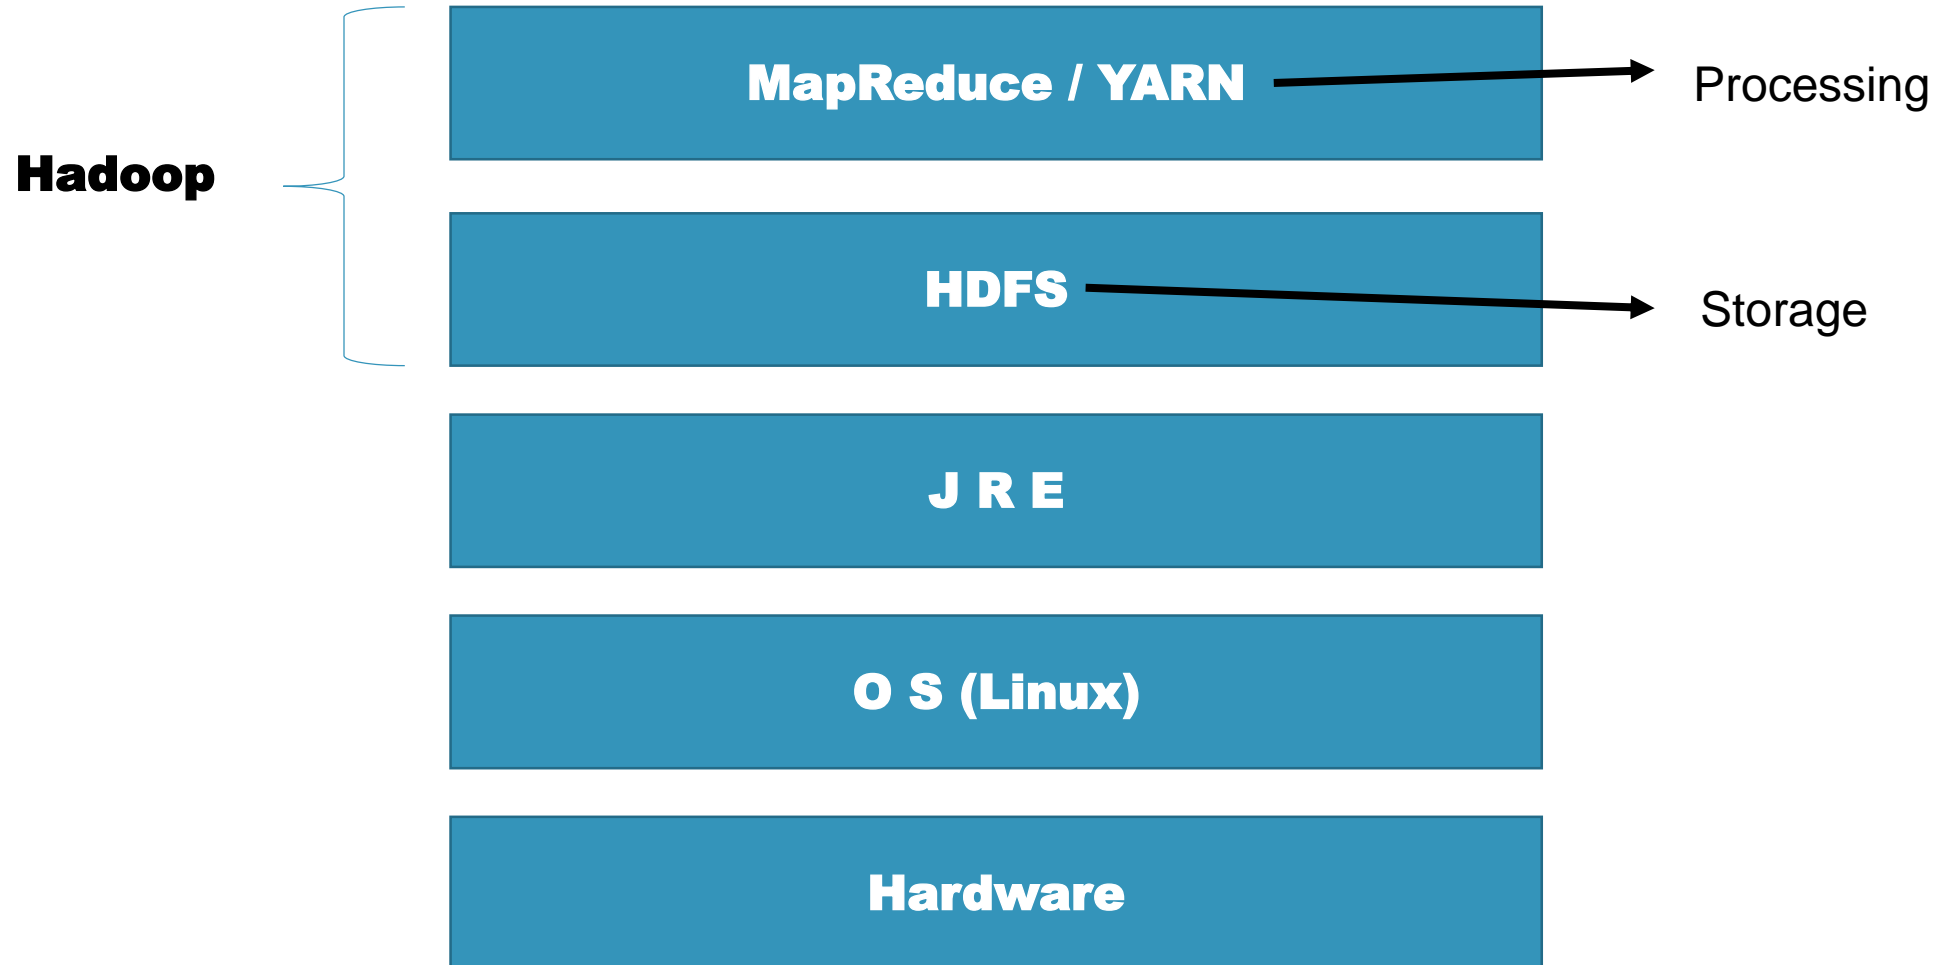
What products are **NOT** selling? **Why?**



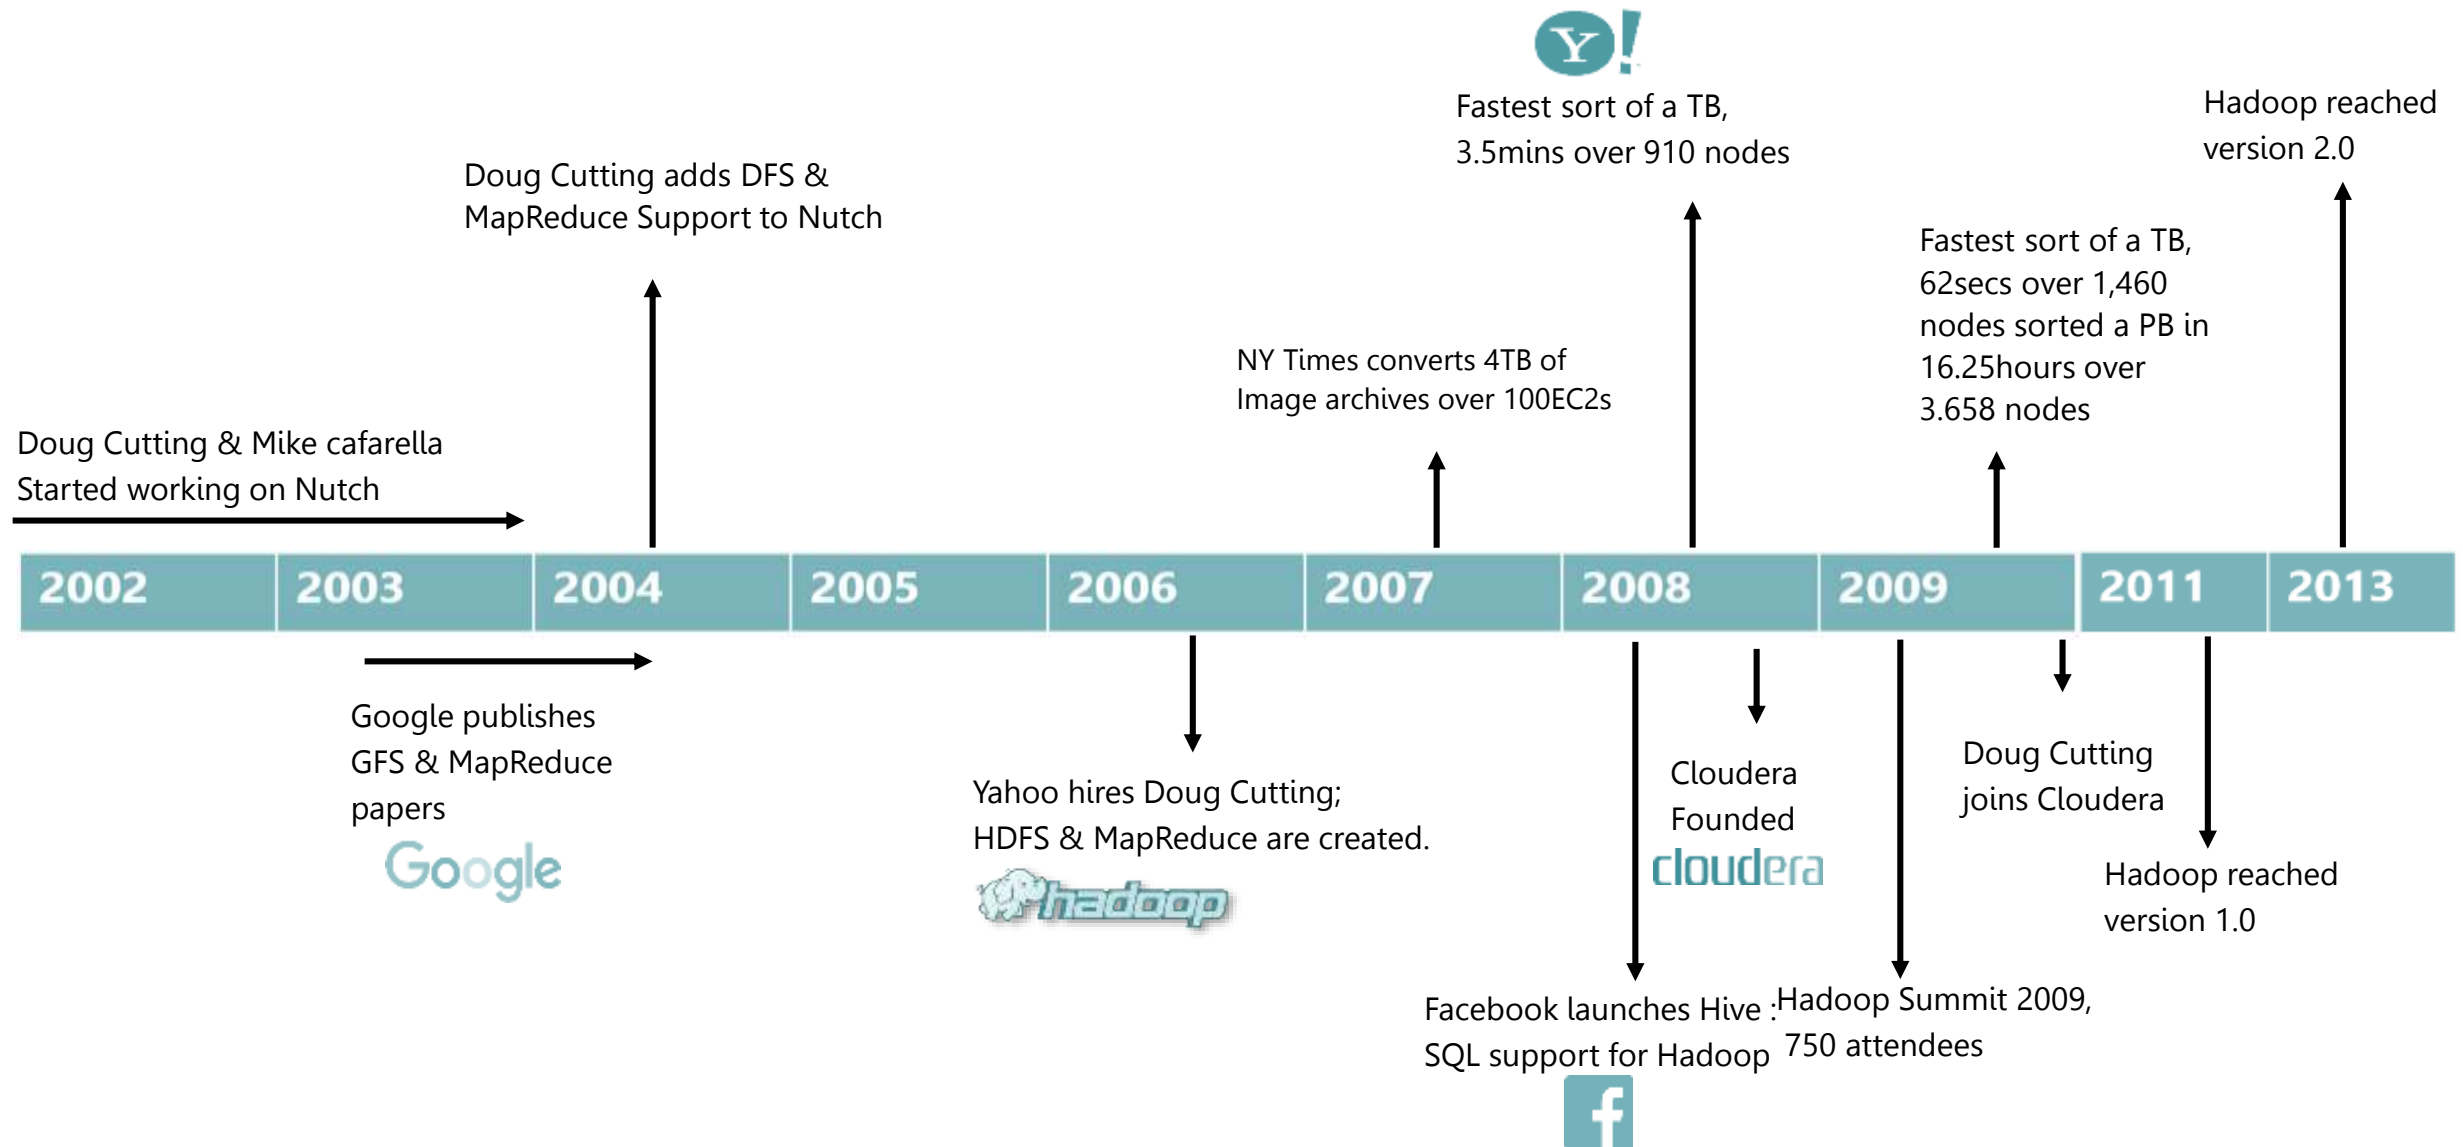
Hadoop Overview – Core & Ecosystem



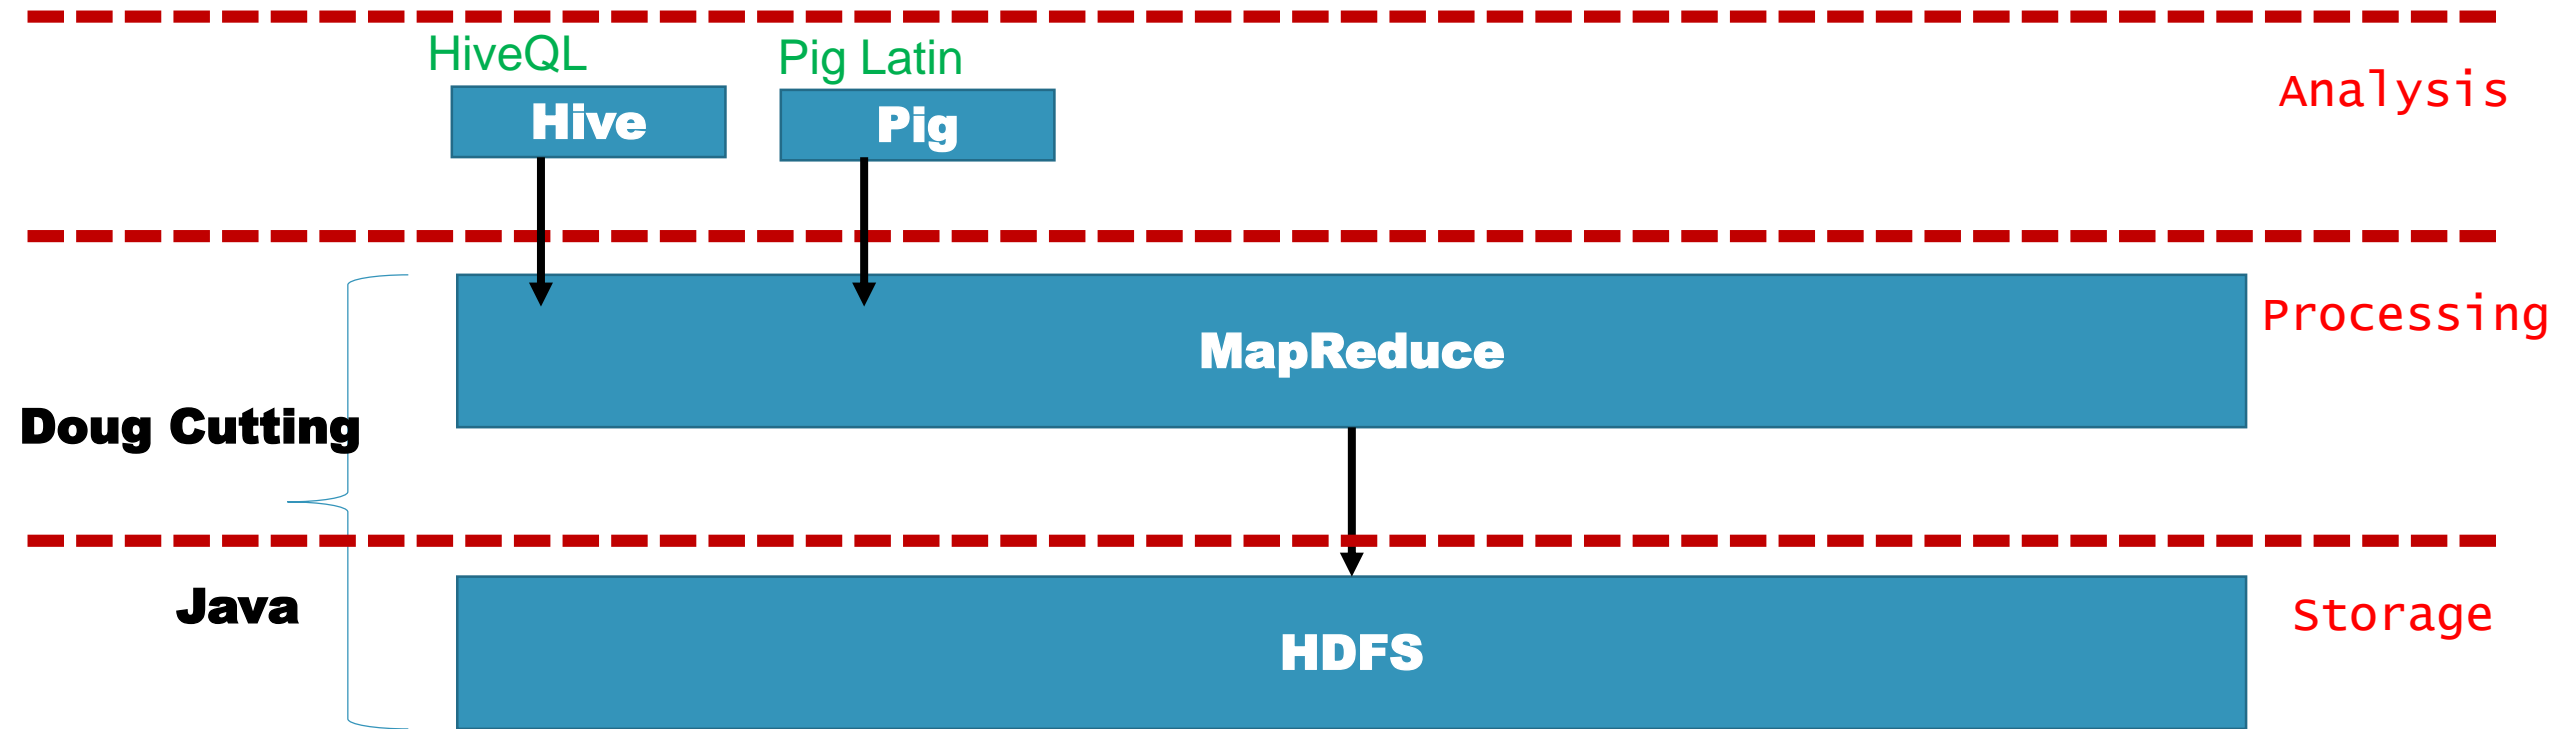
Hadoop's Physical Architecture



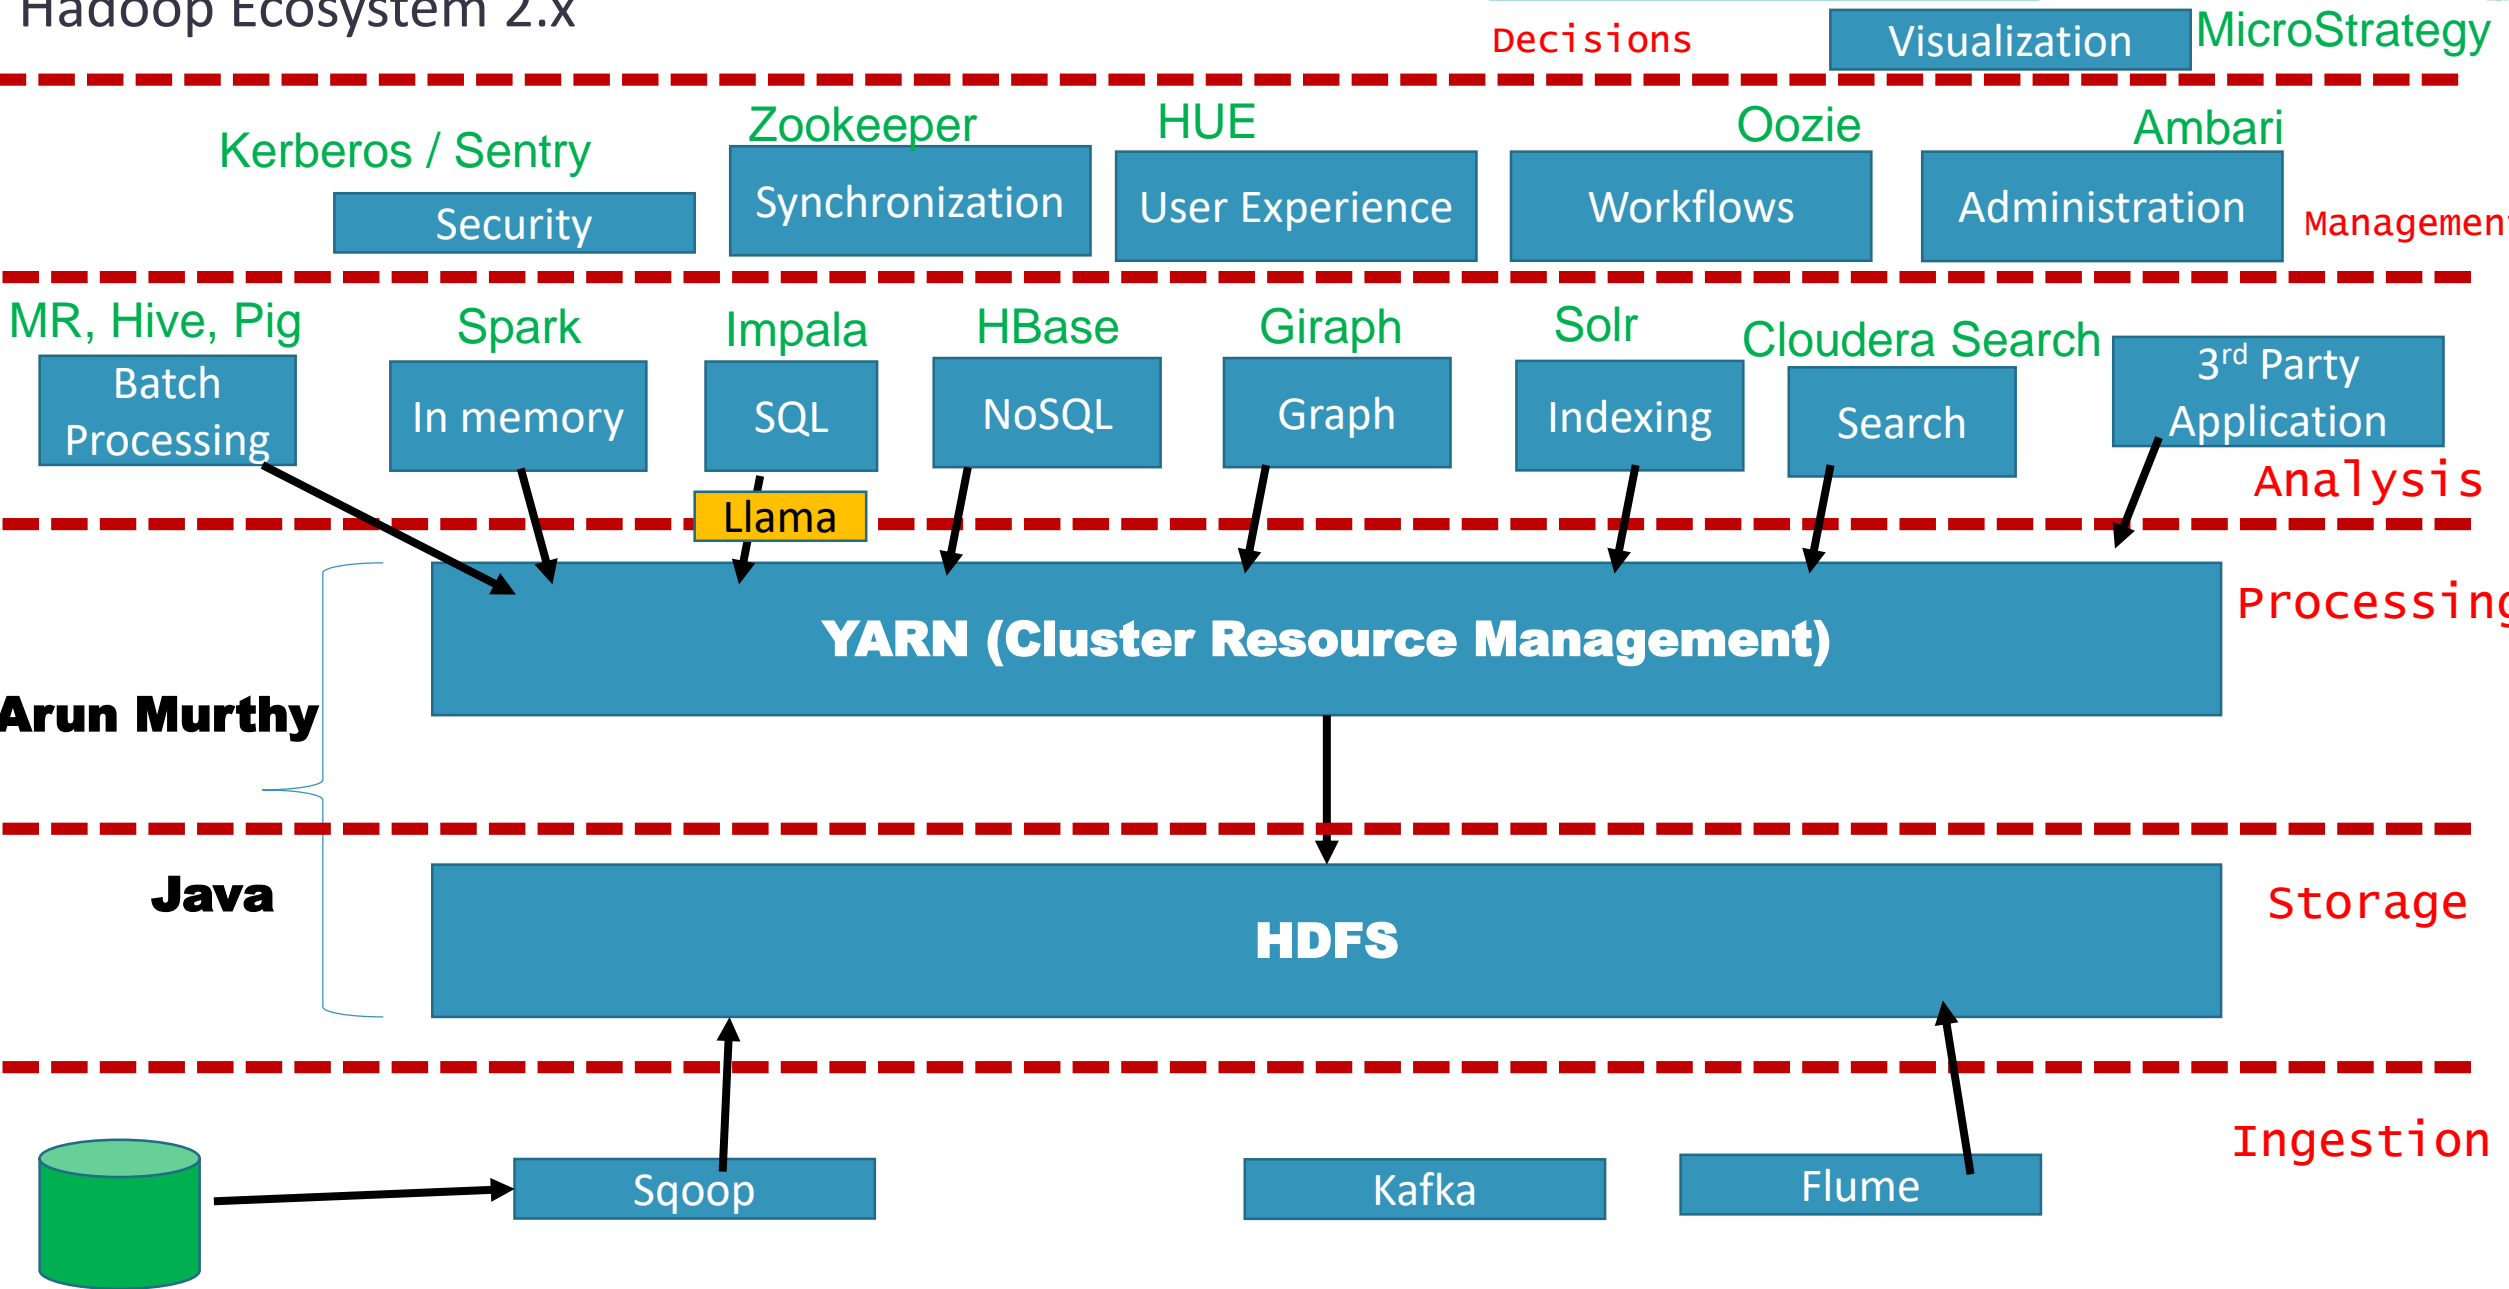
History of Hadoop



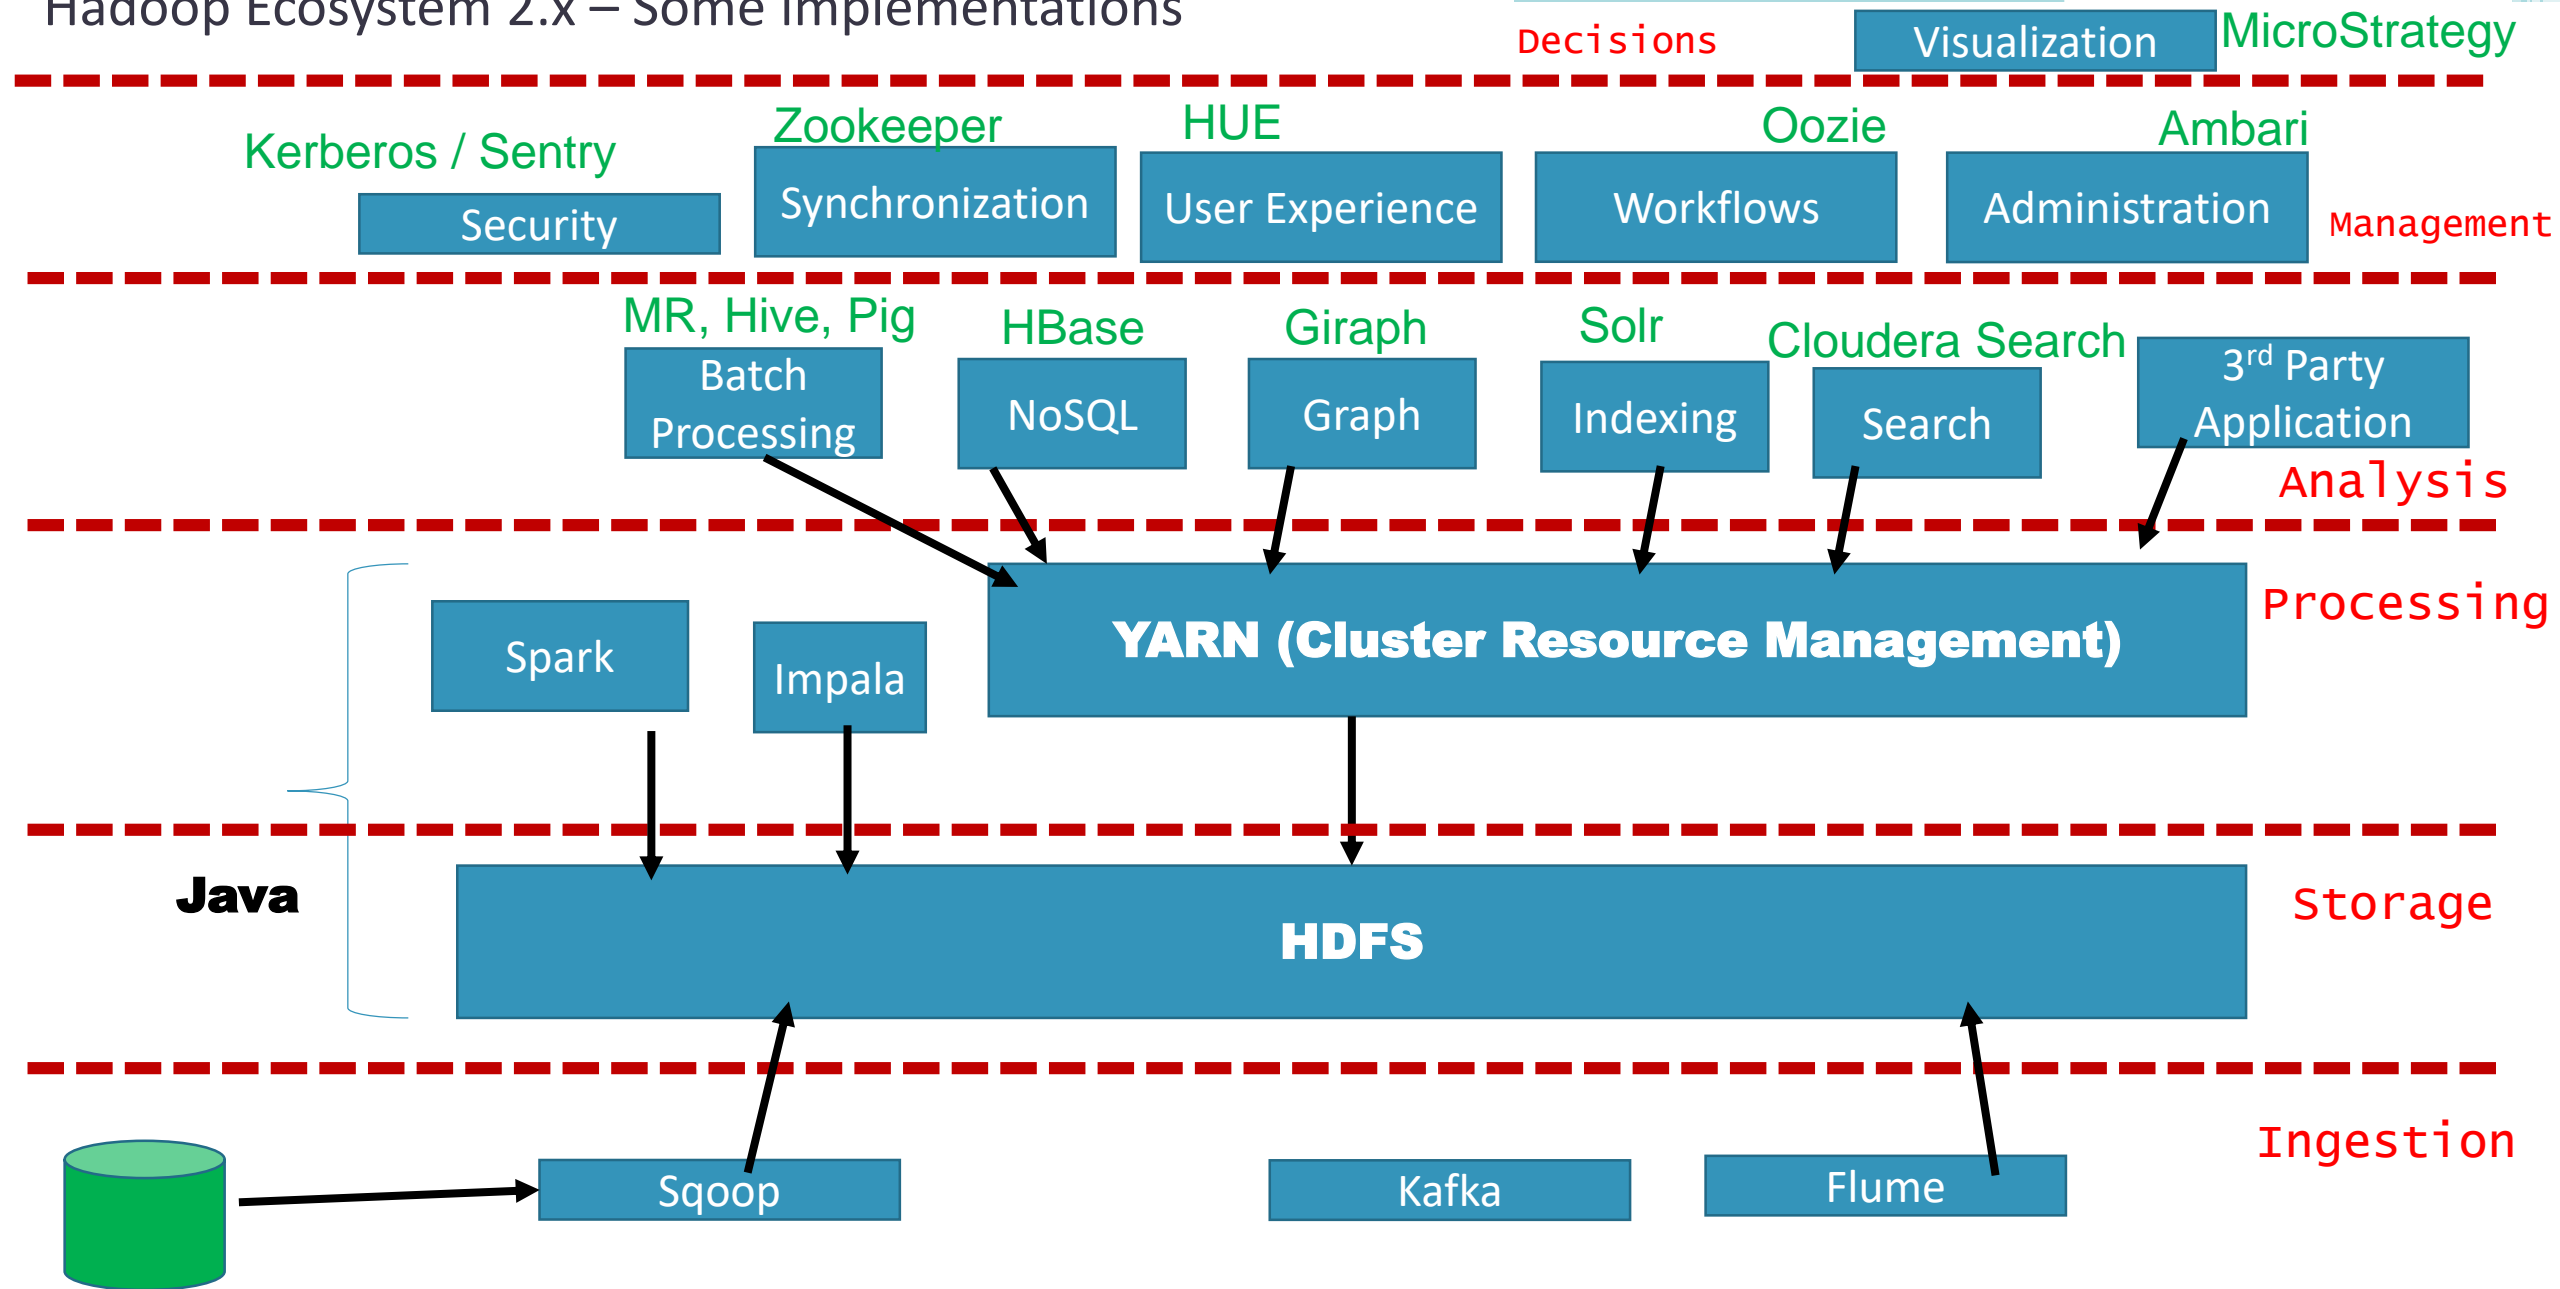
Hadoop Ecosystem 1.x



Hadoop Ecosystem 2.x



Hadoop Ecosystem 2.x – Some implementations



Commercial Distributions

- Cloudera <http://www.cloudera.com/>
- Hortonworks <http://hortonworks.com/>
- MAPR <https://www.mapr.com/>

HDFS Daemons

Master – Slave Architecture

NameNode – Metadata is in 2 file


- fsimage
- edit logs (transaction logs)

NameNode

DataNode

Secondary NameNode

NameNode – orchestrates storage
Manages the File System Metadata



DataNode – Blocks are stored
Listen to NameNode's instructions
Send heartbeats / block report

Checkpoint Node

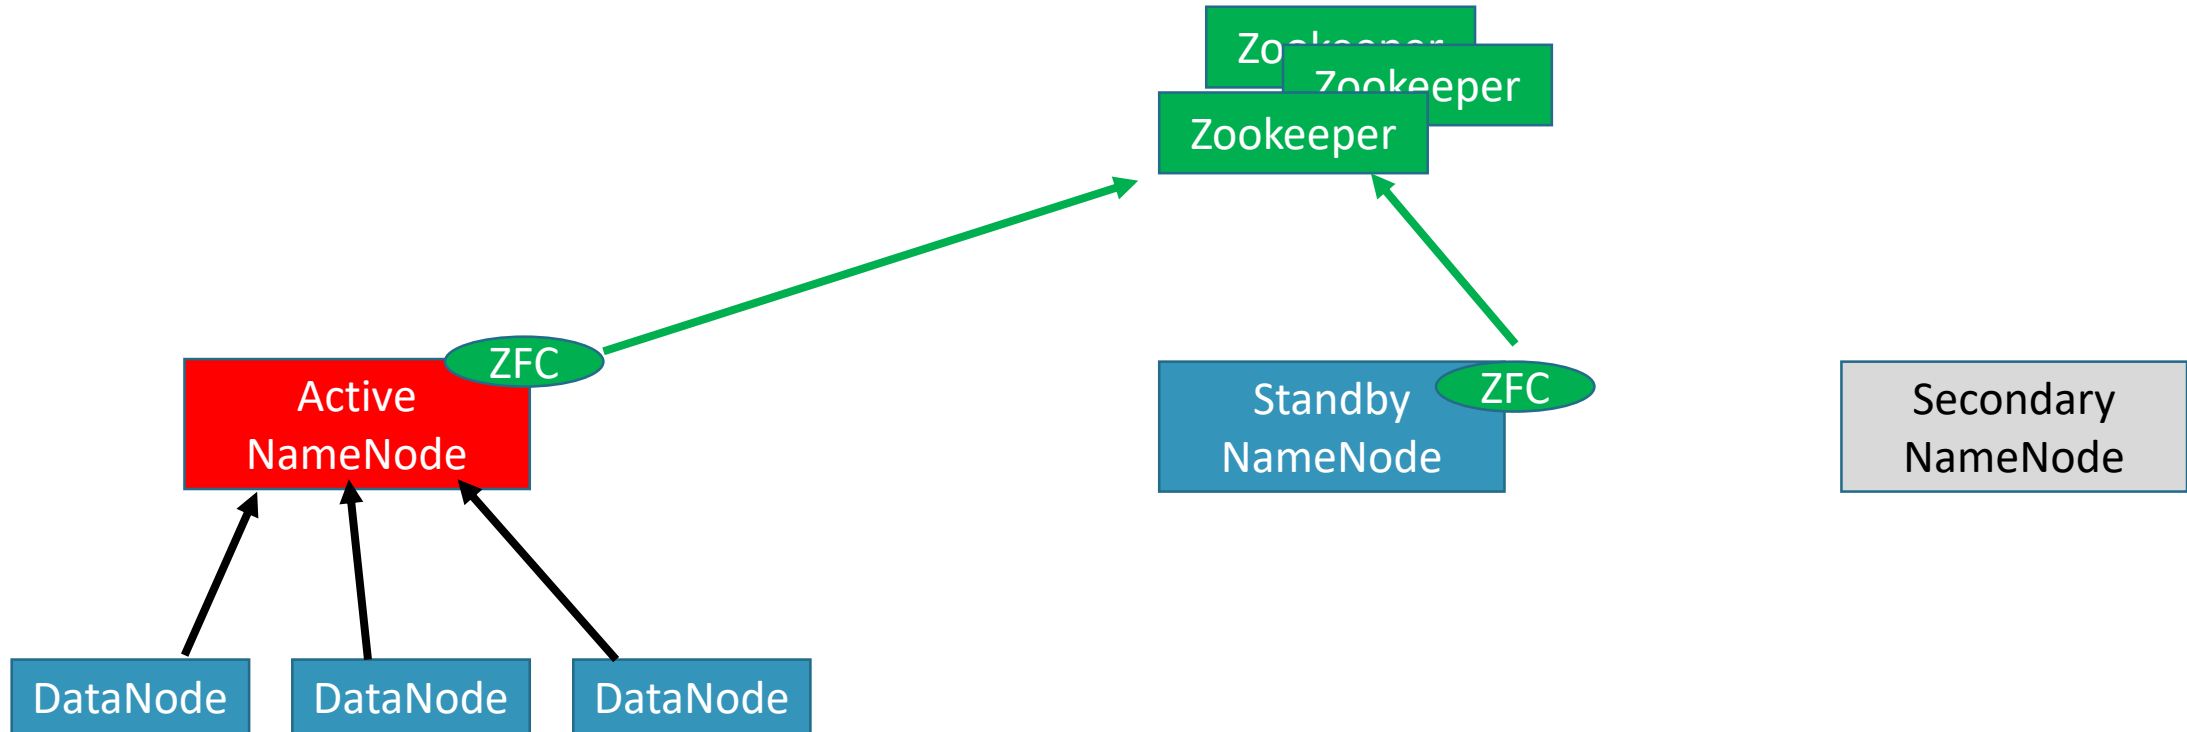
– Regularly backup the metadata

Secondary
NameNode

dfs.namenode.checkpoint.period=3600
dfs.namenode.checkpoint.txns=1000000

Hadoop 1.x → NameNode is Single Point of Failure

HDFS 2 – Introduces NameNode High Availability




YARN Daemons

Master – Slave Architecture

ResourceManager
NodeManager
JobHistoryServer

ResourceManager – Scheduling Jobs




NodeManager – Execution of jobs



Retired Jobs


MRAppMaster → Monitoring jobs

Hadoop Terminologies

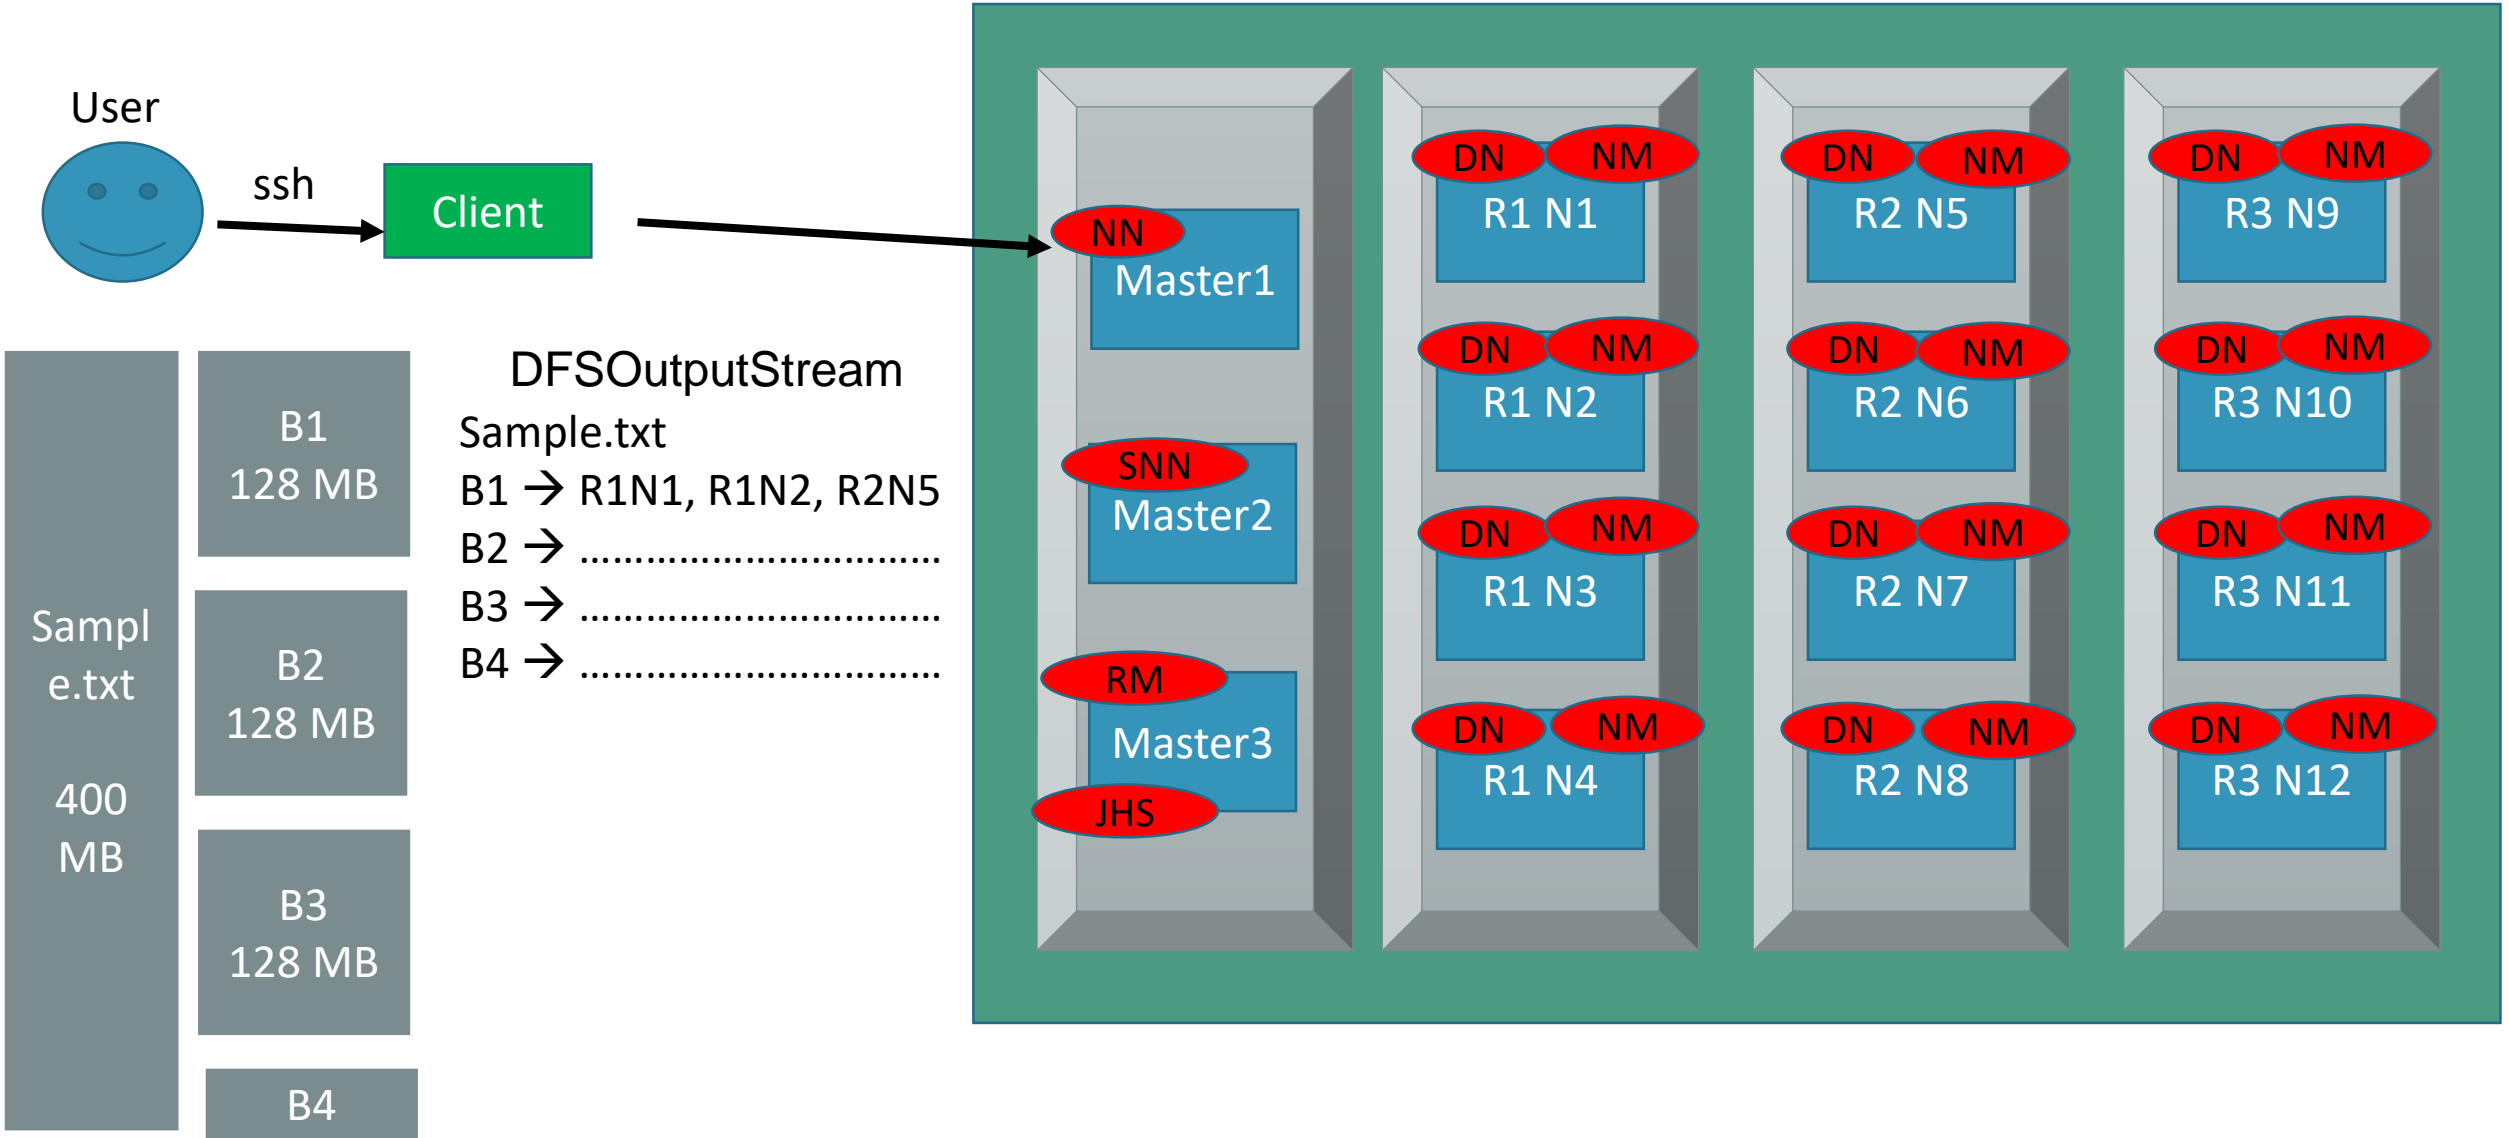


HDFS Daemons Distributed Over a cluster

Client



How to write a File to the cluster



Default Hadoop Configuration

core-default.xml
hdfs-default.xml
mapred-default.xml
yarn-default.xml



dfs.replication=3
dfs.blocksize= 134217728 = 128 MB
dfs.heartbeat.interval=3
dfs.namenode.stale.datanode.interval=30000

Customized Hadoop Configuration

core-site.xml
hdfs-site.xml
mapred-site.xml
yarn-site.xml



dfs.replication=1

`$HADOOP_HOME/etc/hadoop` ➔ Hadoop's conf dir



Hadoop Setup

Hardware

In premise

Virtualization

Cloud – AWS / GCP

Hadoop Flavor

Cloudera

Apache

Hortonworks

MAPR

O S Flavor

RHEL

SUSE

Fedora

Ubuntu

CentOS

.....

Mode of Deployment

Standalone Mode

Pseudo Distributed Mode

Fully Distributed Mode

- Standalone Operation By default, Hadoop is configured to run in a non-distributed mode, as a single Java process. This is useful for debugging
- Pseudo-Distributed Operation: Hadoop can also be run on a single-node in a pseudo-distributed mode where each Hadoop daemon runs in a separate Java process
- Fully Distributed Mode – Multi Node cluster, a typical production environment where each daemons would be distributed on many nodes

Hadoop Setup Steps

- Pre-Requisites
 - JDK
 - ssh (Passphraseless)
- Download and unpack Hadoop packages
- Customize Hadoop
 - core-site.xml
 - hdfs-site.xml
 - mapred-site.xml
 - yarn-site.xml
 - hadoop-env.sh
- Format the NameNode / DataNode
- Start Hadoop Services

Web UI Ports

NameNode – <http://localhost:50070>

ResourceManager – <http://localhost:8088>

HistoryServer – <http://localhost:19888>

Secondary NameNode - <http://localhost:50090>

DataNode - <http://localhost:50075>

Agenda – Day 2

- Multi Node Setup
- Working with HDFS & File System commands
- MapReduce & YARN
- Data Ingestion → Sqoop

Hadoop Setup – Multi Node Demo

Hardware

In premise
Virtualization
Cloud – AWS / GCP

Hadoop Flavor

Cloudera
Apache
Hortonworks
MAPR

O S Flavor

RHEL
SUSE
Fedora
Ubuntu
CentOS

.....

Mode of Deployment

Standalone Mode
Pseudo Distributed Mode
Fully Distributed Mode

Hadoop Setup Steps (AWS)

- Login to AWS with credentials and get the required keys (*.pem) – AWS specific ssh key
 - As we would work on Windows, to connect to the remote machine (AWS), we need a ssh client (Ex: PuTTY)
 - While connecting via PuTTY, the *.pem is not recognized, we need PuTTYgen to convert *.pem to *.ppk
 - Refer <https://docs.aws.amazon.com/console/ec2/instances/connect/putty>
-
1. Subscribe for 3 EC2 instances on <https://aws.amazon.com/> of AMI “Ubuntu Server 14.04 LTS”
 2. Connect to one of the instance and download the Cloudera installer using the command:
wget <http://archive.cloudera.com/cm5/installer/latest/cloudera-manager-installer.bin>

Thank you for choosing Cloudera Manager and CDH.

This installer will install **Cloudera Express 5.9.0** and enable you to later choose packages for the services below (there may be some license implications).

- Apache Hadoop (Common, HDFS, MapReduce, YARN)
- Apache HBase
- Apache ZooKeeper
- Apache Oozie
- Apache Hive
- Hue (Apache licensed)
- Apache Flume
- Cloudera Impala (Apache licensed)
- Apache Sentry
- Apache Sqoop
- Cloudera Search (Apache licensed)
- Apache Spark

You are using Cloudera Manager to install and configure your system. You can learn more about Cloudera Manager by clicking on the **Support** menu above.

Working with HDFS

<http://hadoop.apache.org/docs/r2.7.3/hadoop-project-dist/hadoop-common/FileSystemShell.html>

Read about HDFS Metadata directories from the document within the VM
`/home/user1/Downloads/.05_Programs/15_Misc/HDFS_Metadata_Directories.pdf`

Introduction to MapReduce

- Processing Engine of Hadoop
- Works in 2 phases (Developer)
 - Map
 - Reduce
- The Framework has following phases
 - Input Split
 - **Map**
 - Shuffle & Sort
 - **Reduce**
 - Final Output

MapReduce works on (Key, Value) pairs

Example: (Welcome, 1)

Welcome → Key

1 → Value

Sort → Key



Will be handled by the framework and that makes processing data on a distributed store simpler!

In Map phase – parse, transform, extract

In Reduce phase – statistics, aggregations

MapReduce WordCount Program

Problem Statement: Count the occurrence of each word in the file

/Sample/SampleFile.txt

Welcome to Hadoop
Learning Hadoop is fun
Hadoop Hadoop Hadoop is the buzz

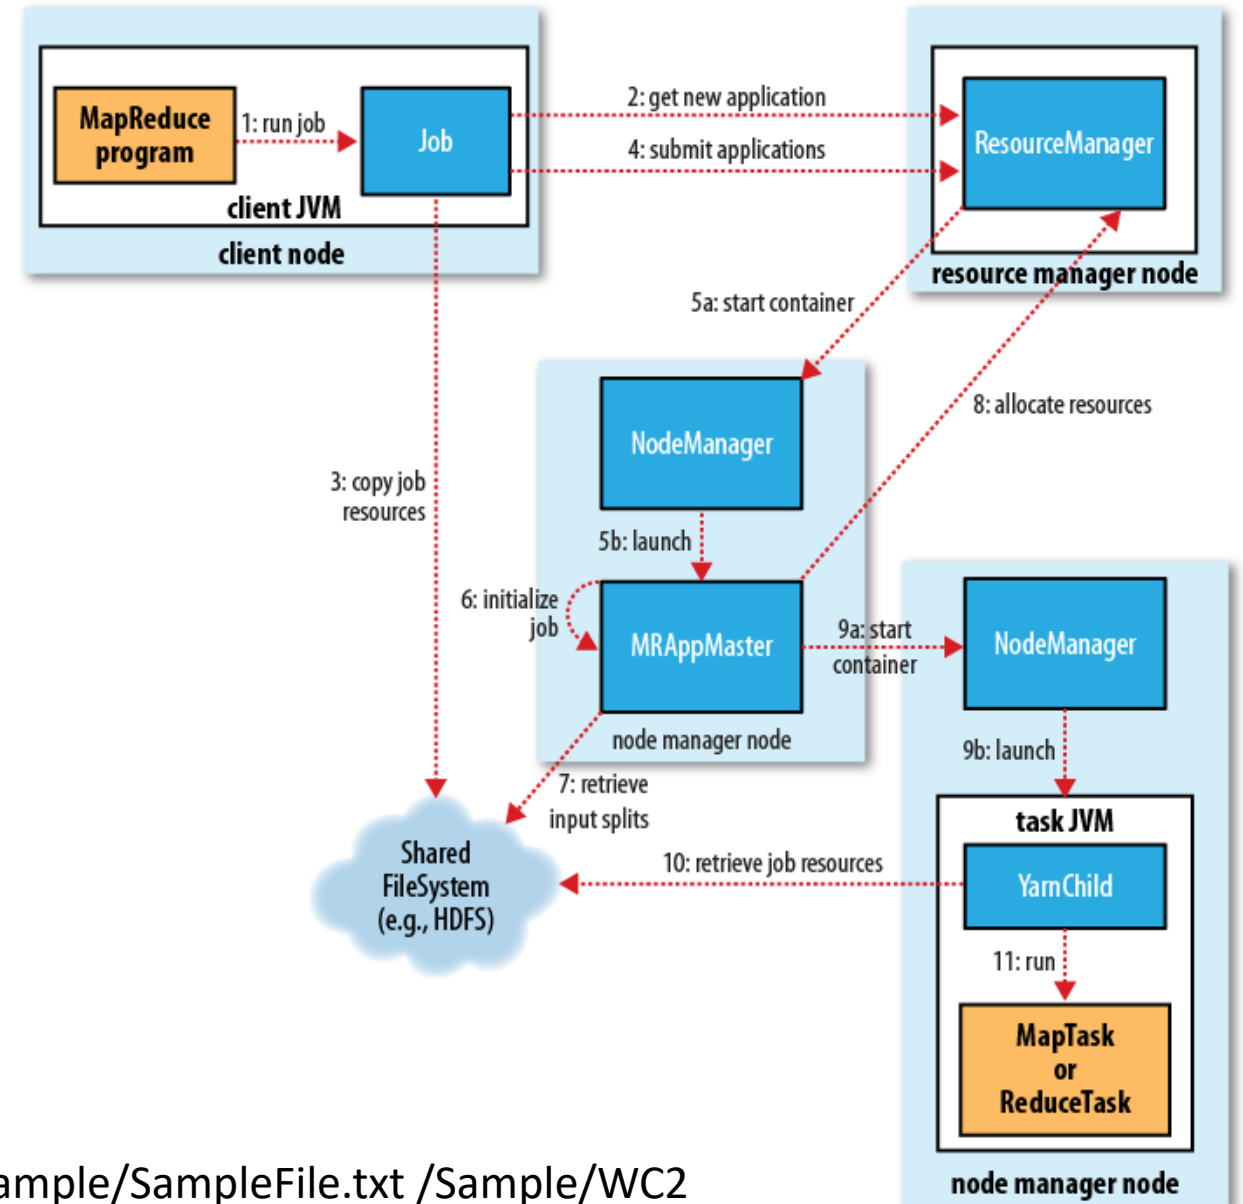
Line to word → Map
Count the words → Reduce

Expected Output

Hadoop, 5
Learning, 1
Welcome, 1
buzz, 1
fun, 1
is, 2
the, 1
to, 1

The Anatomy of a Job Run on YARN


ResourceManager – 1 / cluster
MRAppMaster – 1 / job
YarnChild – 1 / task



yarn jar /home/user1/Documents/wc.jar WordCount /Sample/SampleFile.txt /Sample/WC2


Limitations of Hadoop 1.x

- NameNode is the Single Point of Failure
- JobTracker is the Single Point of Failure
- JobTracker is overburdened → Scheduling + Monitoring
- TaskTracker(s) have slots (2 map slots and 2 reduce slots by default)
 - `mapreduce.tasktracker.map.tasks.maximum=2`
 - `mapreduce.tasktracker.reduce.tasks.maximum=2`
- Inefficient Resource Utilization at the cluster level
- Scalability is at a threshold of 4000+ nodes
- The TaskTracker(s) were only capable of Map/Reduce
- FIFO scheduler on JobTracker



Features of Hadoop 2.x

- NameNode HA → Active NameNode + Standby NameNode
- Resource Manager HA → Active ResourceManager + Standby ResourceManager
- JobTracker is now ResourceManager (1 / cluster) and Application Master (1 / job)
- TaskTracker(s) are NodeManagers and have generic containers capable of running applications beyond Map/Reduce
- ResourceManager + Application Manager + Application Master + History Server = YARN (Yet Another Resource Negotiator) is a new cluster resource management layer which promises efficient cluster resource utilization even at a scale of 25000+ nodes in a cluster
- Capacity Scheduler
- Multi Tenancy



MapReduce Detailed

Welcome to Hadoop
Learning Hadoop is fun
Hadoop Hadoop Hadoop is the buzz

R1 N1

Map

parse, transform, extract

R1 N2

Map

statistics, aggregations

Reduce

R2 N7

Map

R2 N6

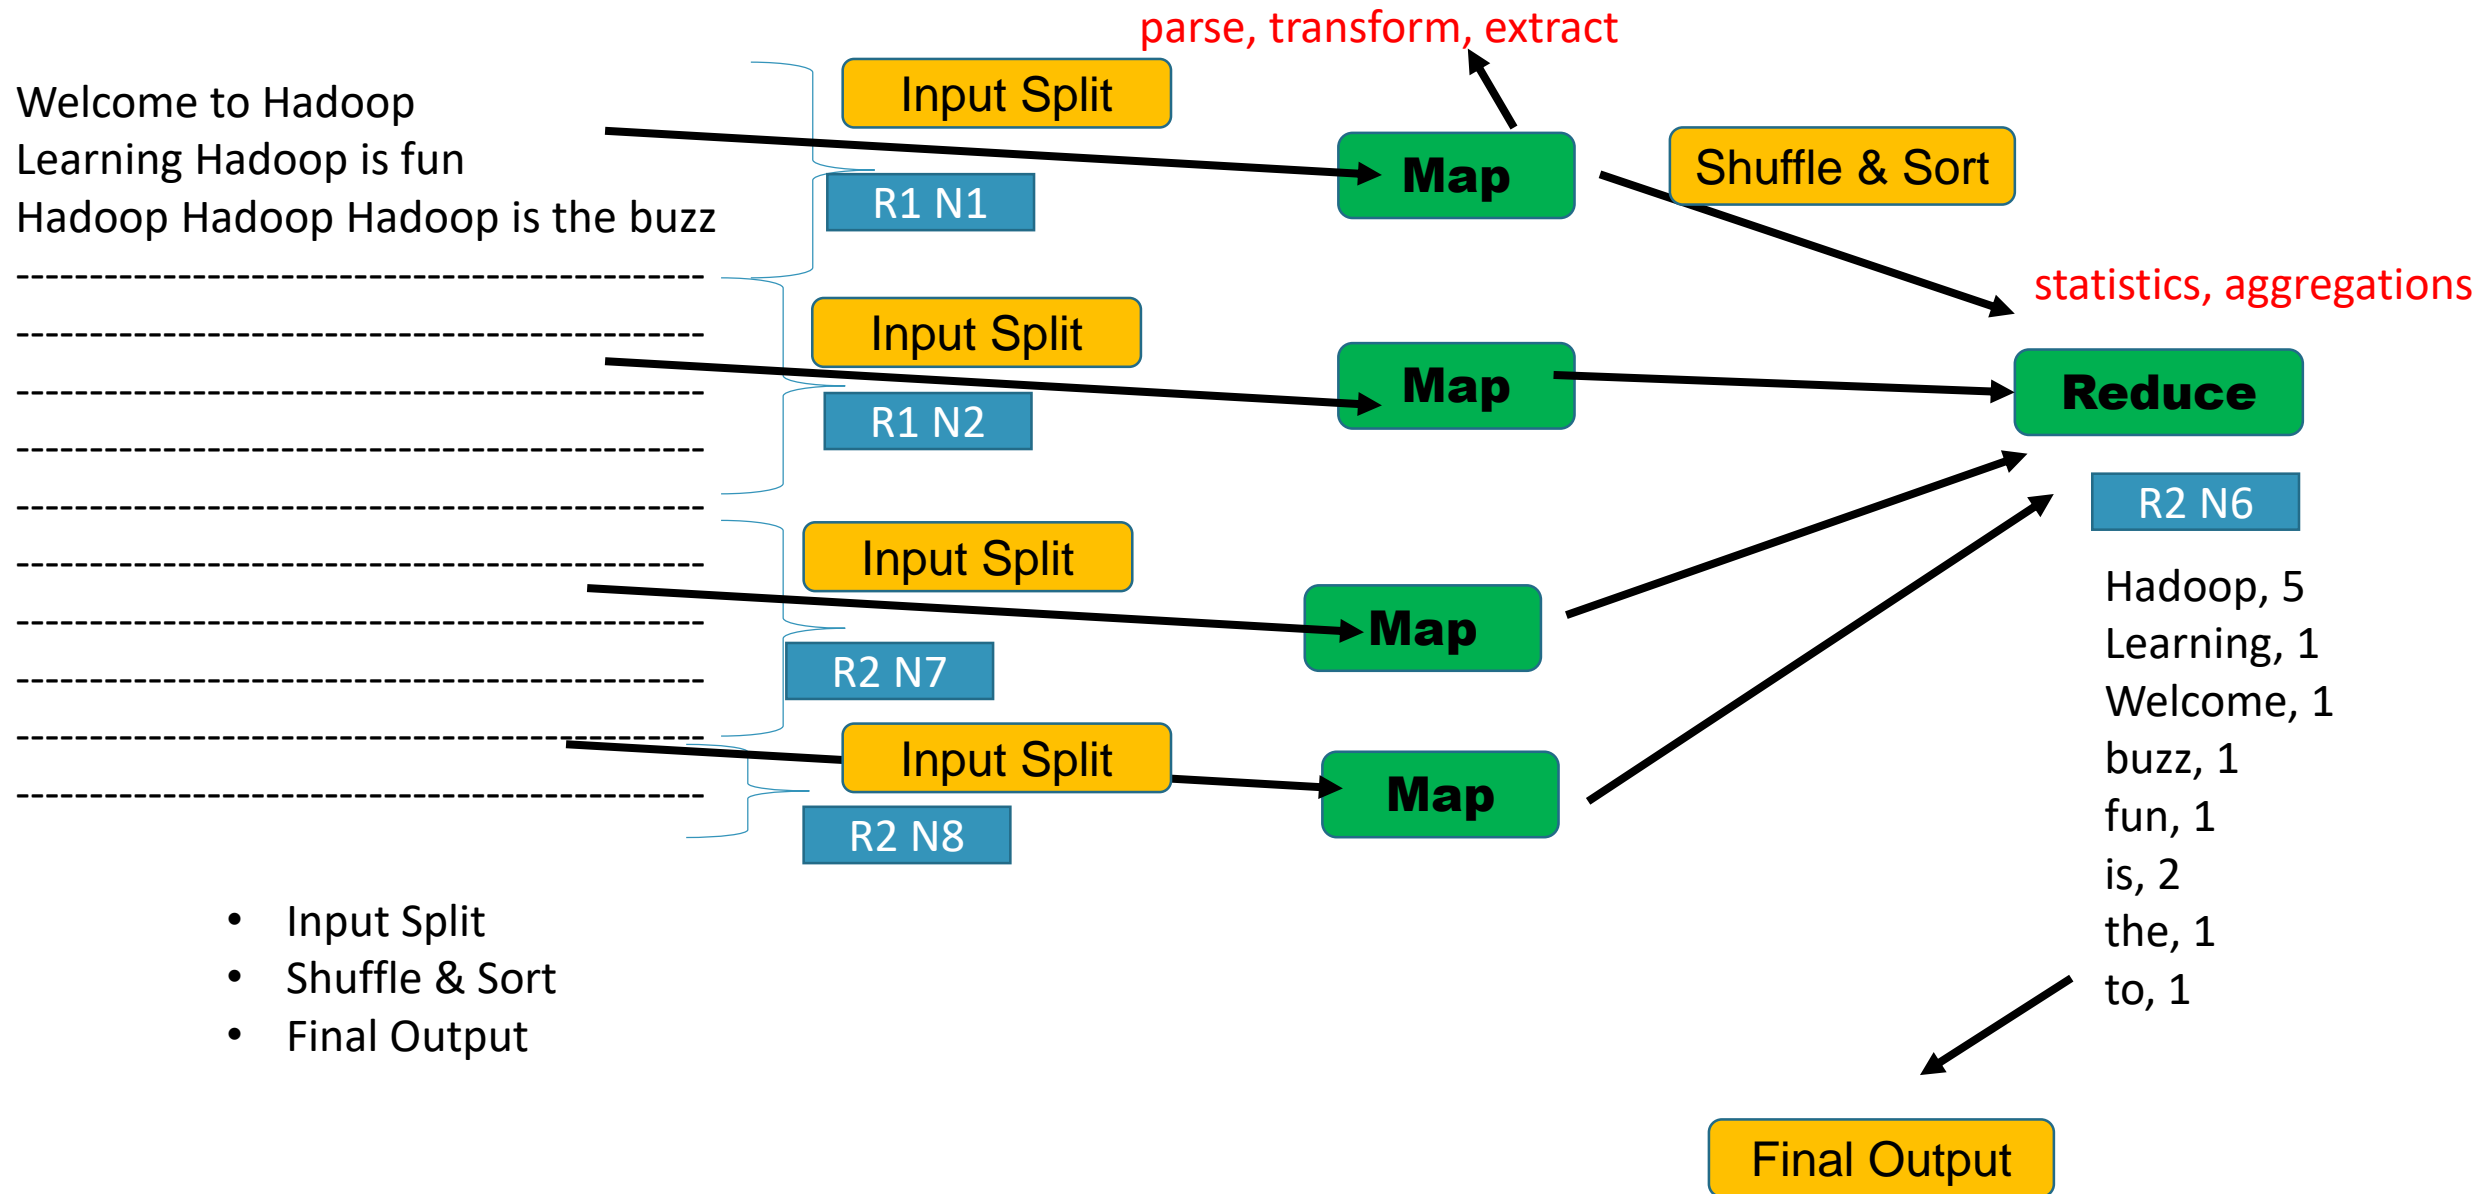
R2 N8

Map

Hadoop, 5
Learning, 1
Welcome, 1
buzz, 1
fun, 1
is, 2
the, 1
to, 1

MapReduce Detailed Steps

Welcome to Hadoop
Learning Hadoop is fun
Hadoop Hadoop Hadoop is the buzz



MapReduce API Overview

</home/user1/HadoopInstallations/hadoop-2.7.1/share/doc/hadoop/api/index.html>

Map Signature

```
Mapper Class {  
    map (key, value, context) {  
        -----  
        -----  
        -----  
        Logic for Transformation  
        -----  
        -----  
        -----  
    }  
}
```

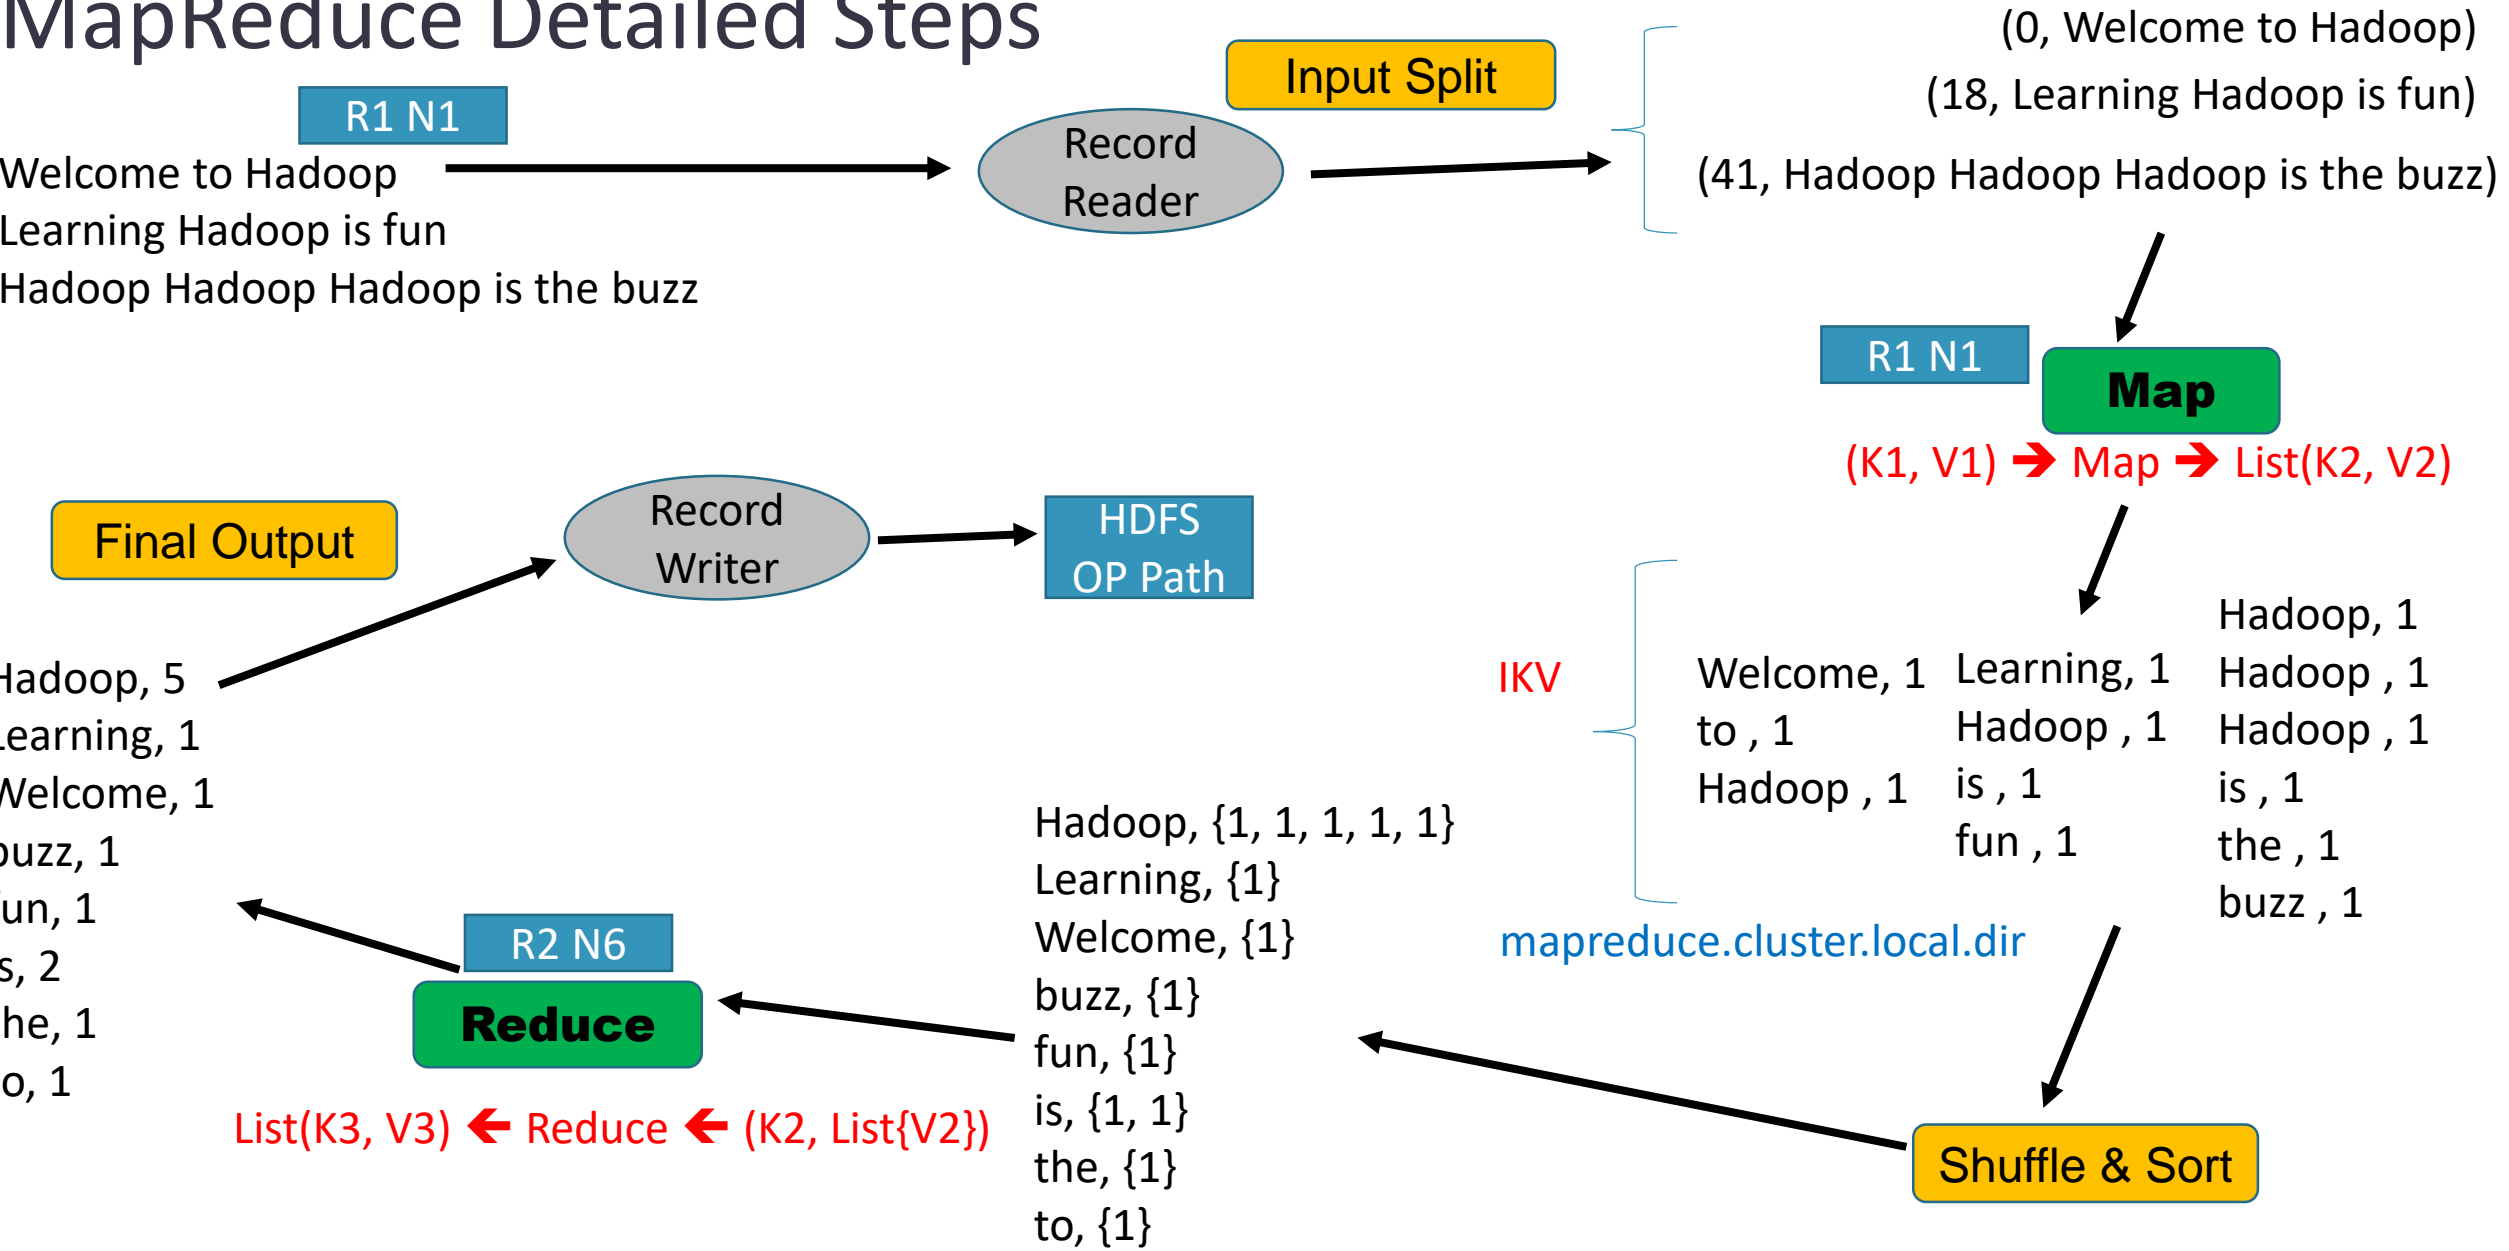
(K1, V1) → Map → List(K2, V2)

Reduce Signature

```
Reducer Class {  
    reduce (key, list{values}, context) {  
        -----  
        -----  
        -----  
        Logic for Aggregation  
        -----  
        -----  
        -----  
    }  
}
```

(K2, List{V2}) → Reduce → List(K3, V3)

MapReduce Detailed Steps



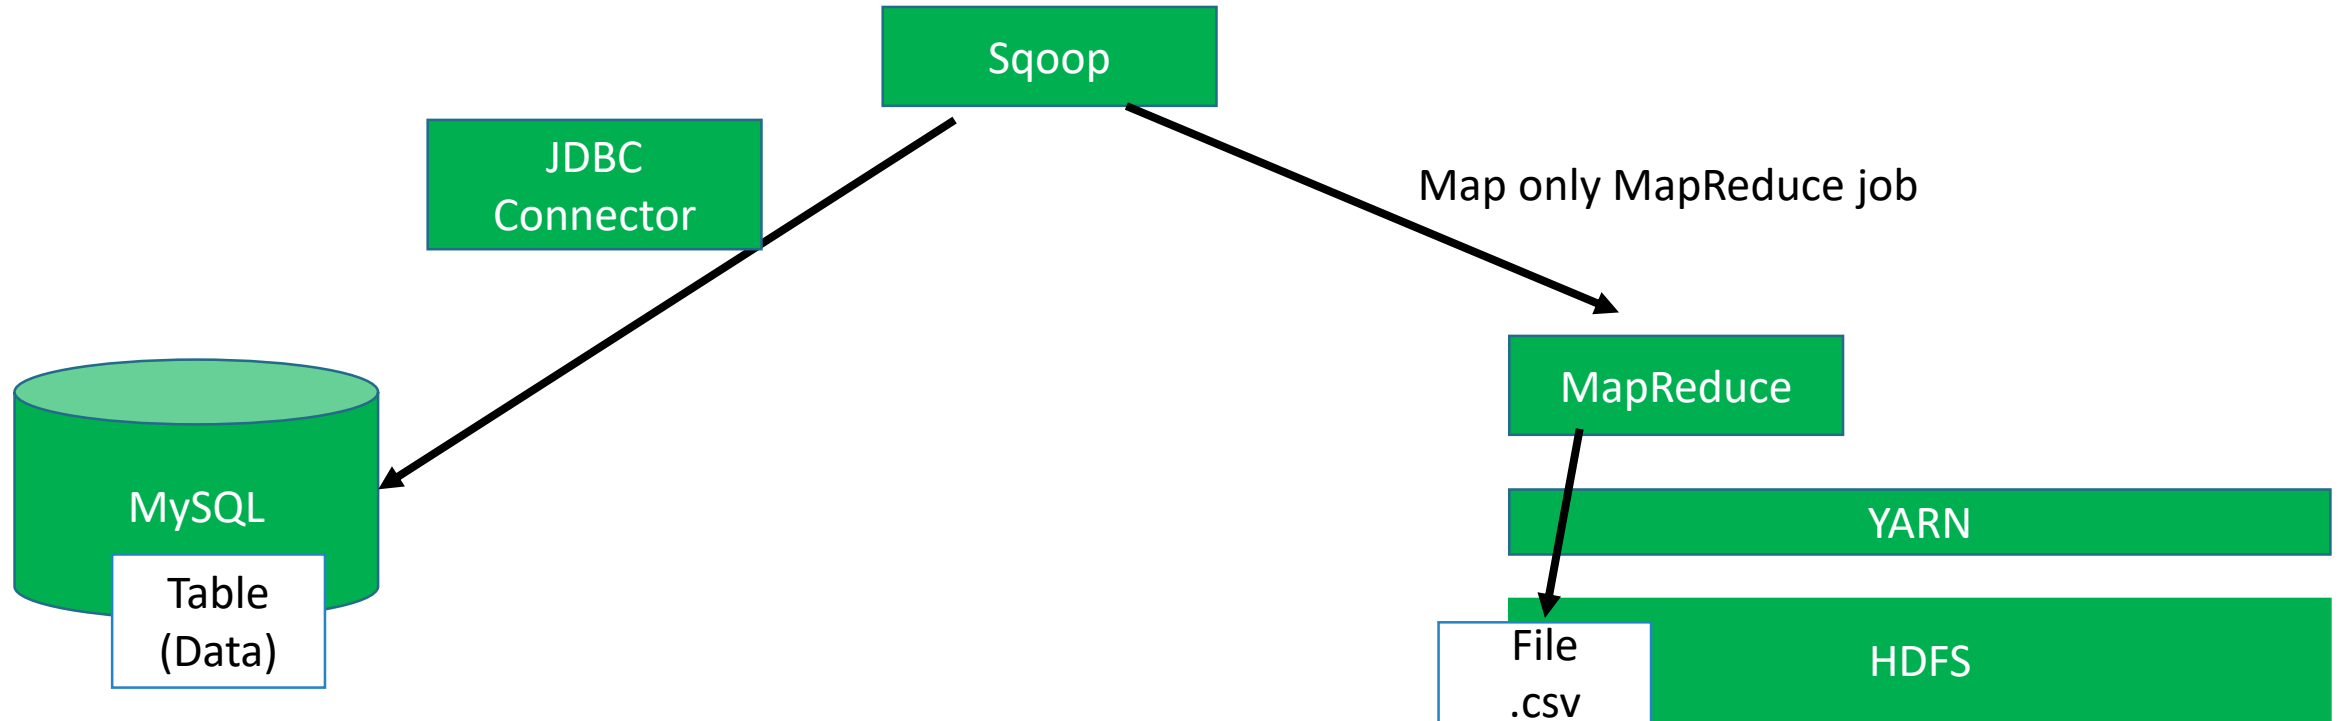
Apache Sqoop

<http://sqoop.apache.org/>

Apache Sqoop is a tool designed for efficiently transferring bulk data between Apache Hadoop and structured data stores such as relational databases

Sqoop Import → From RDBMS to HDFS

Sqoop Export → From HDFS to RDBMS



Agenda – Day 3

- Advanced MapReduce
- Data Ingestion → Flume
- Hive
- Advanced Hive

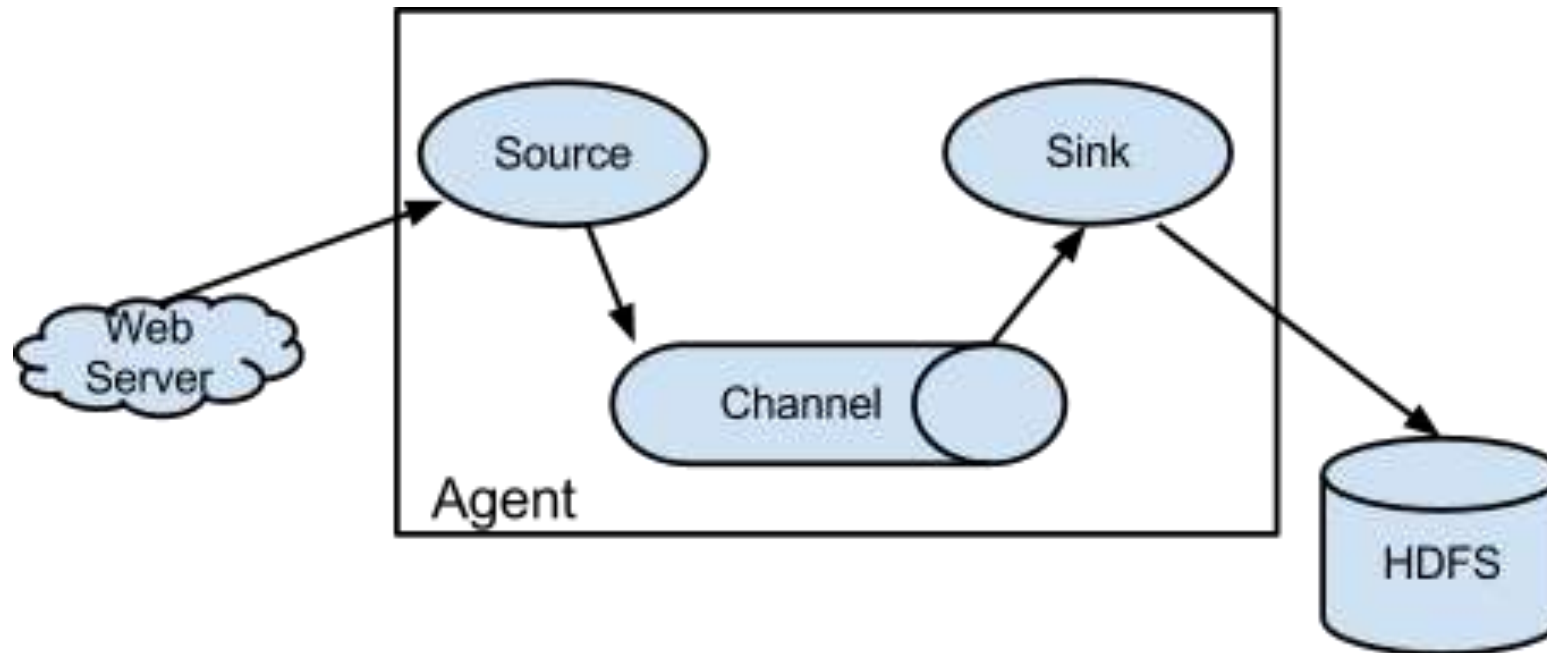
- Advanced MapReduce
 - Combiners & Partitioners
 - Chaining Multiple Maps
 - Joining datasets in MR
 - Working with images, sequence files
 - MapReduce datatypes

Apache Flume

<http://flume.apache.org/>

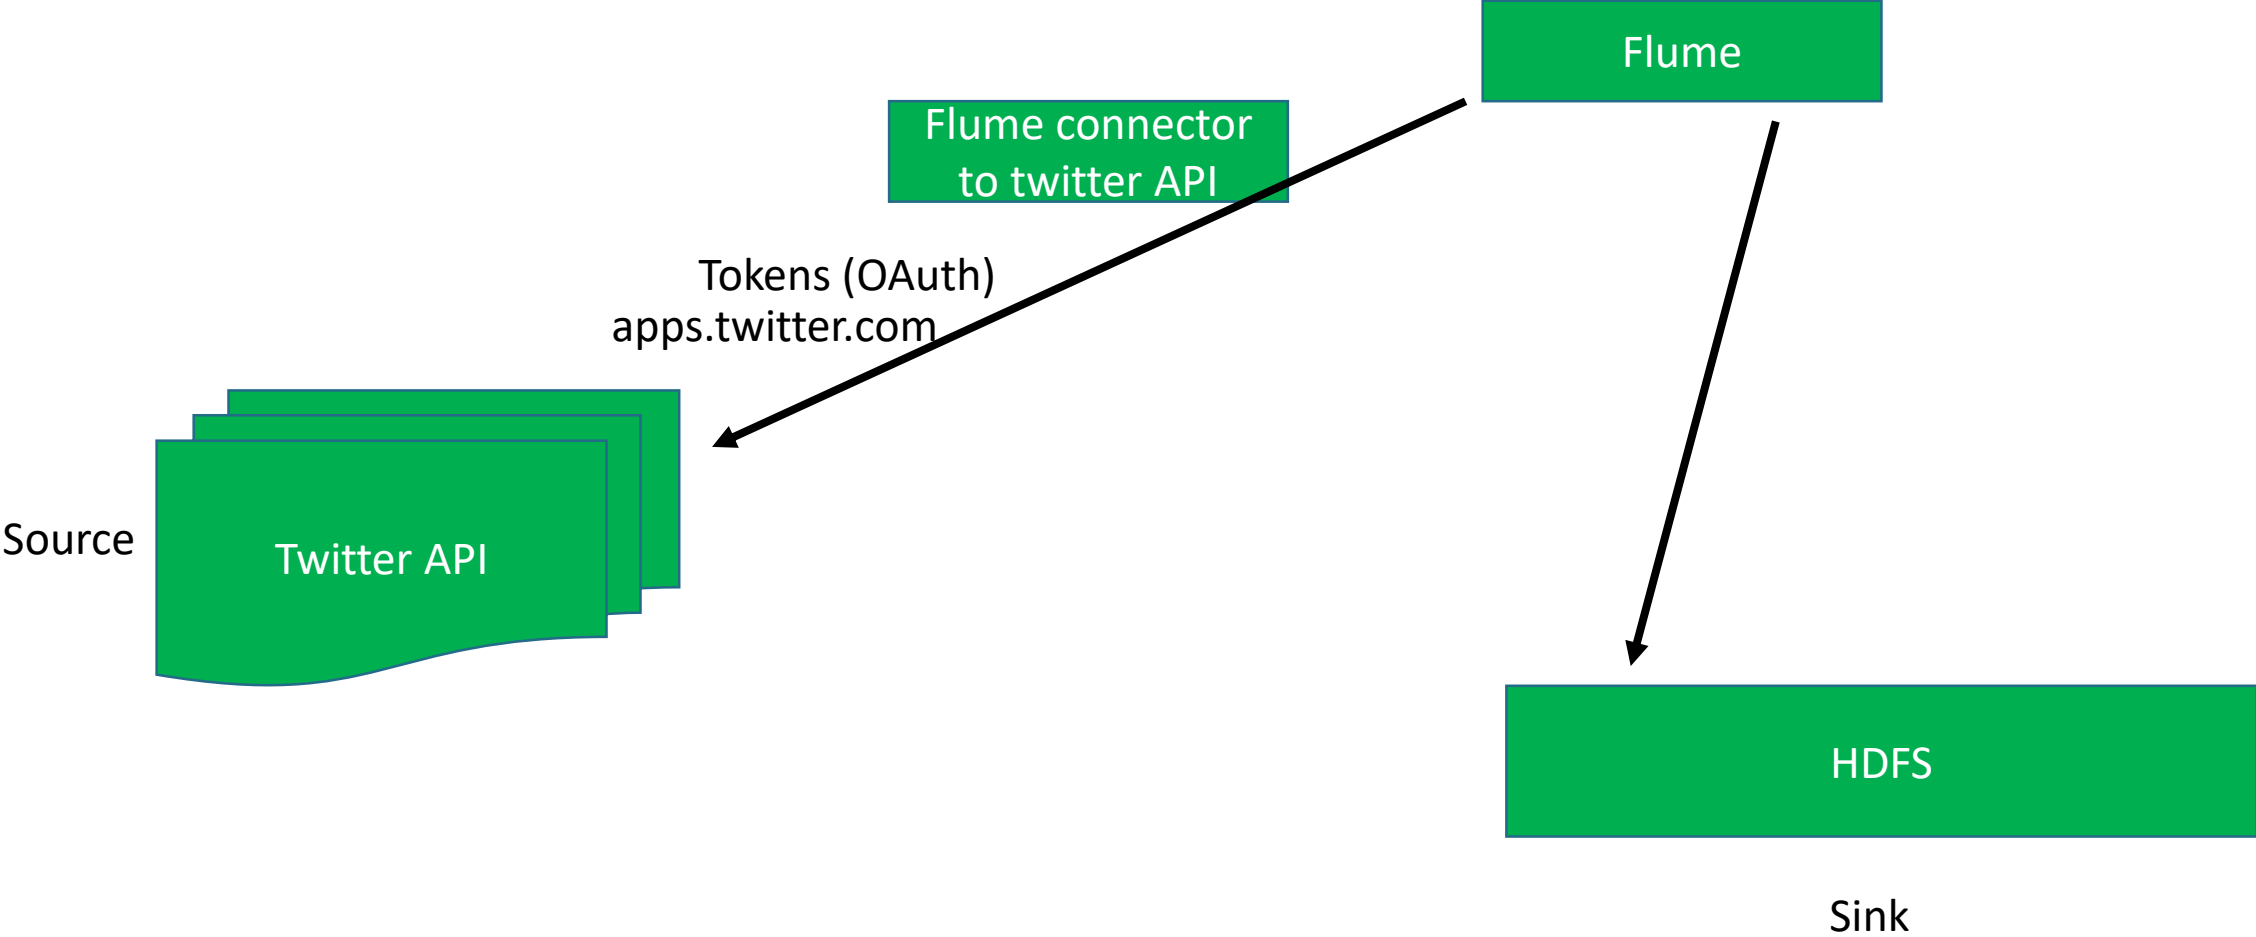
Flume is a distributed service for efficiently collecting and moving large amounts of log data

Example:



<http://www.cloudera.com/documentation/kafka/latest.html>

Apache Flume - Twitter



Apache Hive

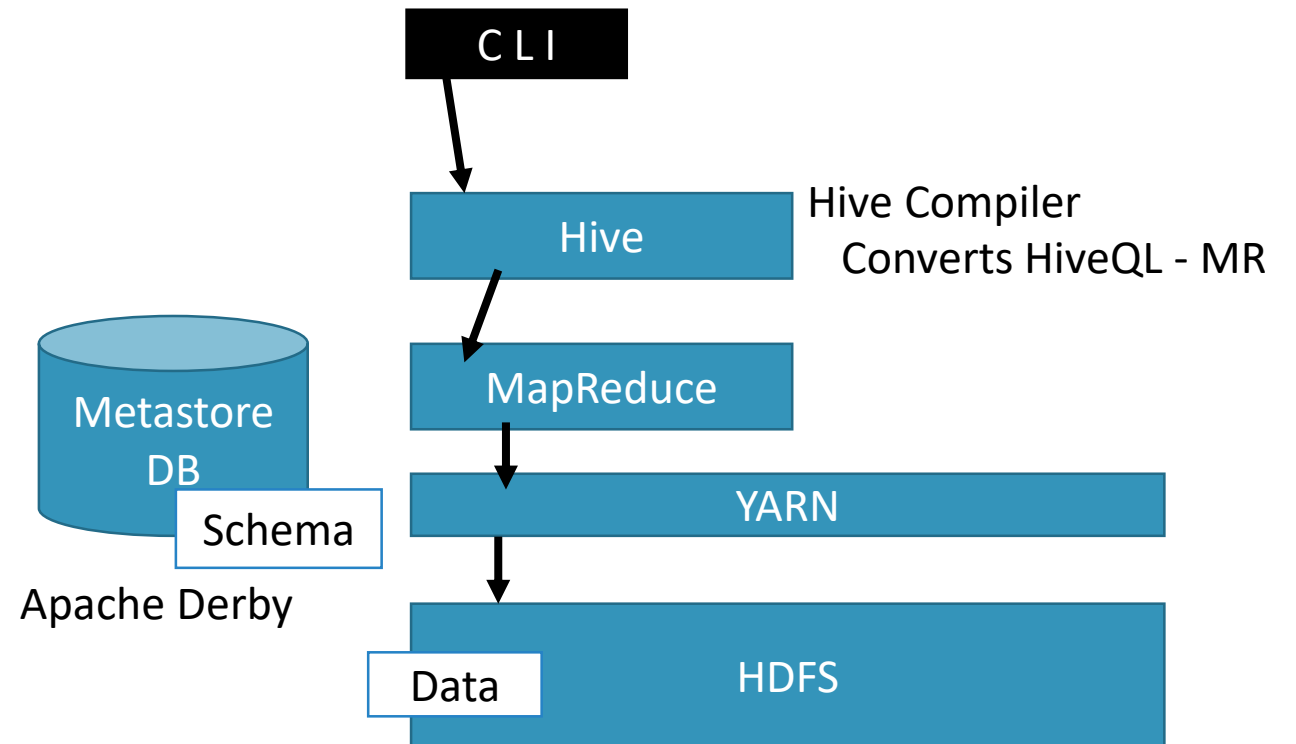
<http://hive.apache.org/>

- The Apache Hive data warehouse software facilitates reading managing large datasets residing in distributed storage using SQL
- Structure can be projected onto data already in storage → Schema on Read

/user/hive/warehouse → Hive's default warehouse directory

Hive Tables

- Managed Tables → Default → Hive manages data & table
- External Tables → External Keyword → Hive manages table



HDFS → Hadoop File System (Read Only)

MapR FS → MapR Distribution (RW)

KUDU → Cloudera Distribution (RW)

Further Reading

<http://hadoop.apache.org/>

<http://hortonworks.com/blog/impala-vs-hive-performance-benchmark/>

<https://cwiki.apache.org/confluence/display/Hive/LanguageManual>