

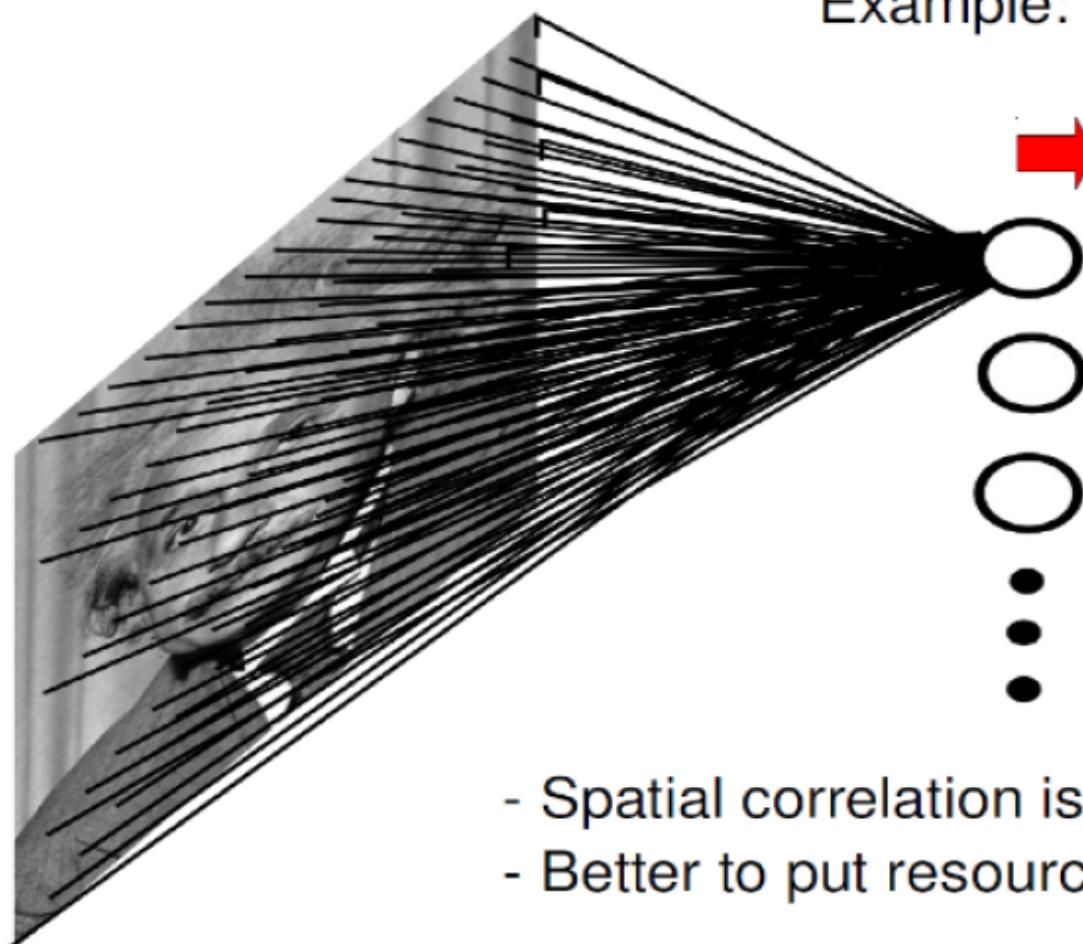
Convoluted Neural Network Internals

By Mohit Kumar

Convulated Neural Network:Why?

Why Convolution?

FULLY CONNECTED NEURAL NET



Example: 1000x1000 image

1M hidden units

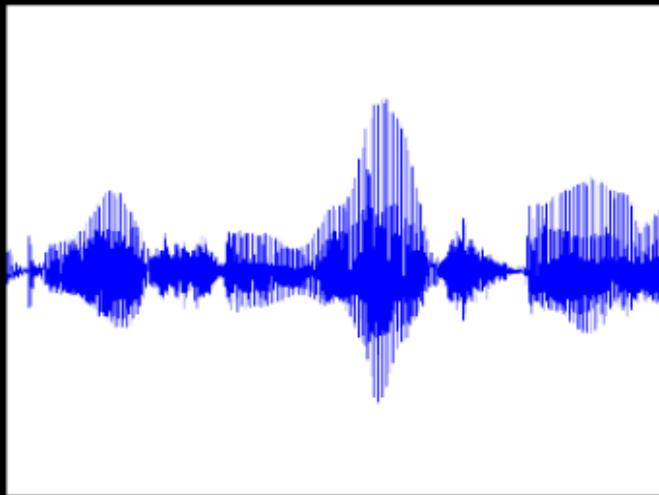
→ **10¹² parameters!!!**

- Spatial correlation is local
- Better to put resources elsewhere!

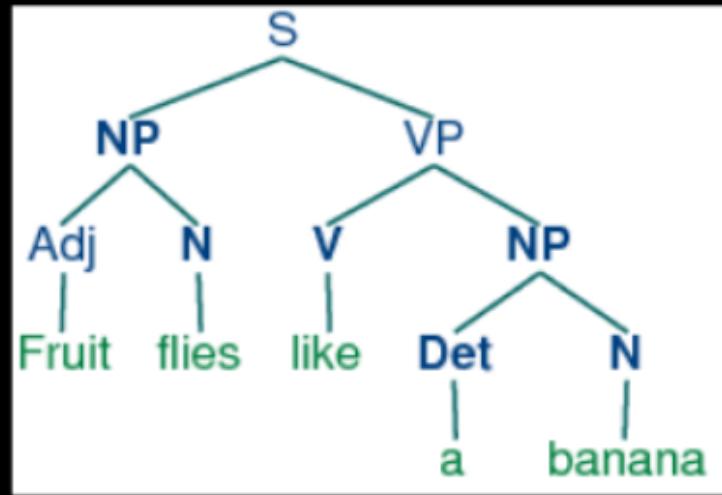
Convolved Neural Network: Why?: Detection or classification



Computer vision

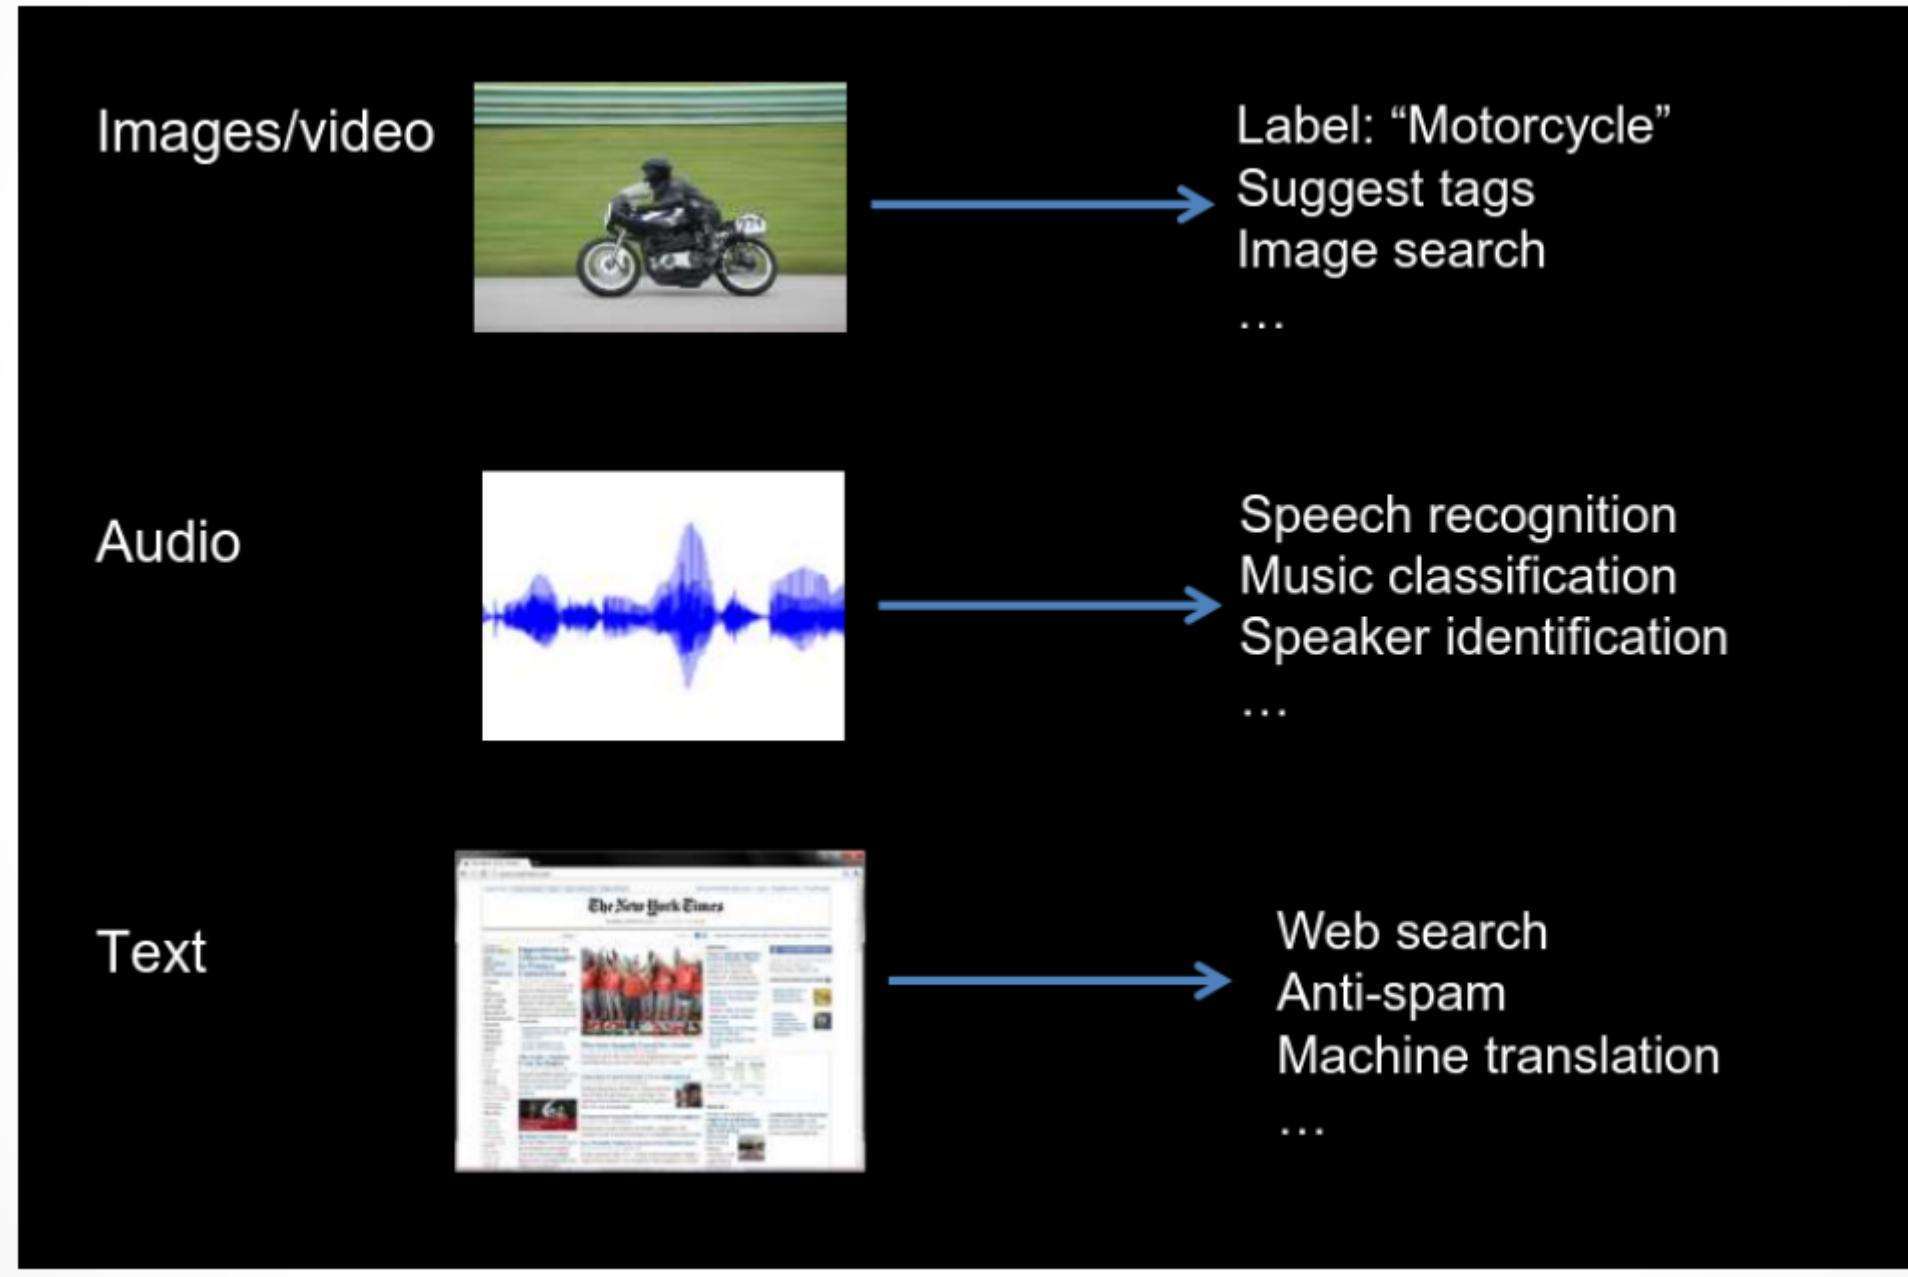


Audio



Text

Convolved Neural Network: Why?: What to do with this data?



Convulated Neural Network:Why?:Feature Selection



Input



Learning
algorithm

Convolved Neural Network:Why?:Feature Selection



Input

Feature
Representation

E.g., SIFT, HoG, etc.

Learning
algorithm

Convolved Neural Network: Why?: How is computer Perception done?

Object detection



Image



Vision features

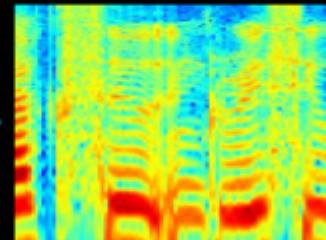


Detection

Audio classification



Audio



Audio features

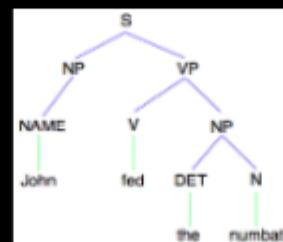


Speaker ID

NLP



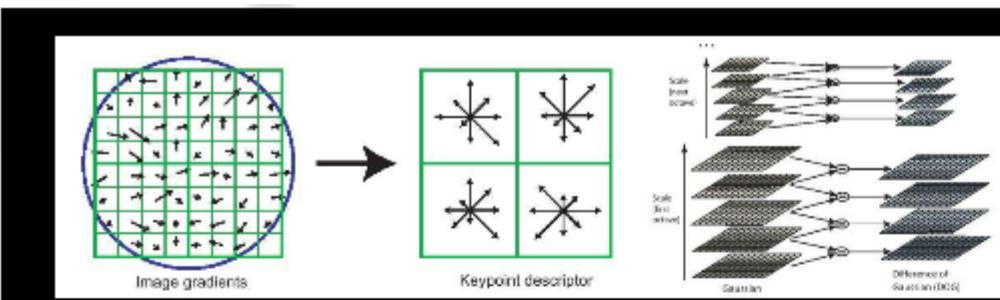
Text



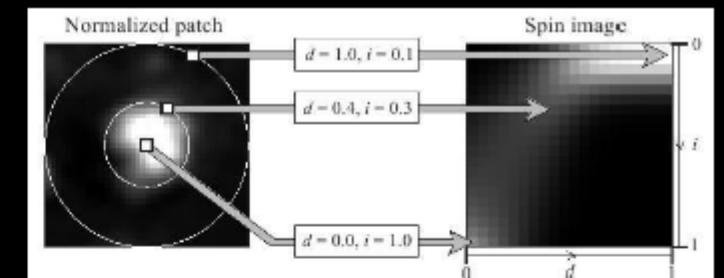
Text features

Text classification,
Machine translation,
Information retrieval,
etc.

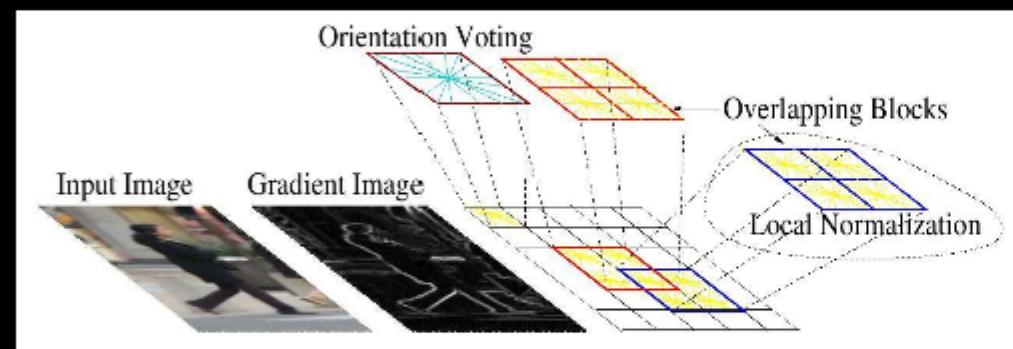
Convolved Neural Network: Why?: Computer Vision features



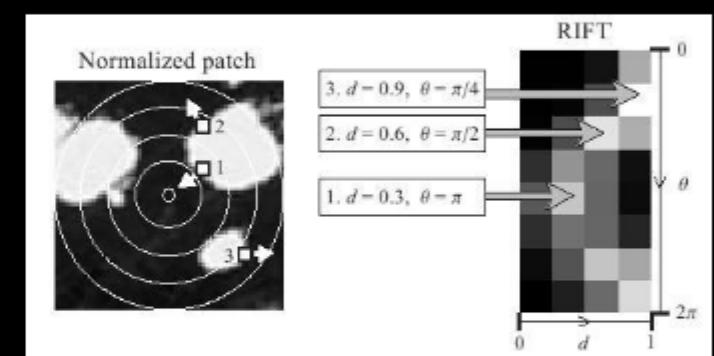
SIFT



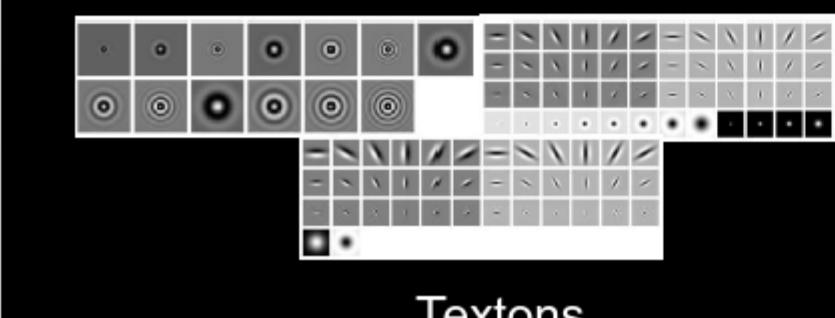
Spin image



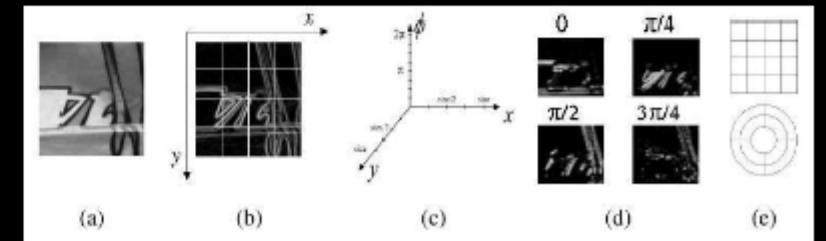
HoG



RIFT

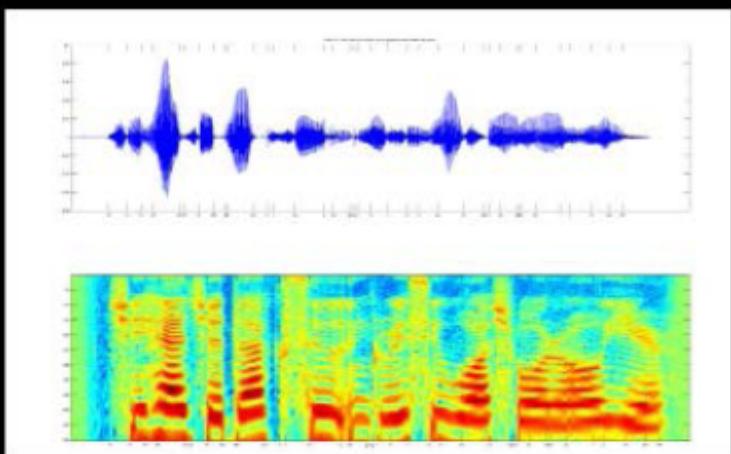


Textons

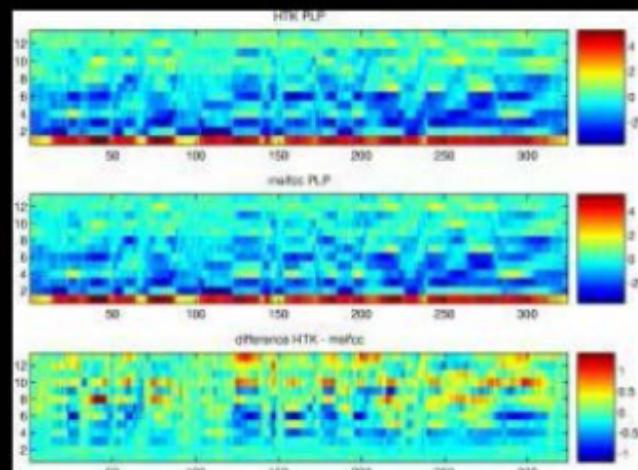


GLOH

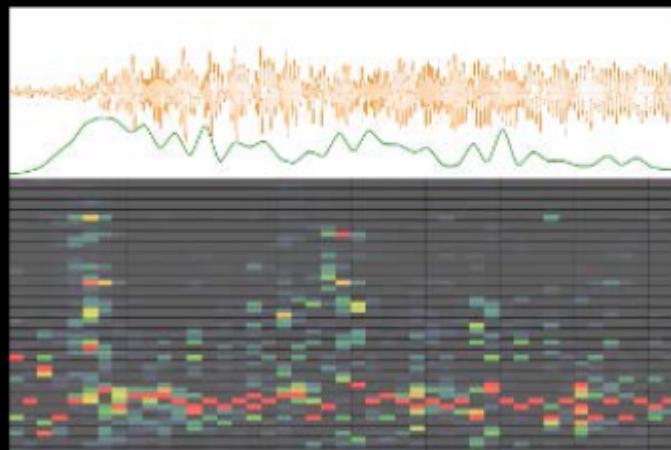
Convolved Neural Network: Why?: Audio Features



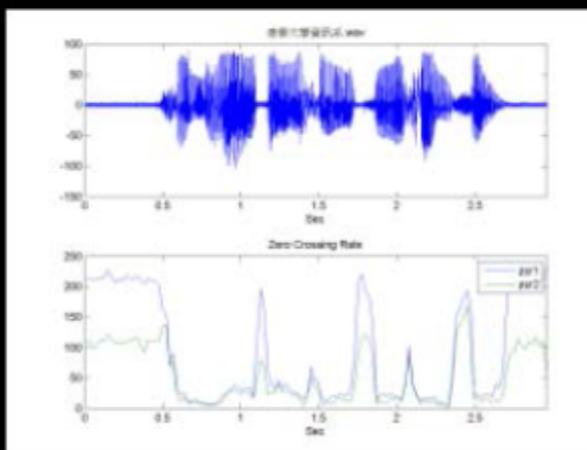
Spectrogram



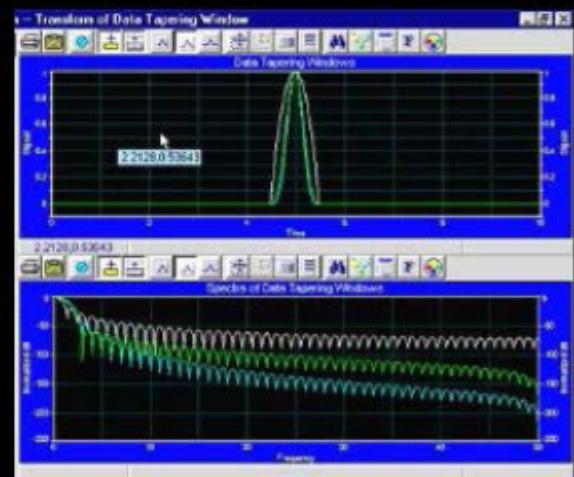
MFCC



Flux

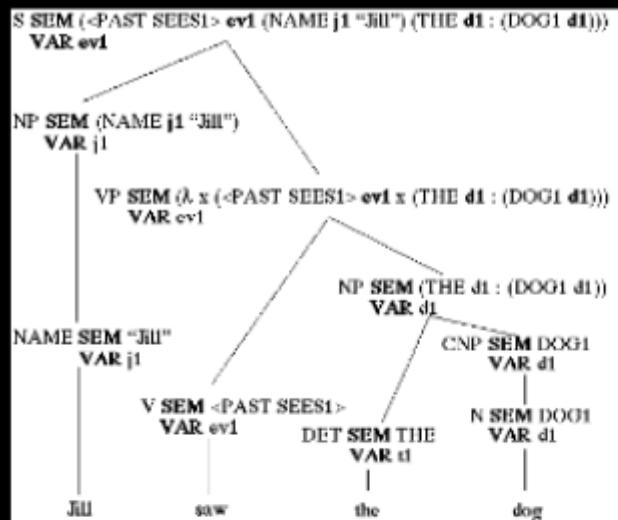


ZCR



Rolloff

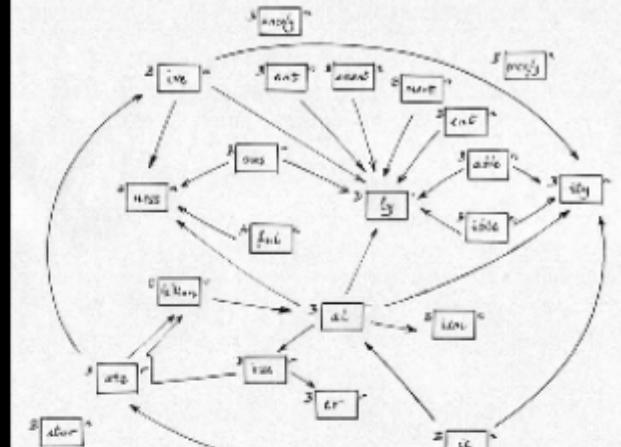
Convolved Neural Network: Why?: NLP Features



Parser features

<DOC>
<DOCID> waj94_008_0212 </DOCID>
<DOCNO> 940413-0062. </DOCNO>
<HL> Who's News:
<H1> Burns Fry Ltd. </H1>
<DD> 04/13/94 </DD>
<SO> WALL STREET JOURNAL (J), PAGE B10 </SO>
<CO> MER </CO>
<IND> SECURITIES (SCR) </IND>
<TXT>
<p> BURNS FRY LTD., **Toronto** -- Donald Wright, 46 years old, was named executive vice president and director of fixed income at this brokerage firm. Mr. Wright resigned as president of **Merrill Lynch Canada Inc.**, a unit of **Merrill Lynch & Co.**, to succeed Mark Bassiner, 48, who left Burns Fry last month. A **Merrill Lynch** spokeswoman said it hasn't named a successor to Mr. Wright, who is expected to begin his new position by the end of the month.
</p>
</TXT>
</DOC>

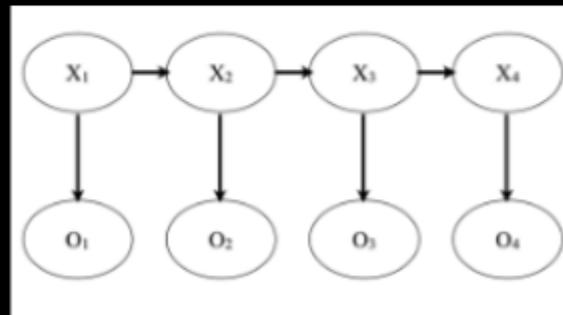
NER/SRL



Stemming

His father, Nick Begich, **won an election**
posthumously, only they didn't know for sure that **it**
was posthumous because **his plane** just disappeared.
It still hasn't turned up. **It's** why locators are now
required in all US planes.

Anaphora



POS tagging



Figure 2. 'win' relation example

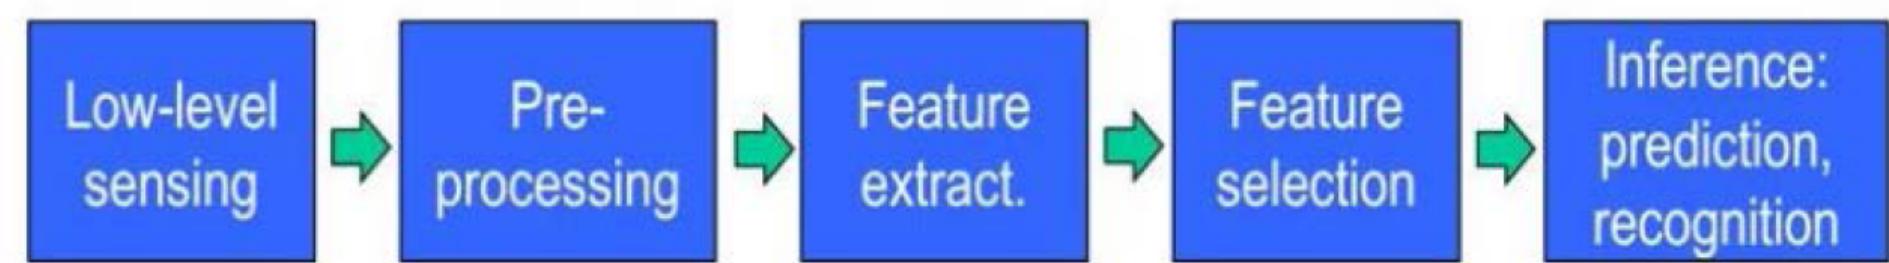
WordNet features

Convoluted Neural Network:Why?:Challenges

- Coming up with features is difficult, time consuming and requires expert knowledge in the domain.
- A lot of time is spent tuning features which are often handcrafted.

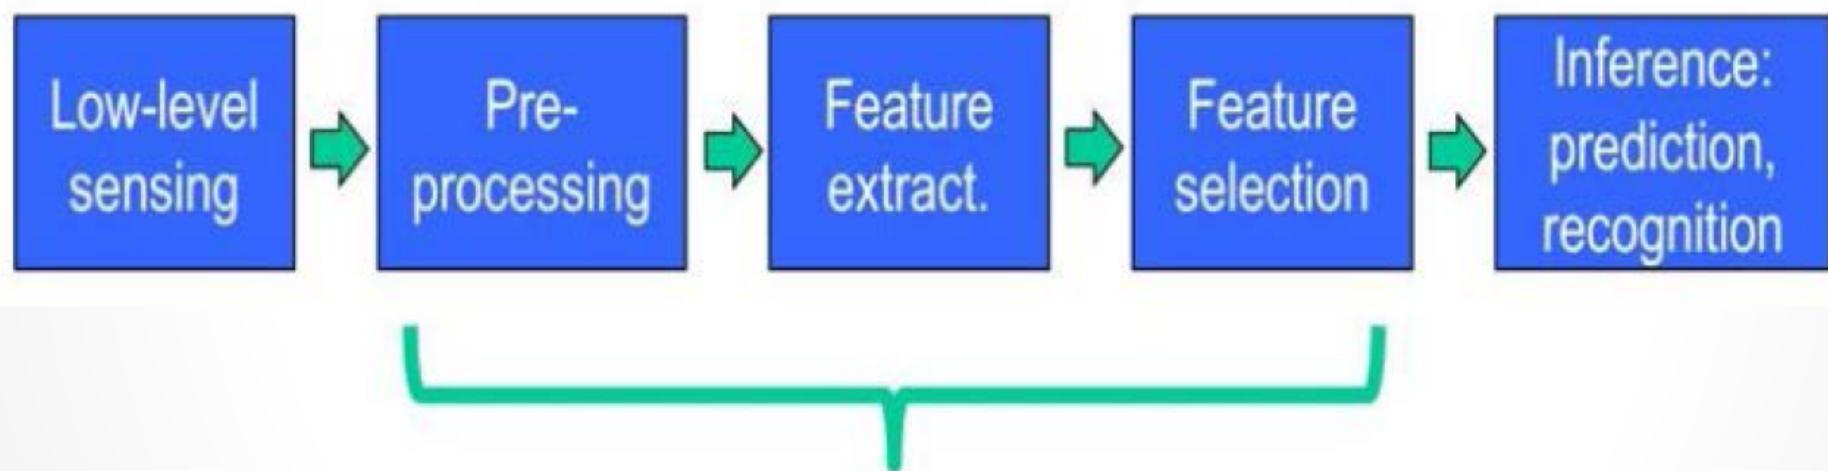
Convolved Neural Network:Why?:Feature Representation

Machine Learning



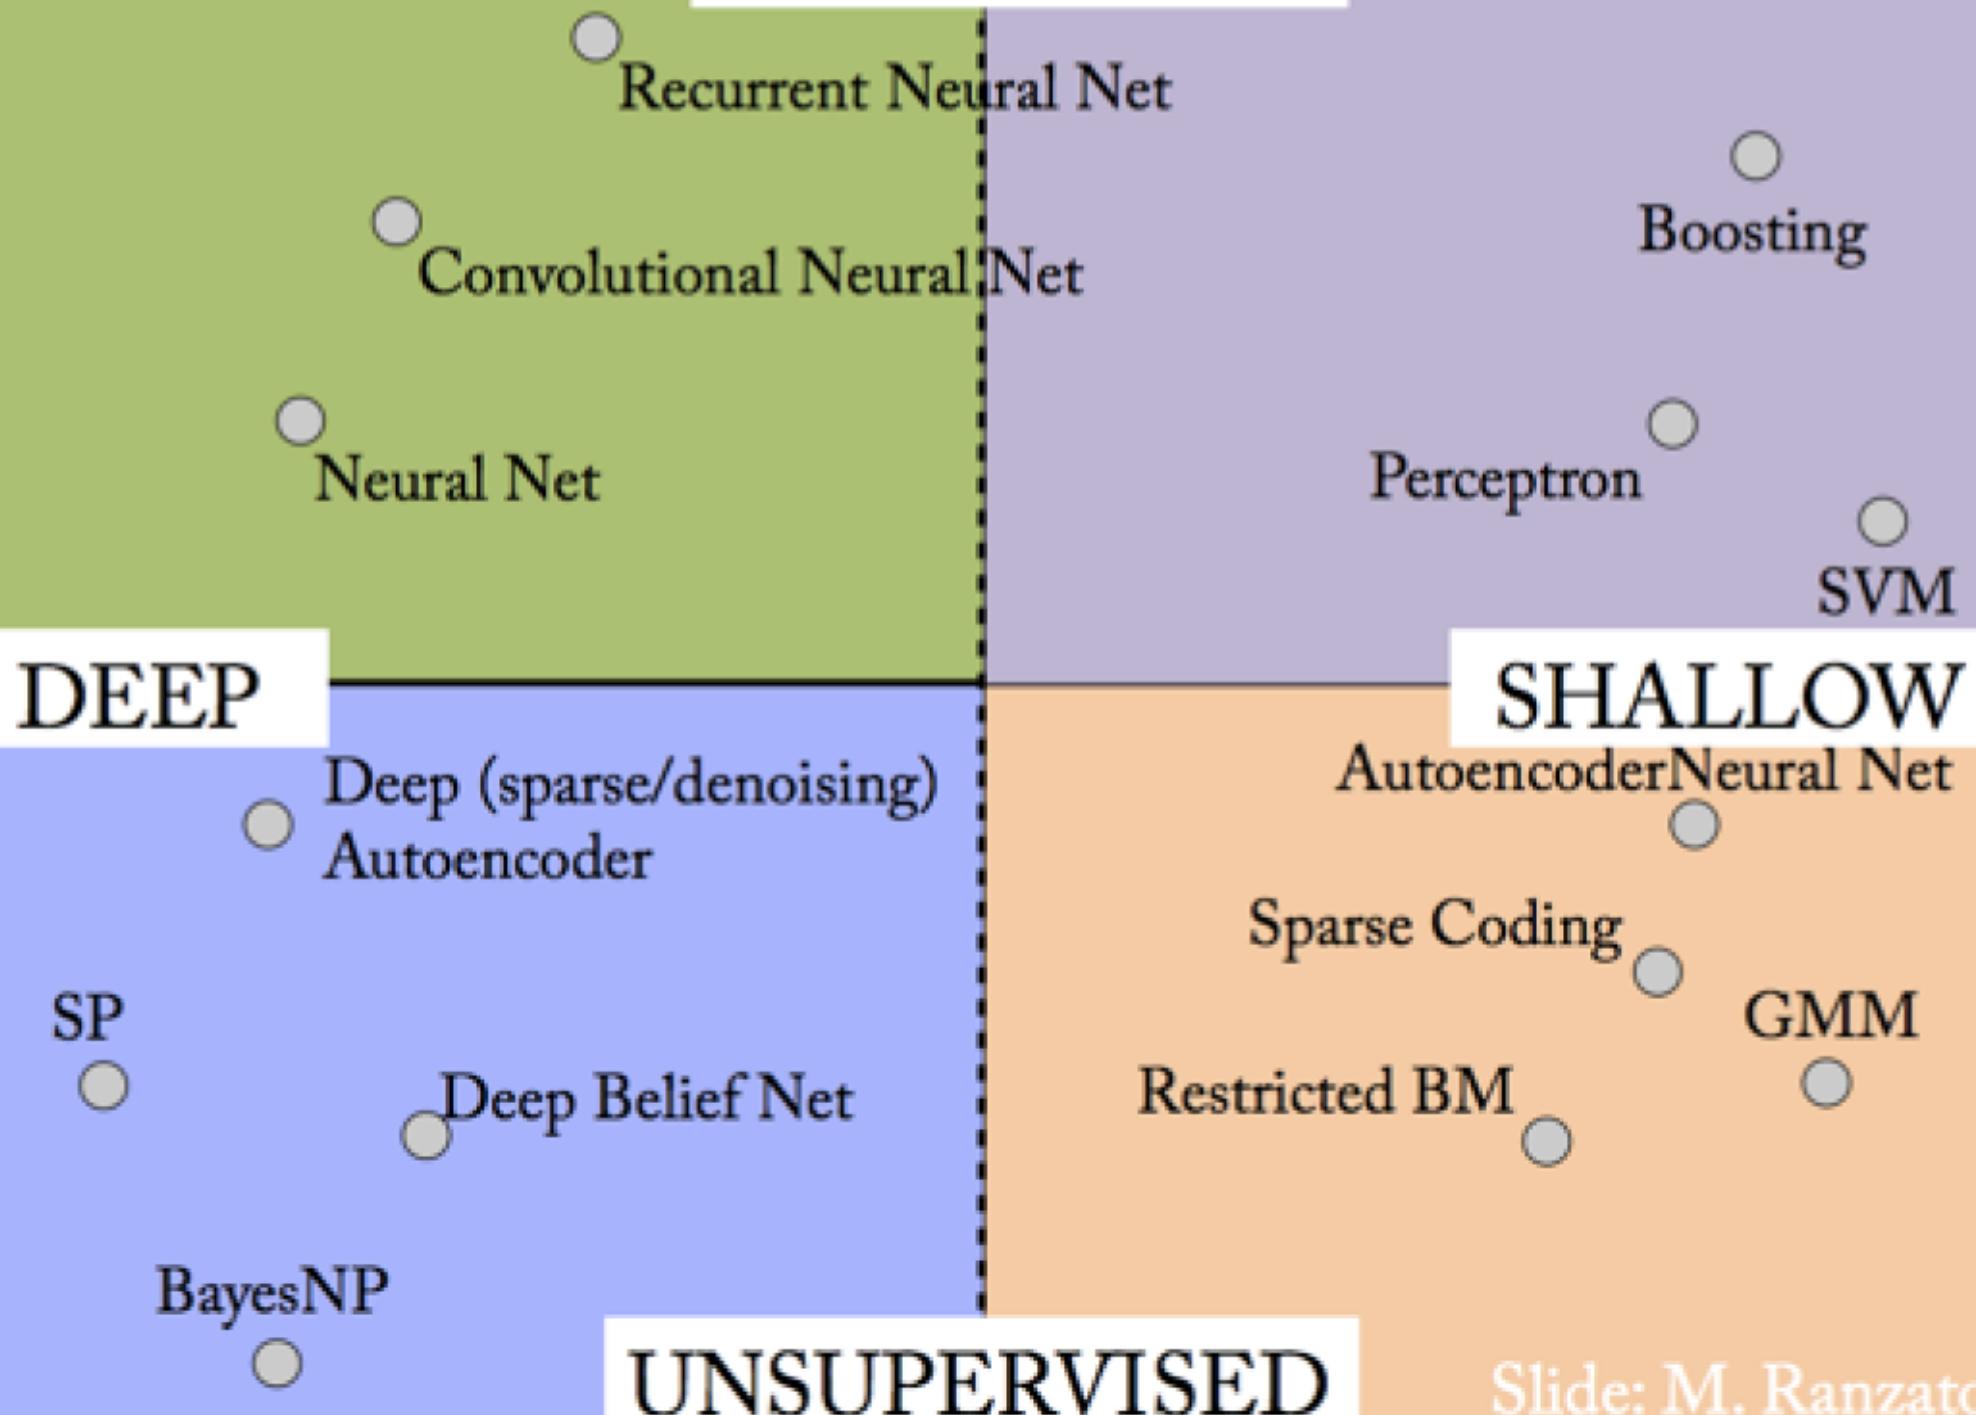
Convolved Neural Network: Why?: Feature Representation

Machine Learning



Feature Learning:
instead of design features,
let's design feature learners

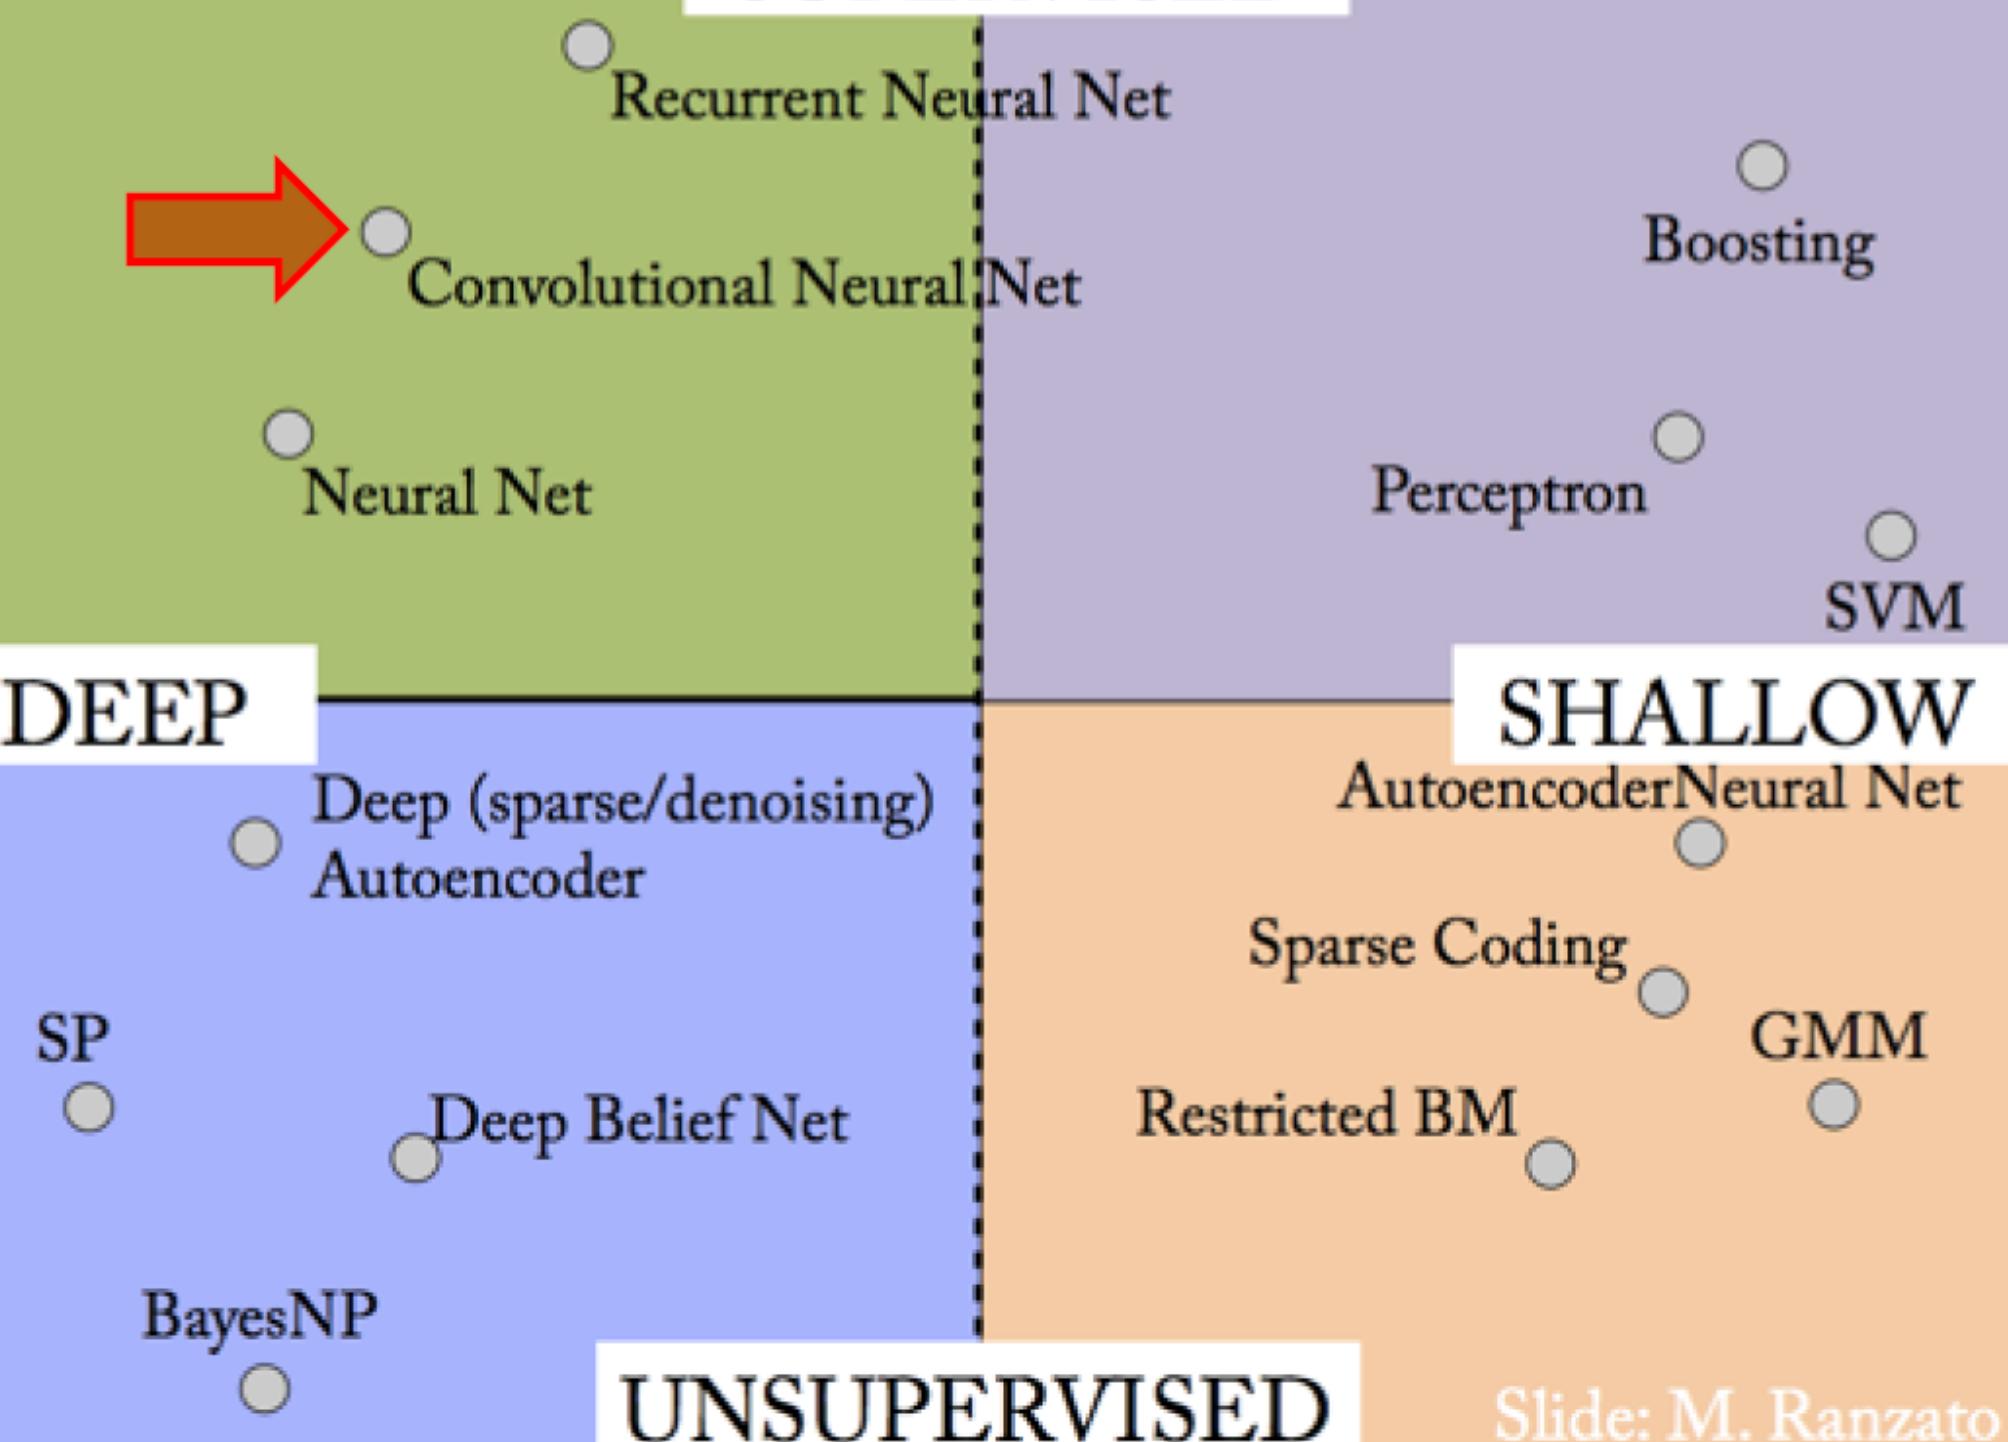
SUPERVISED



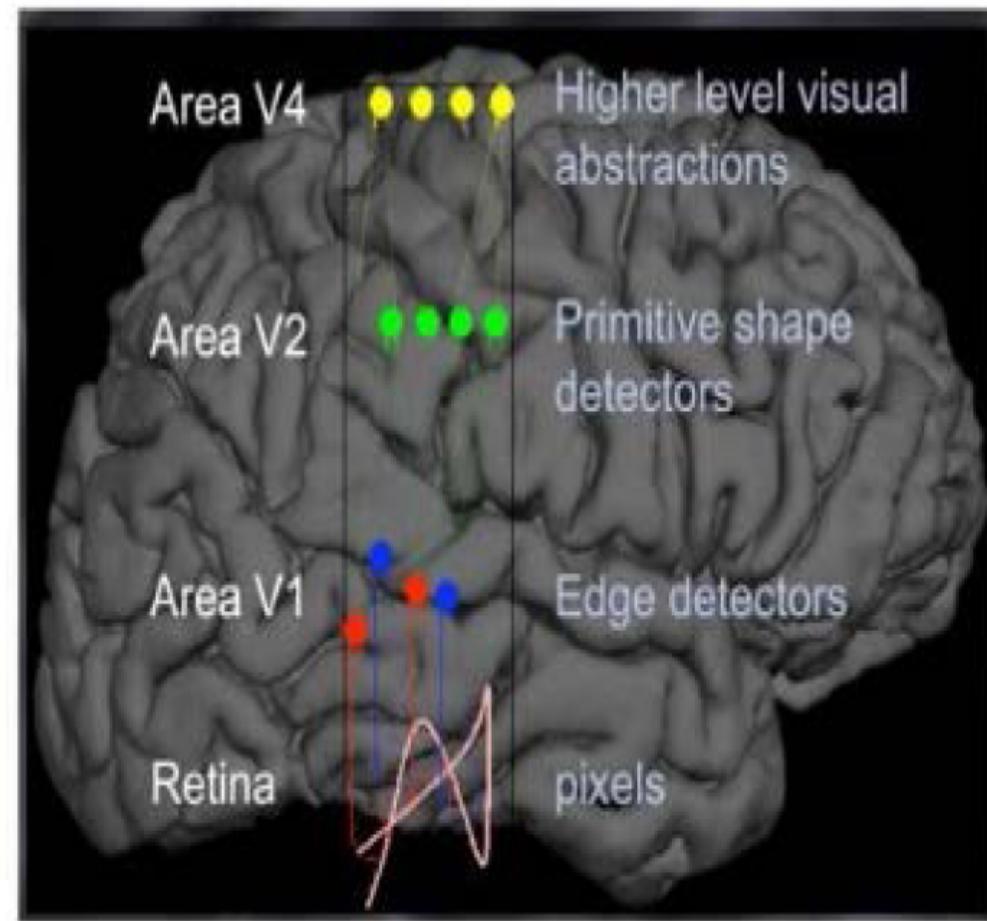
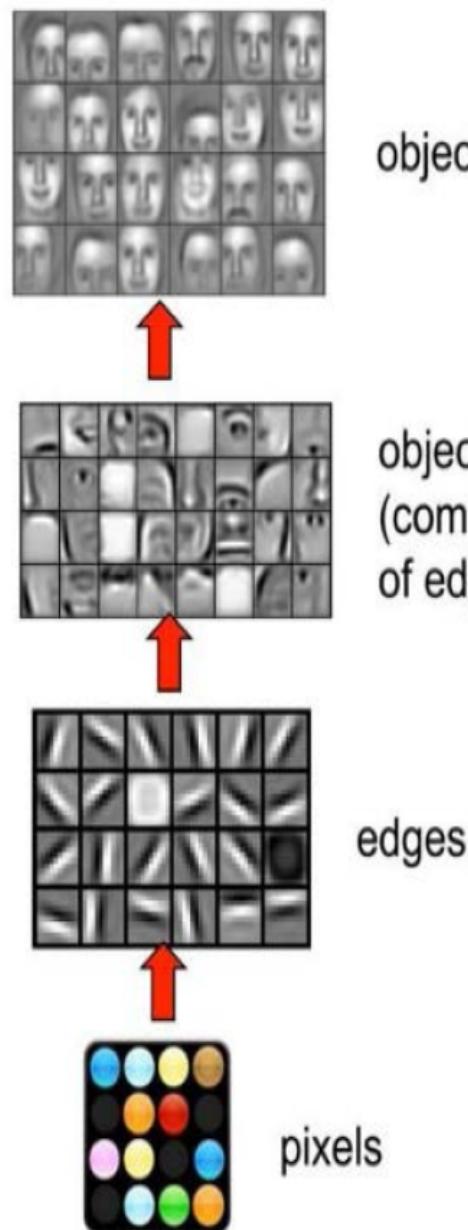
UNSUPERVISED

Slide: M. Ranzato

SUPERVISED



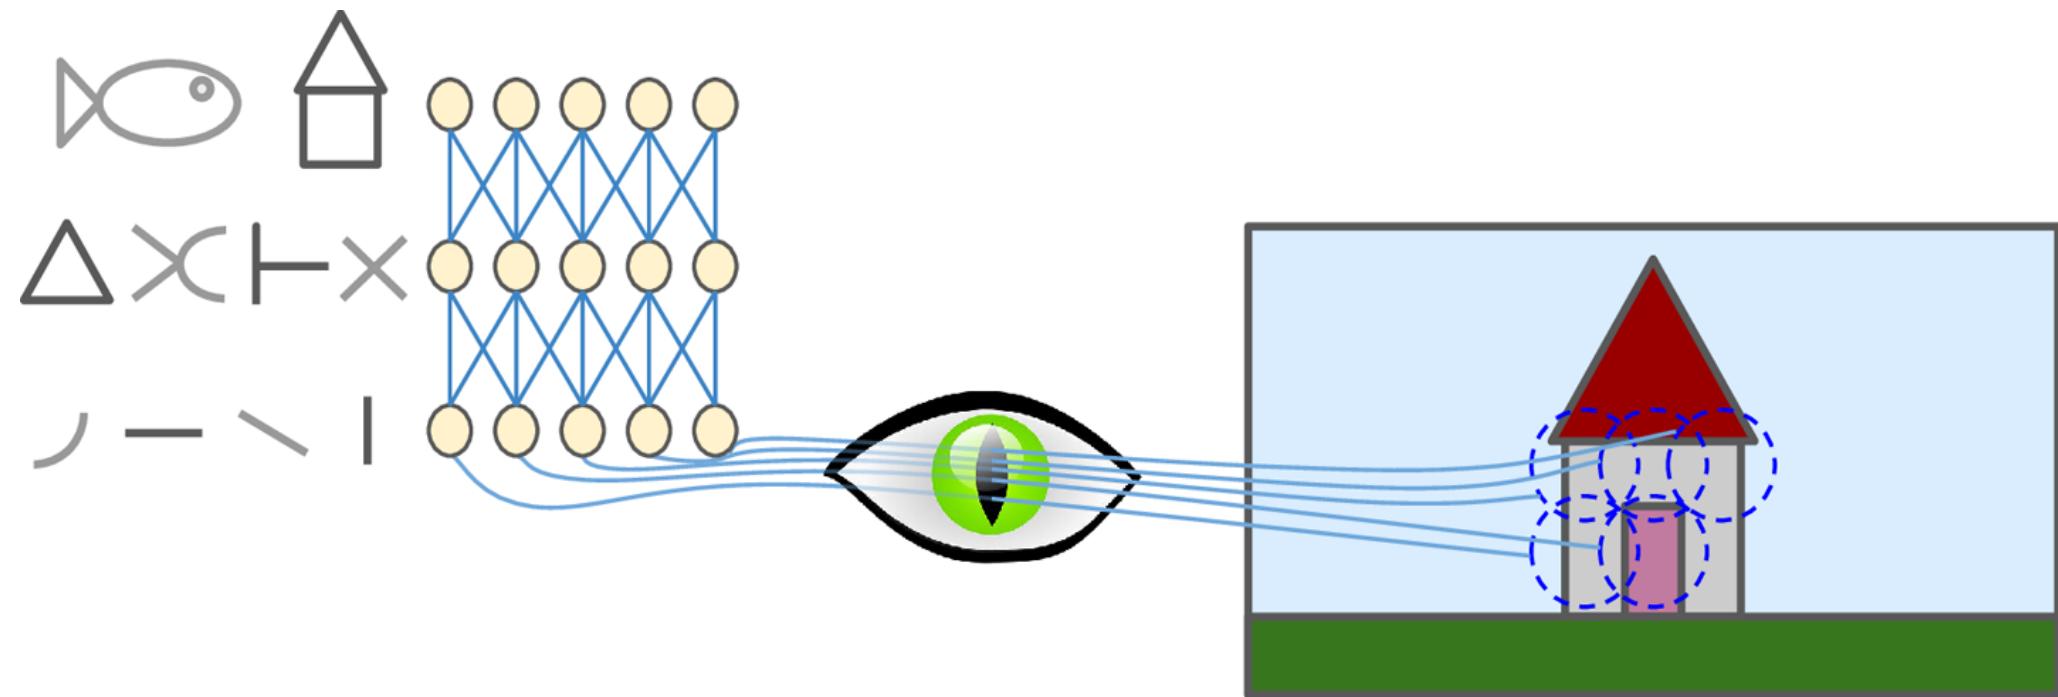
Convolved Neural Network: Why?: Biological Inspiration



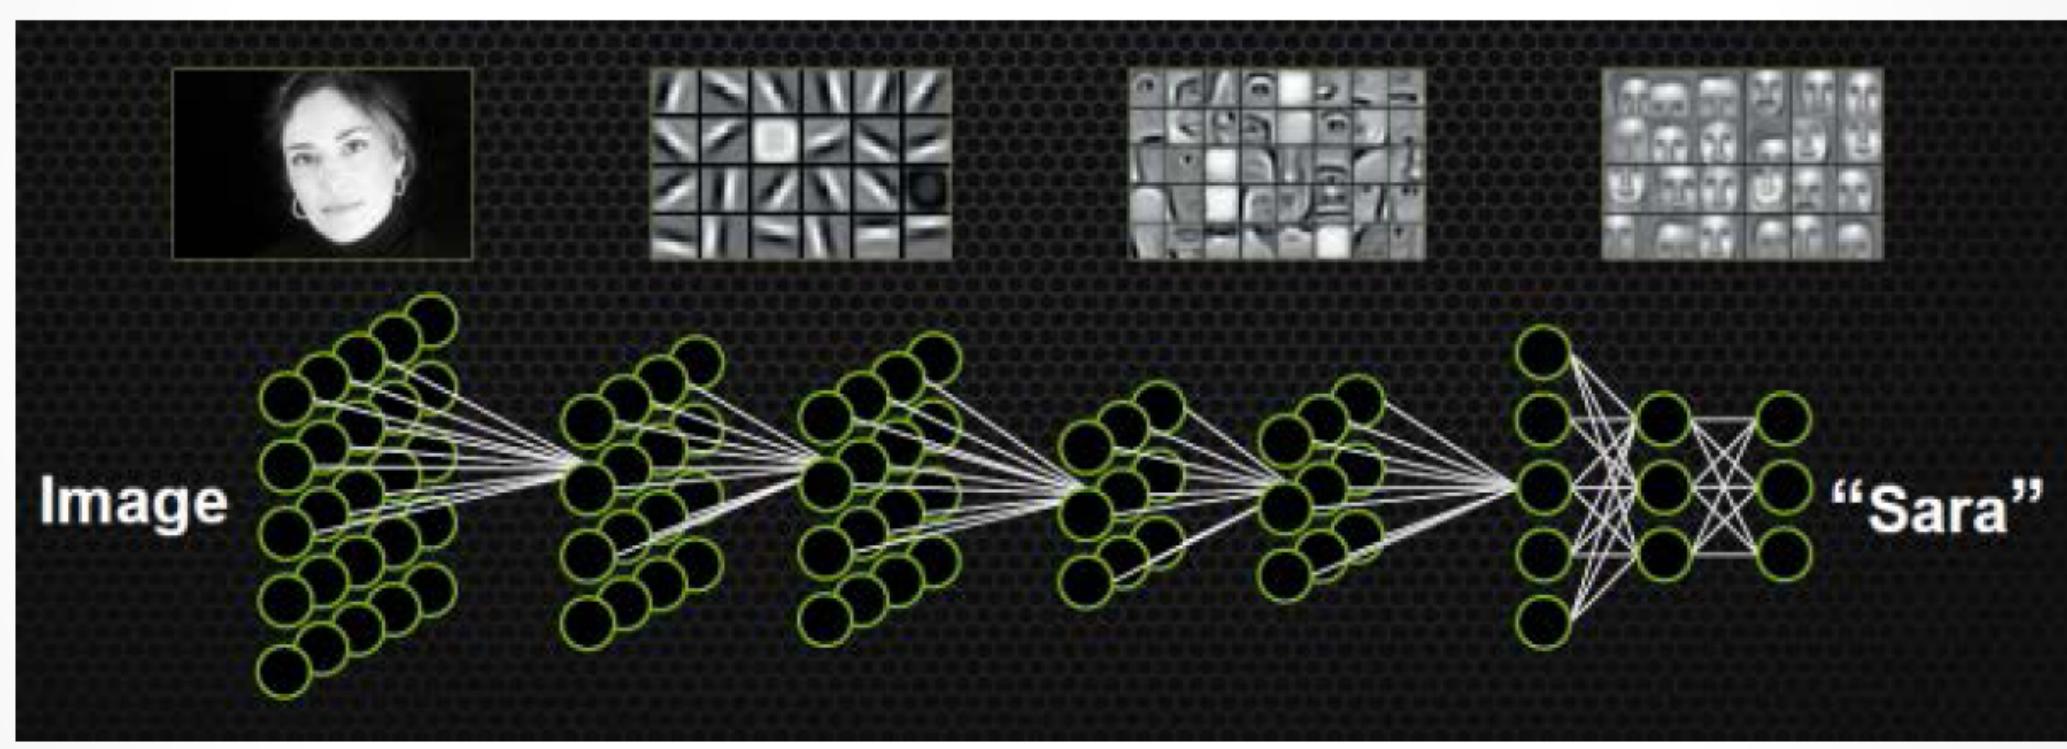
Convolved Neural Network:Why?Biological Inspiration

- David H. Hubel and Torsten Wiesel performed a series of experiments on cats in 19581 and 19592 (and a few years later on monkeys3), giving crucial insights on the structure of the visual cortex (the authors received the Nobel Prize in Physiology or Medicine in 1981 for their work).
- In particular, they showed that many neurons in the visual cortex have a small local receptive field, meaning they react only to visual stimuli located in a limited region of the visual field .
- The receptive fields of different neurons may overlap, and together they tile the whole visual field. Moreover, the authors showed that some neurons react only to images of horizontal lines, while others react only to lines with different orientations (two neurons may have the same receptive field but react to different line orientations).
- They also noticed that some neurons have larger receptive fields, and they react to more complex patterns that are combinations of the lower-level patterns. These observations led to the idea that the higher-level neurons are based on the outputs of neighboring lower-level neurons . This powerful architecture is able to detect all sorts of complex patterns in any area of the visual field.

Convulated Neural Network: Why? Biological Inspiration



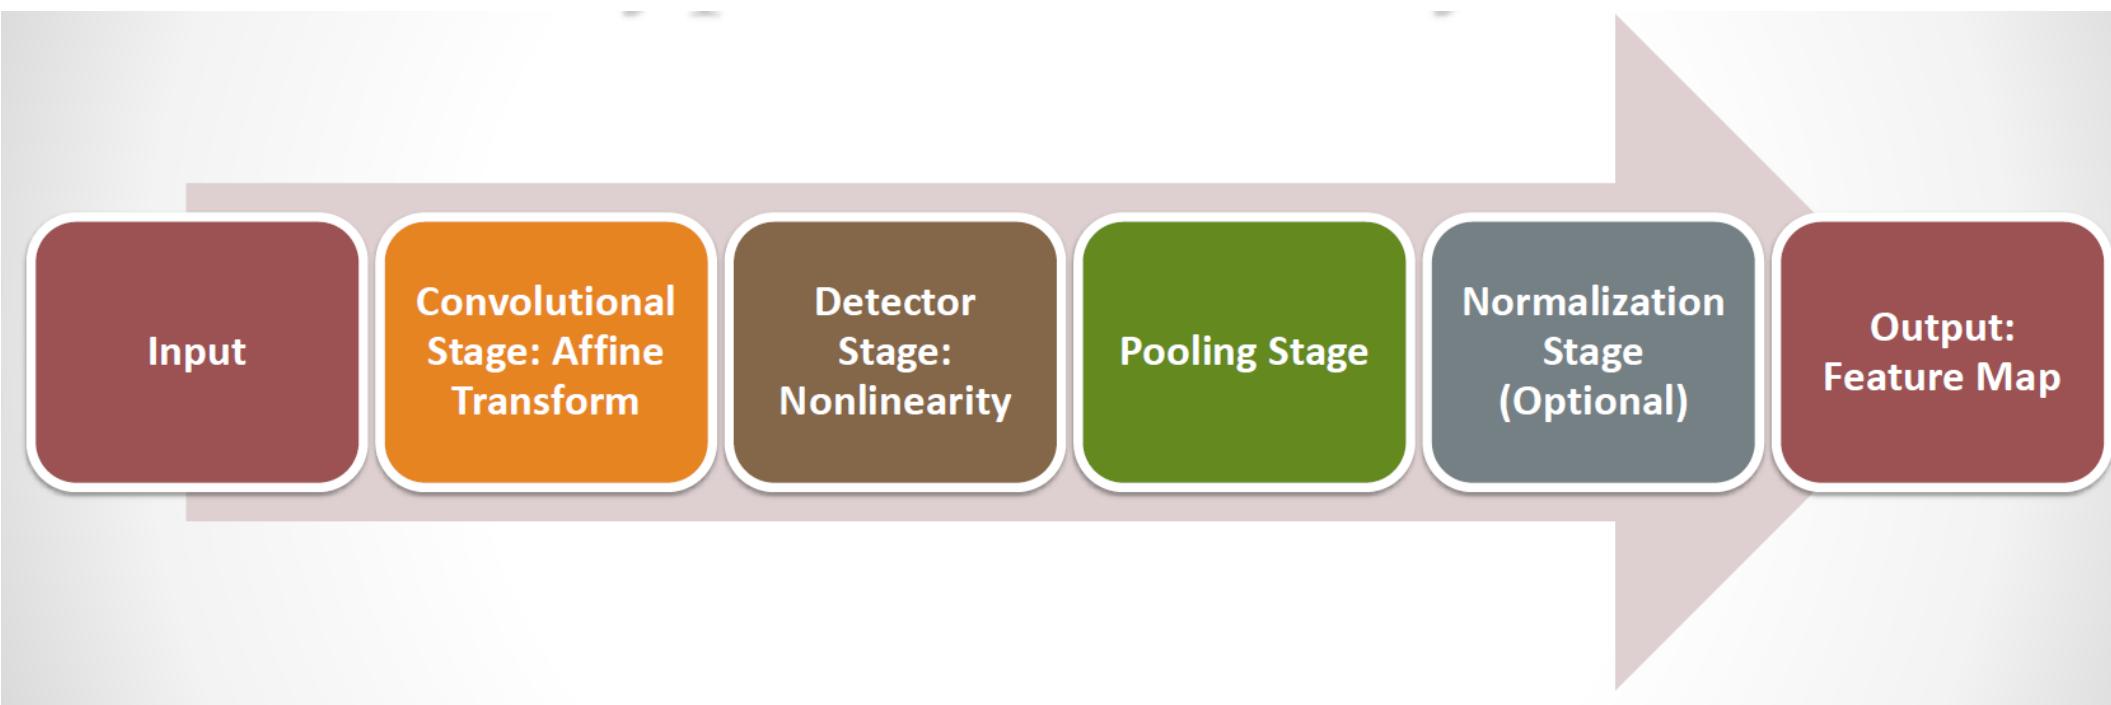
Convolved Neural Network: Why? Features learnt by Deep training



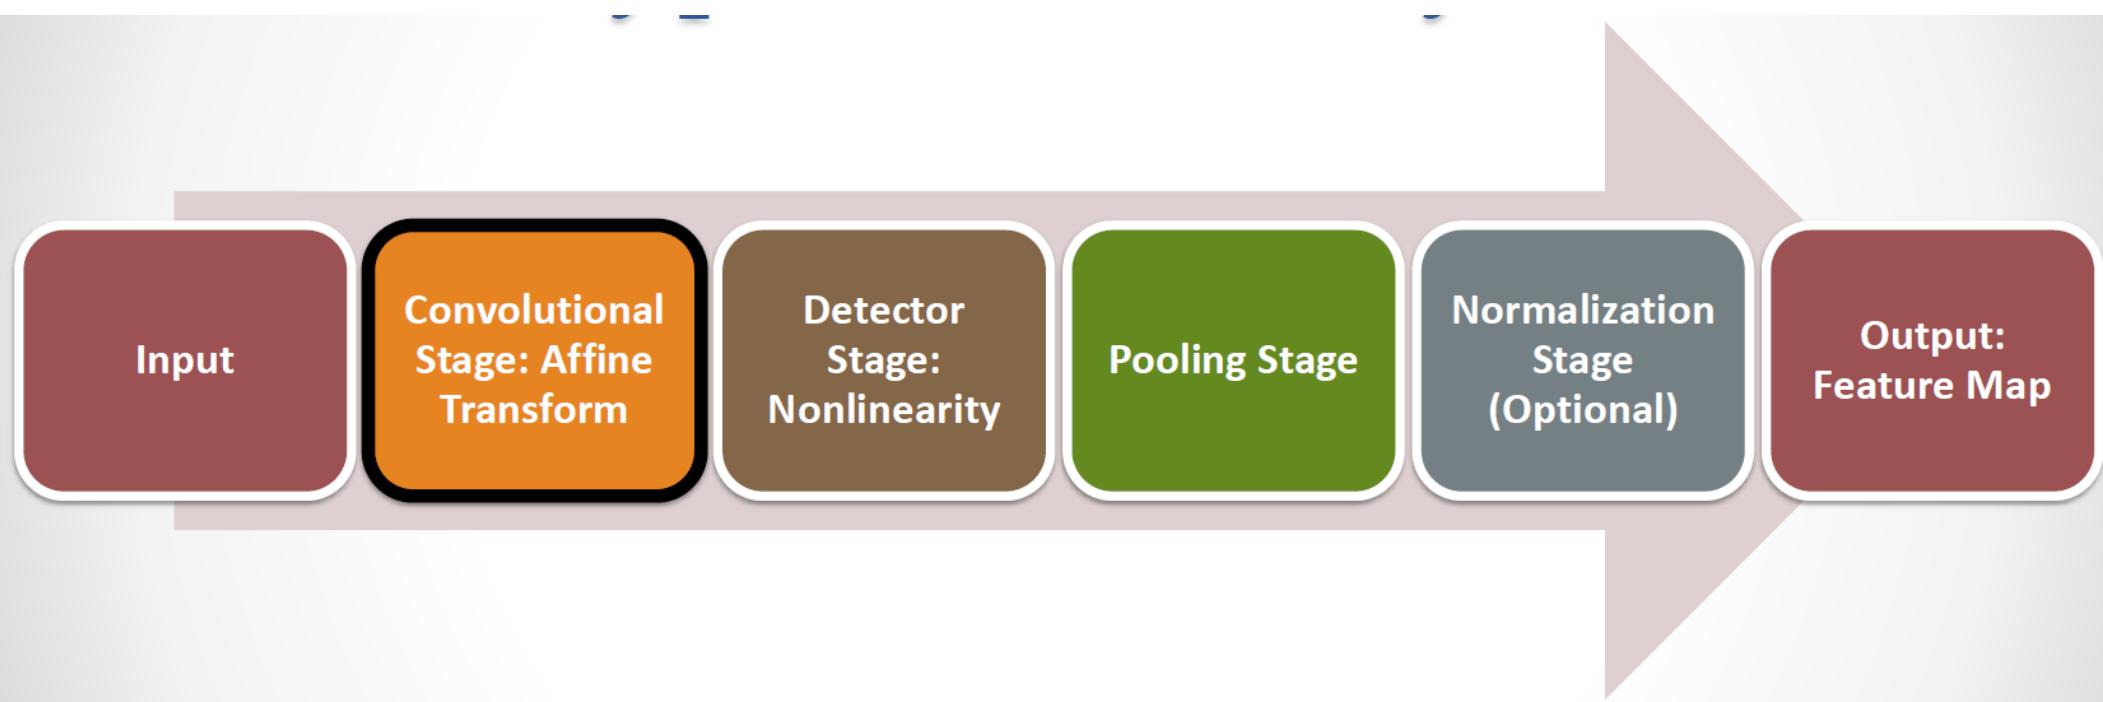
Convoluted Neural Network:What is it?

- Essentially neural networks that use convolution in place of general matrix multiplication in at least one of their layers.
 - Locally Receptive Fields
 - Shared Weights
 - Spatial or Temporal Sub-sampling

Convoluted Neural Network:Typical Architecture



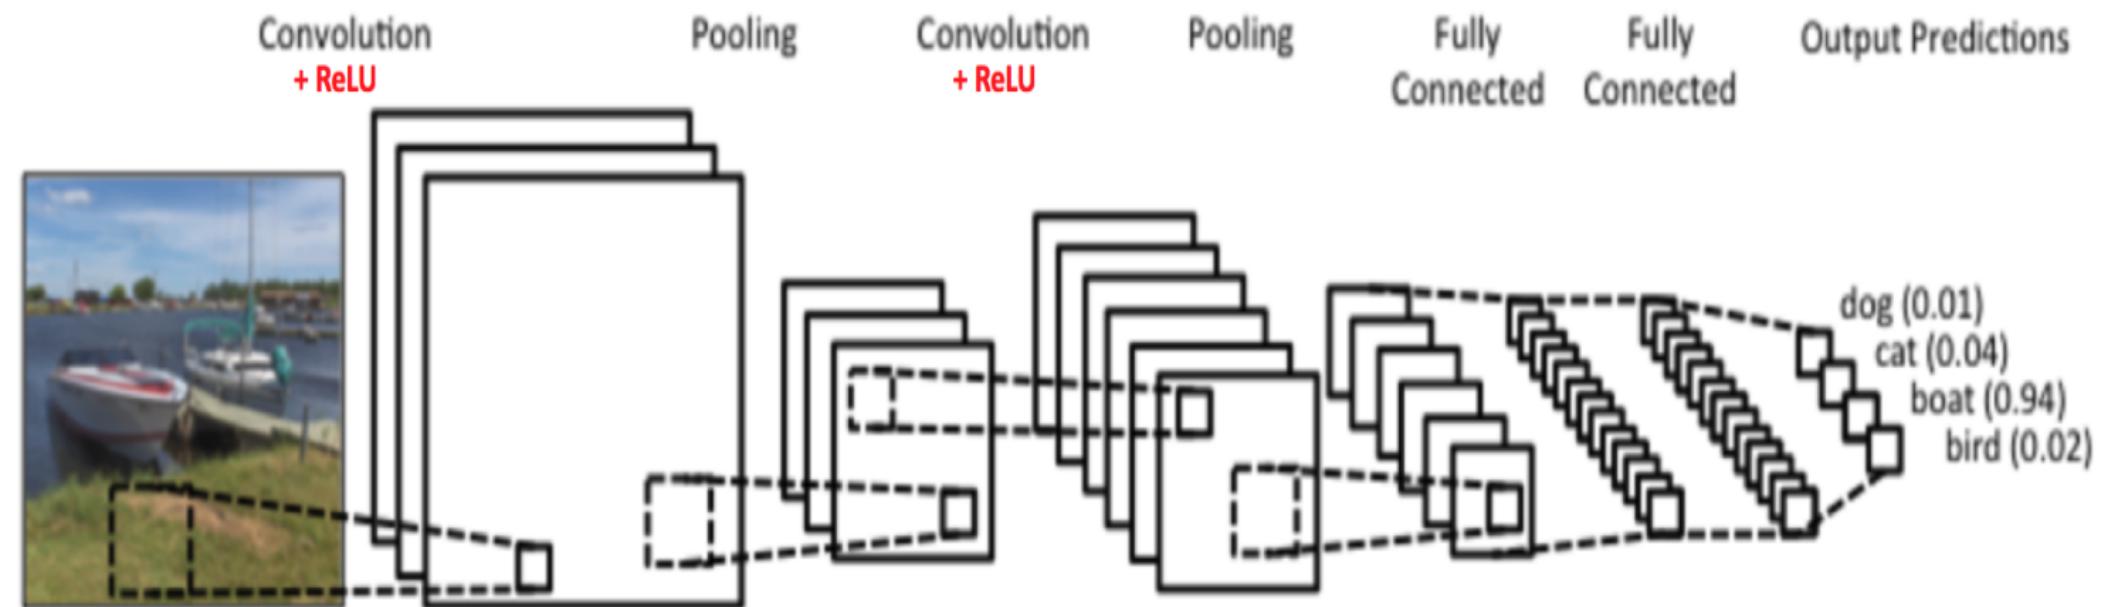
Convoluted Neural Network:Typical Architecture



Convulated Neural Network:Typical Architecture

- LeNet was one of the very first convolutional neural networks which helped propel the field of Deep Learning.
- This pioneering work by Yann LeCun was named **LeNet5** after many previous successful iterations since the year 1988 [3].
- At that time the LeNet architecture was used mainly for character recognition tasks such as reading zip codes, digits, etc.

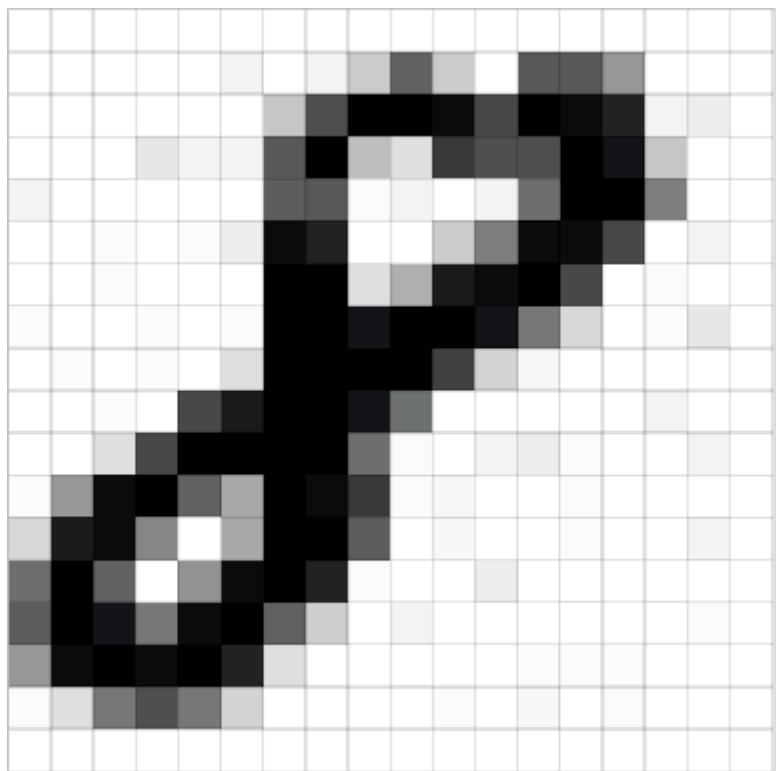
Convulated Neural Network:Lenet



Convulated Neural Network:Typical Architecture

- There are four main operations in the ConvNet:
 - Convolution
 - Non Linearity (ReLU)
 - Pooling or Sub Sampling
 - Classification (Fully Connected Layer)

Convoluted Neural Network:Lenet



Convolved Neural Network:Lenet:Convolutions

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

- Consider a 5×5 image whose pixel values are only 1 or 0 as a simplification.
- Also consider 3×3 matrix, called a filter, or kernel, or feature detector.

1	0	1
0	1	0
1	0	1

Convolved Neural Network:Lenet:Convolutions

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	

Convolved Feature

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved Feature

- Called a stride, 1 pixel in this case.

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved Feature

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved Feature

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved Feature

- Also called Activation map or feature map

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved Feature

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved Feature

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved Feature

- Different filters obviously produce different features.

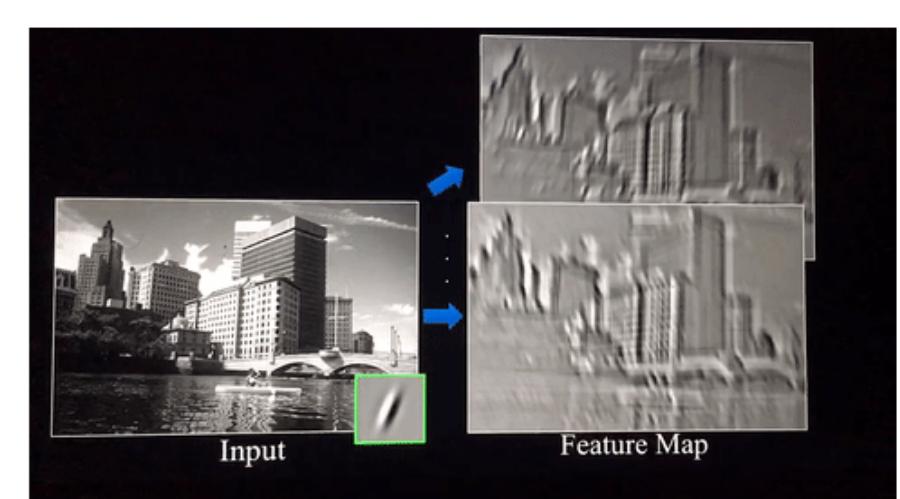
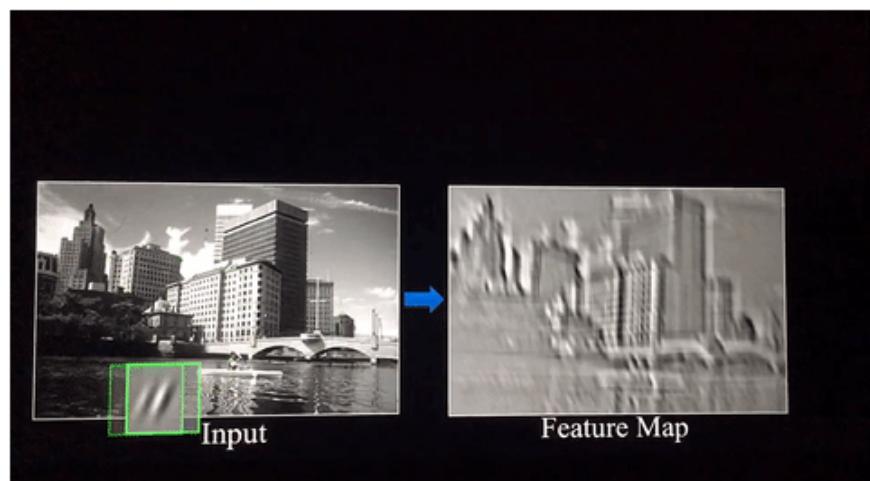
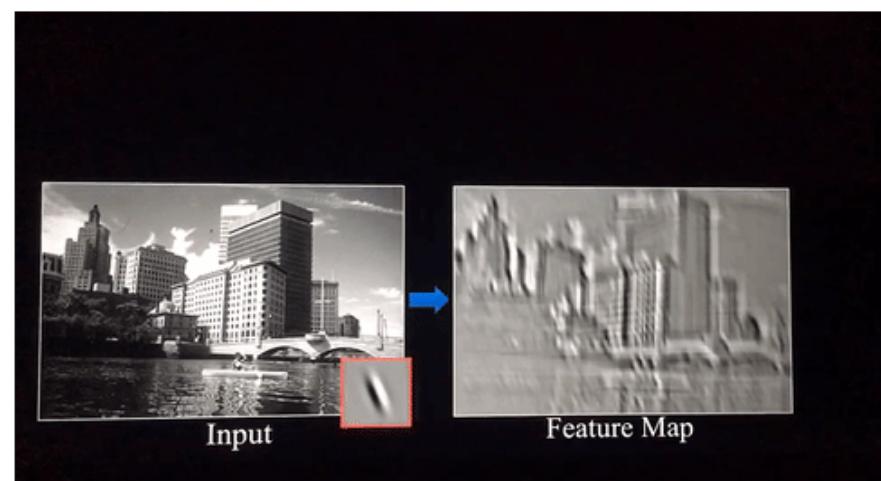
Convolved Neural Network:Lenet:Convolutions



original image

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	The original image of the deer head.
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	A black and white image where edges of the deer's features are highlighted.
Sharpen	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	The original image with some high-frequency noise added.
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	The original image with a smooth, uniform blur applied.
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	The original image with a smooth, non-uniform blur applied, simulating a Gaussian distribution.

Convolved Neural Network:Lenet:Convolutions



• More filters we have, more feature maps we have.

Convulated Neural Network:Lenet:Convolutions

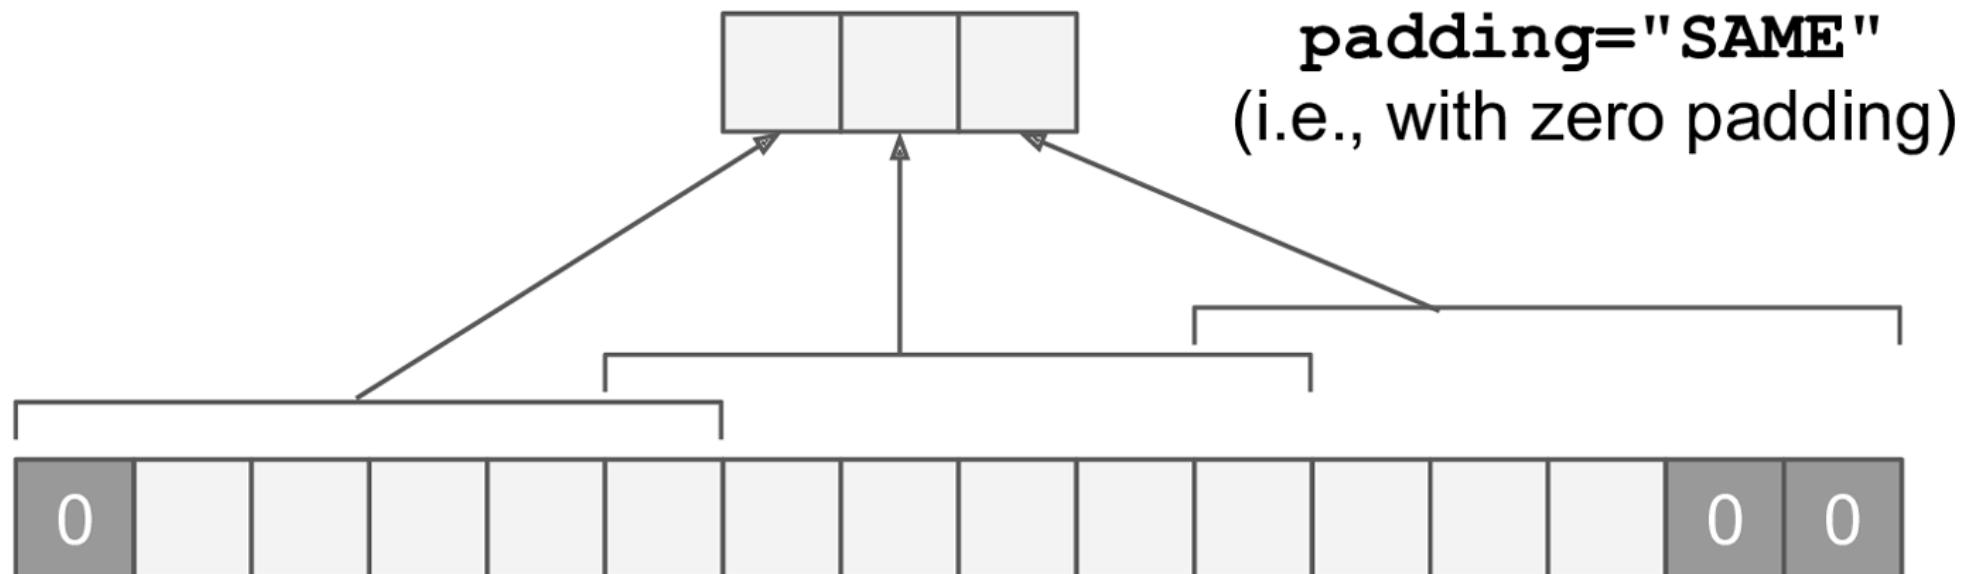
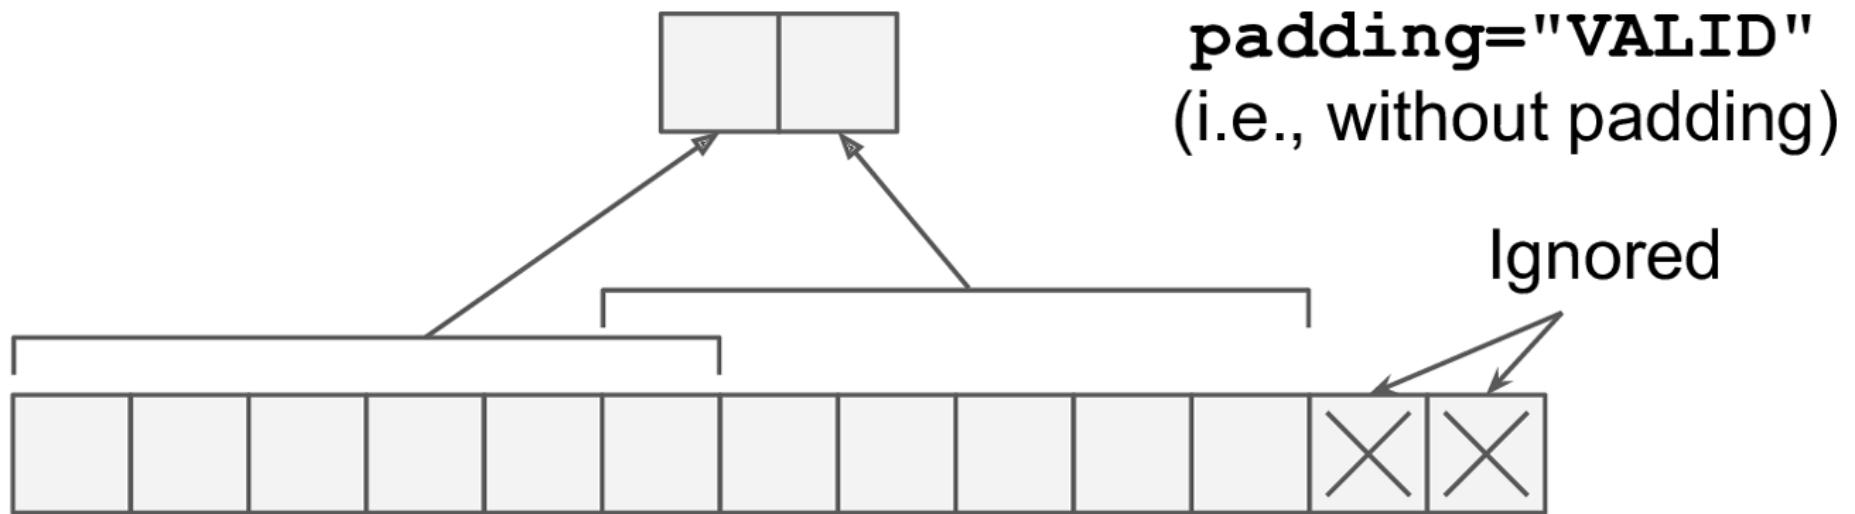


Feature Map having depth of 3 (since 3 filters have been used)

Convolution Operation

- Depth corresponds to the number of features.

Convolved Neural Network:Lenet:Convolutions



Convolved Neural Network:Lenet:Convolutions:Relu

Input Feature Map



Black = negative; white = positive values

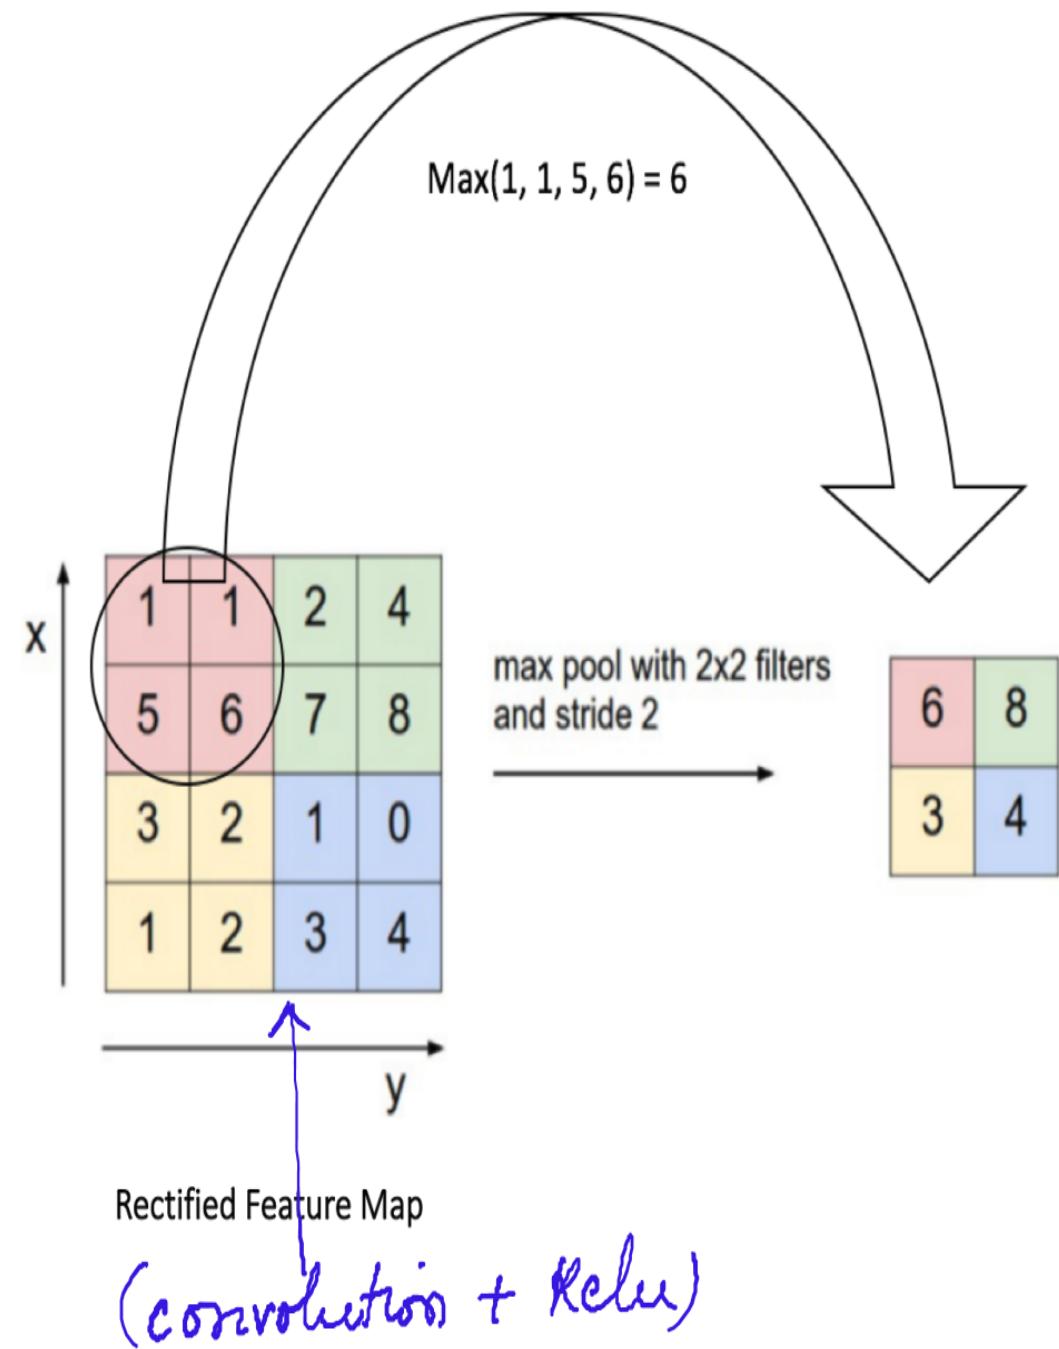
Rectified Feature Map



ReLU
→

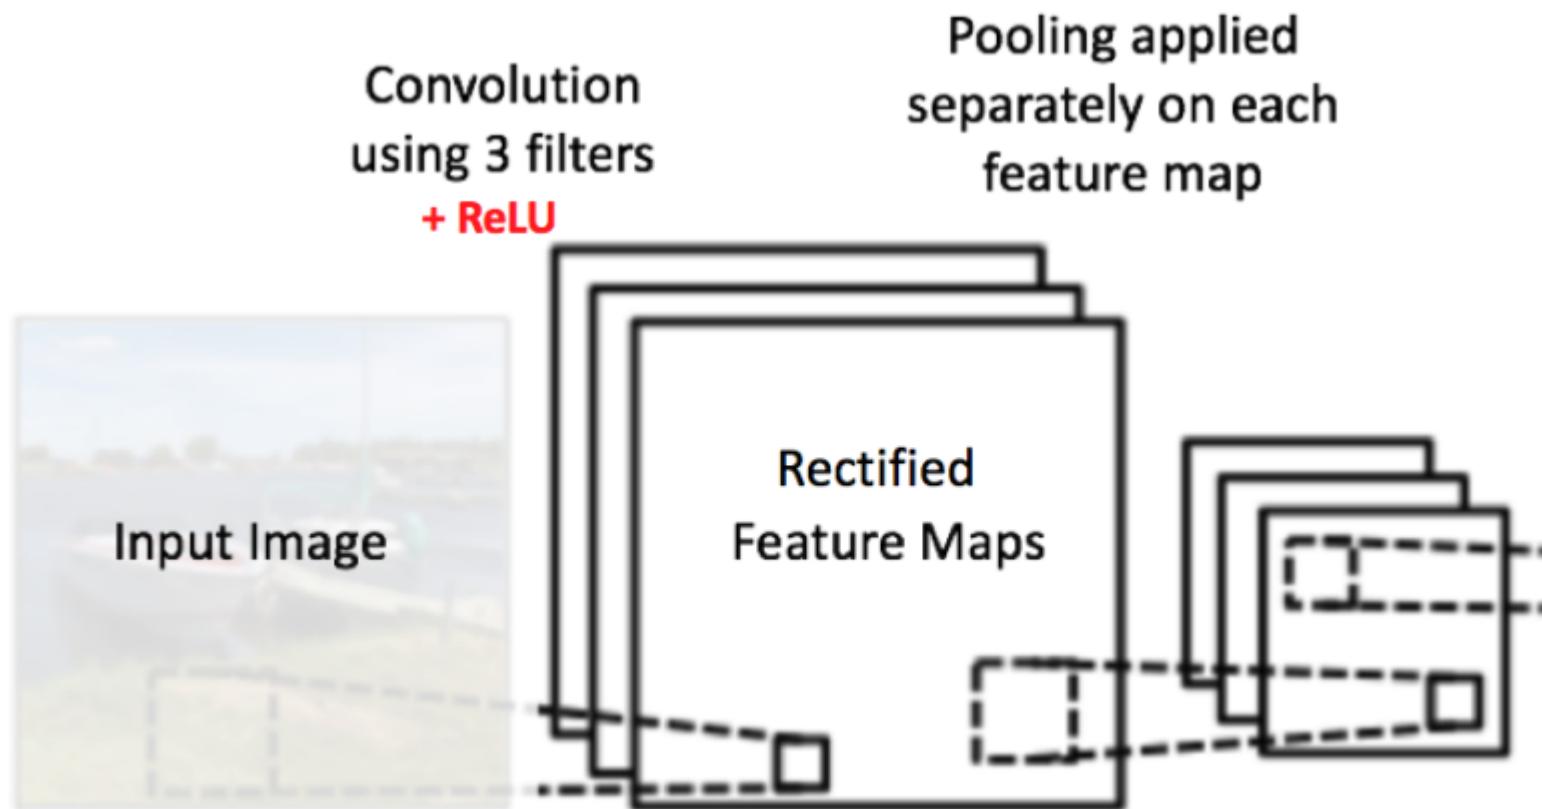
Only non-negative values

Convolved Neural Network:Lenet:Pooling



- Spatial pooling also called subsampling or down sampling.
- Different kinds of pooling can be done, Max, Average, sum.
- In practice max pooling works better.

Convulated Neural Network:Lenet:Pooling

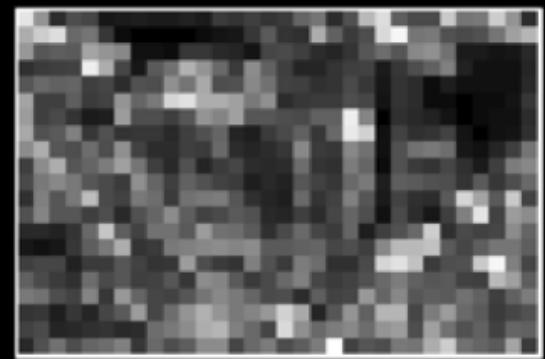


Convulated Neural Network:Lenet:Pooling

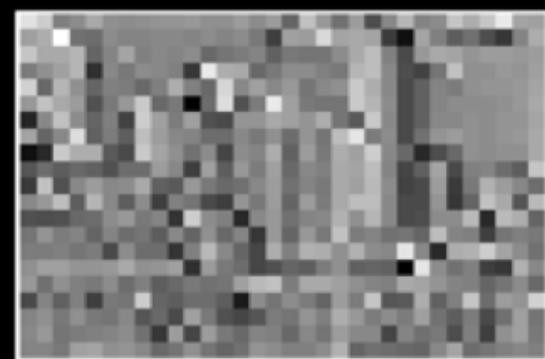


Pooling

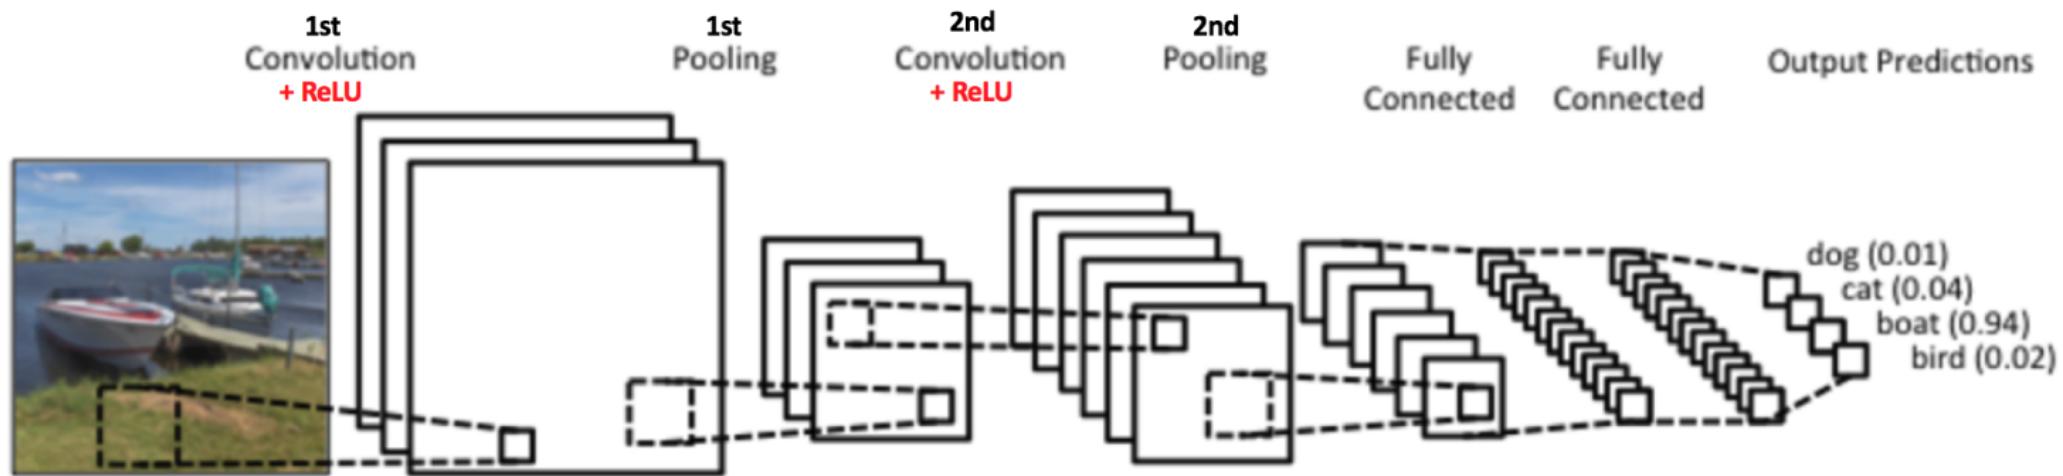
Max



Sum

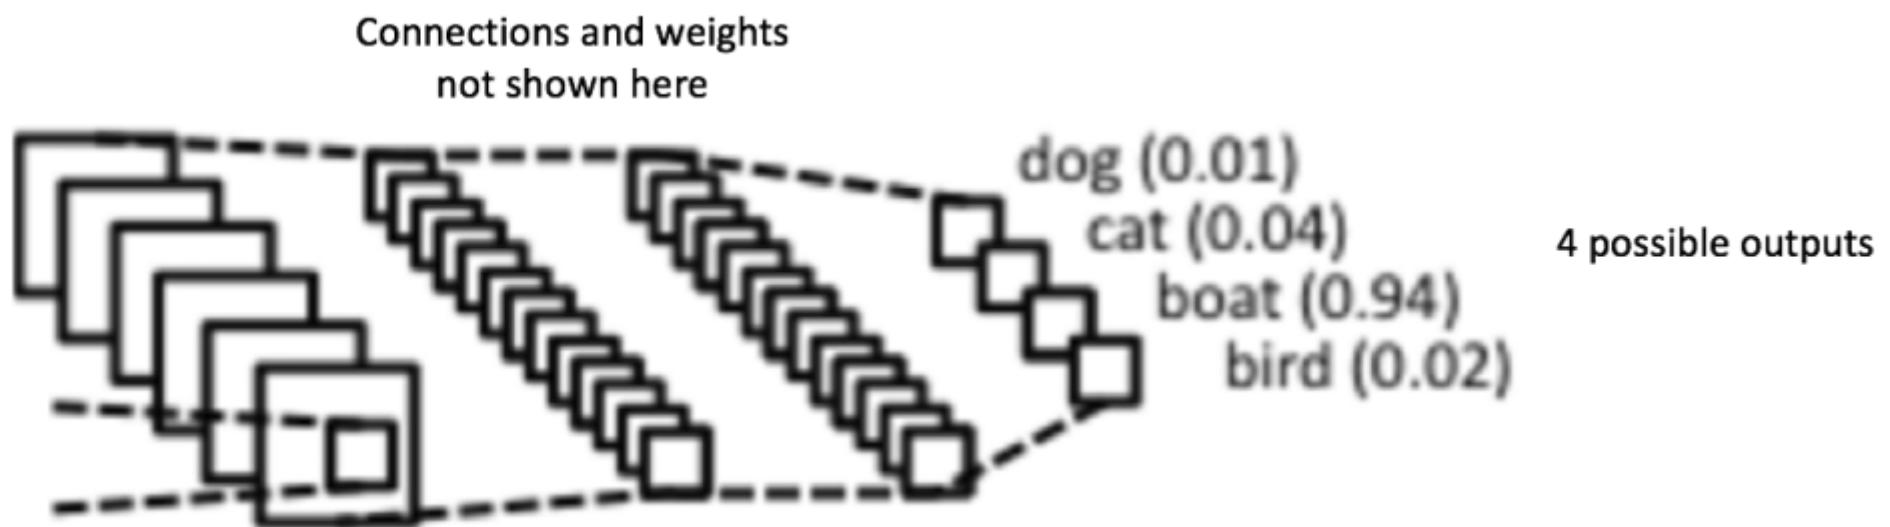


Convulated Neural Network:Lenet

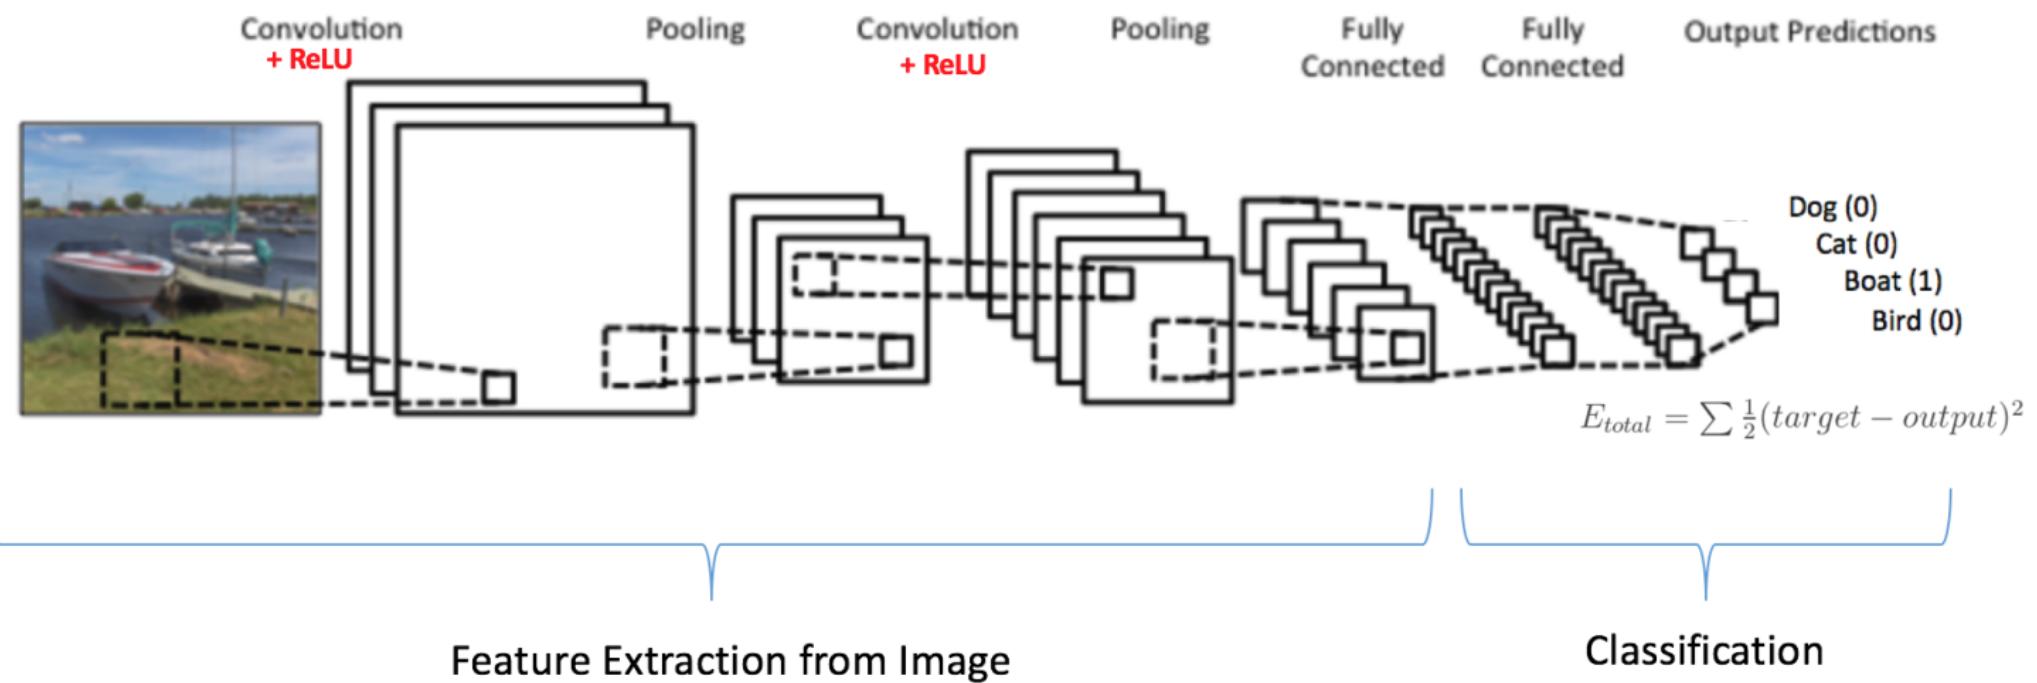


- 1st Convolution creates 3 feature maps.
- pooling
- 2nd Convolution creates 6 feature map
- pooling
- Finally FC layer and a softmax for classification.

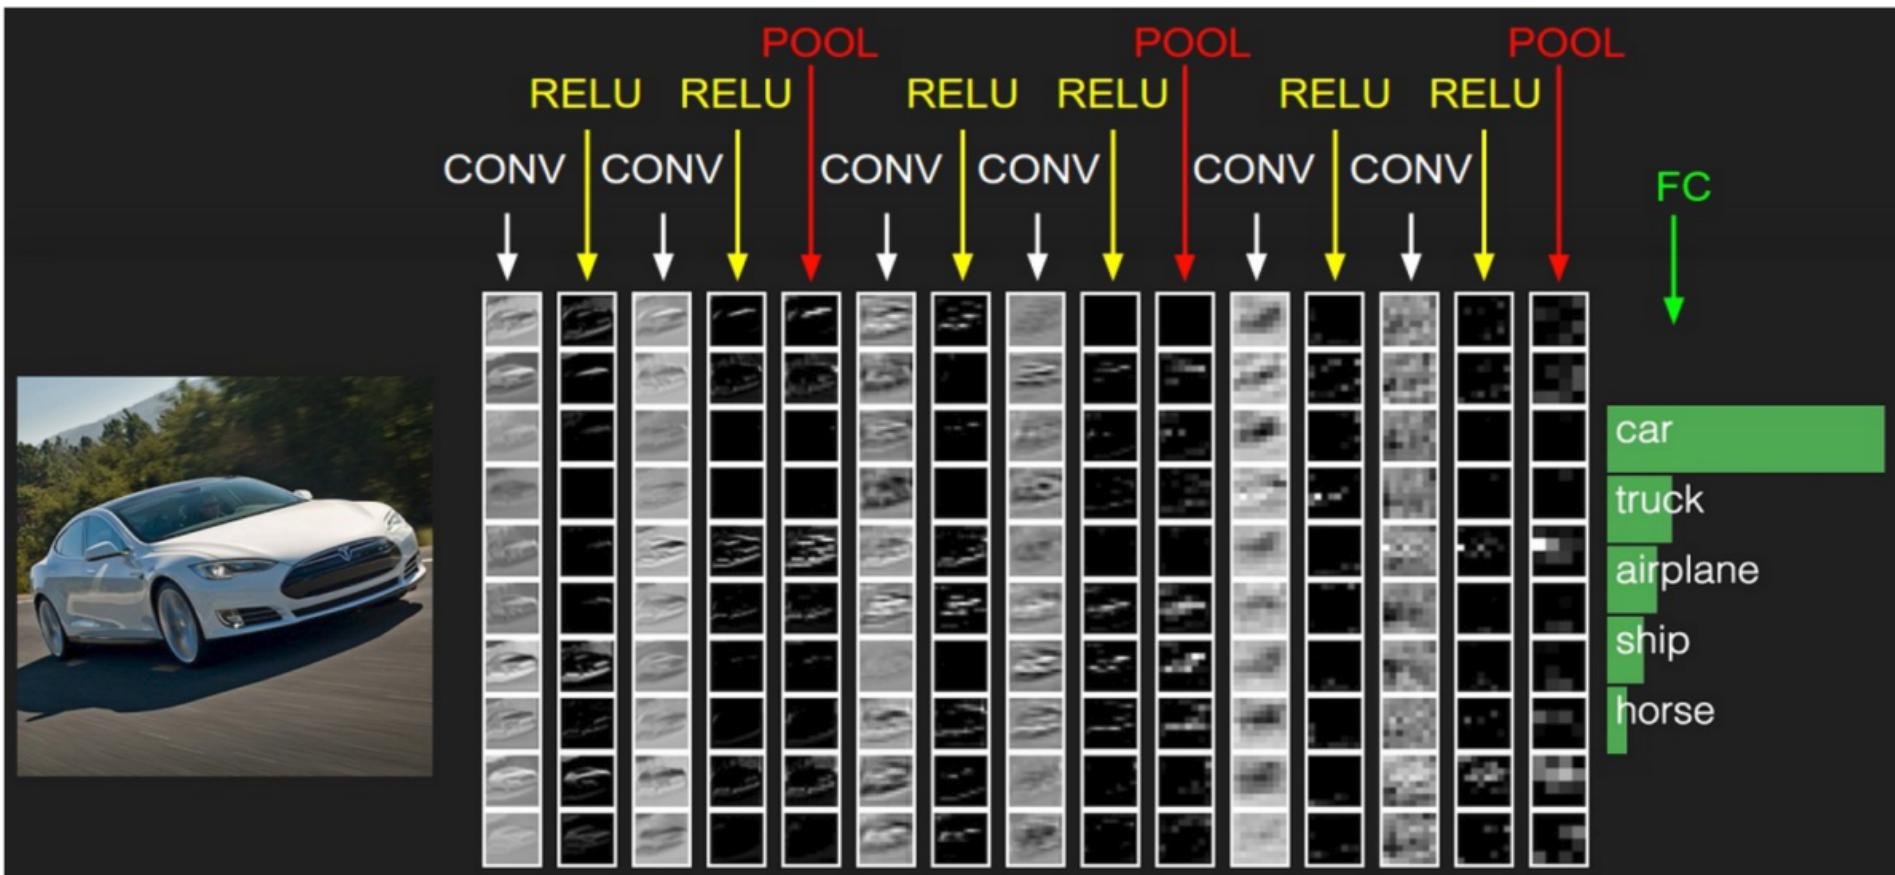
Convulated Neural Network:Lenet:Fully Connected



Convulated Neural Network:Lenet

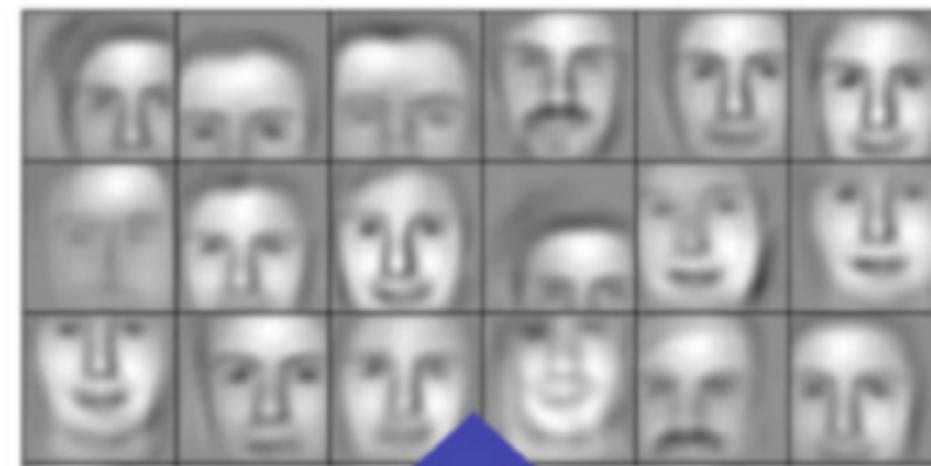


Convulated Neural Network: Visualization

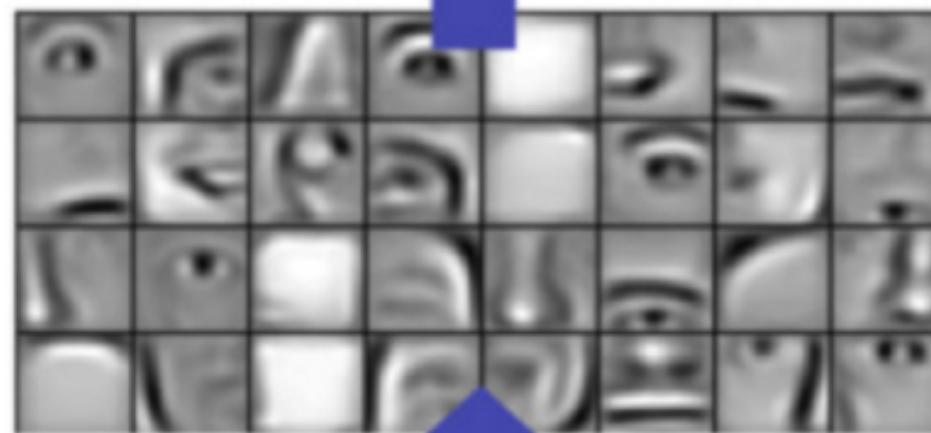


- These operations can be repeated any number of times
- Best performing CNNs have 10s of convolution and pooling layer.

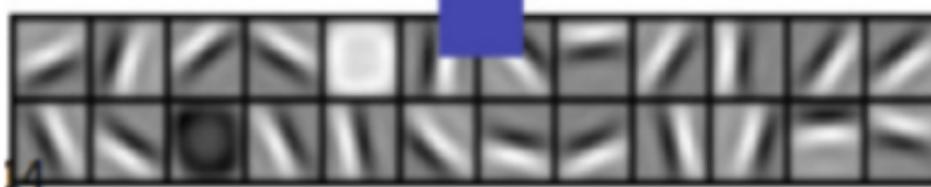
Convulated Neural Network: Visualization



Layer 3

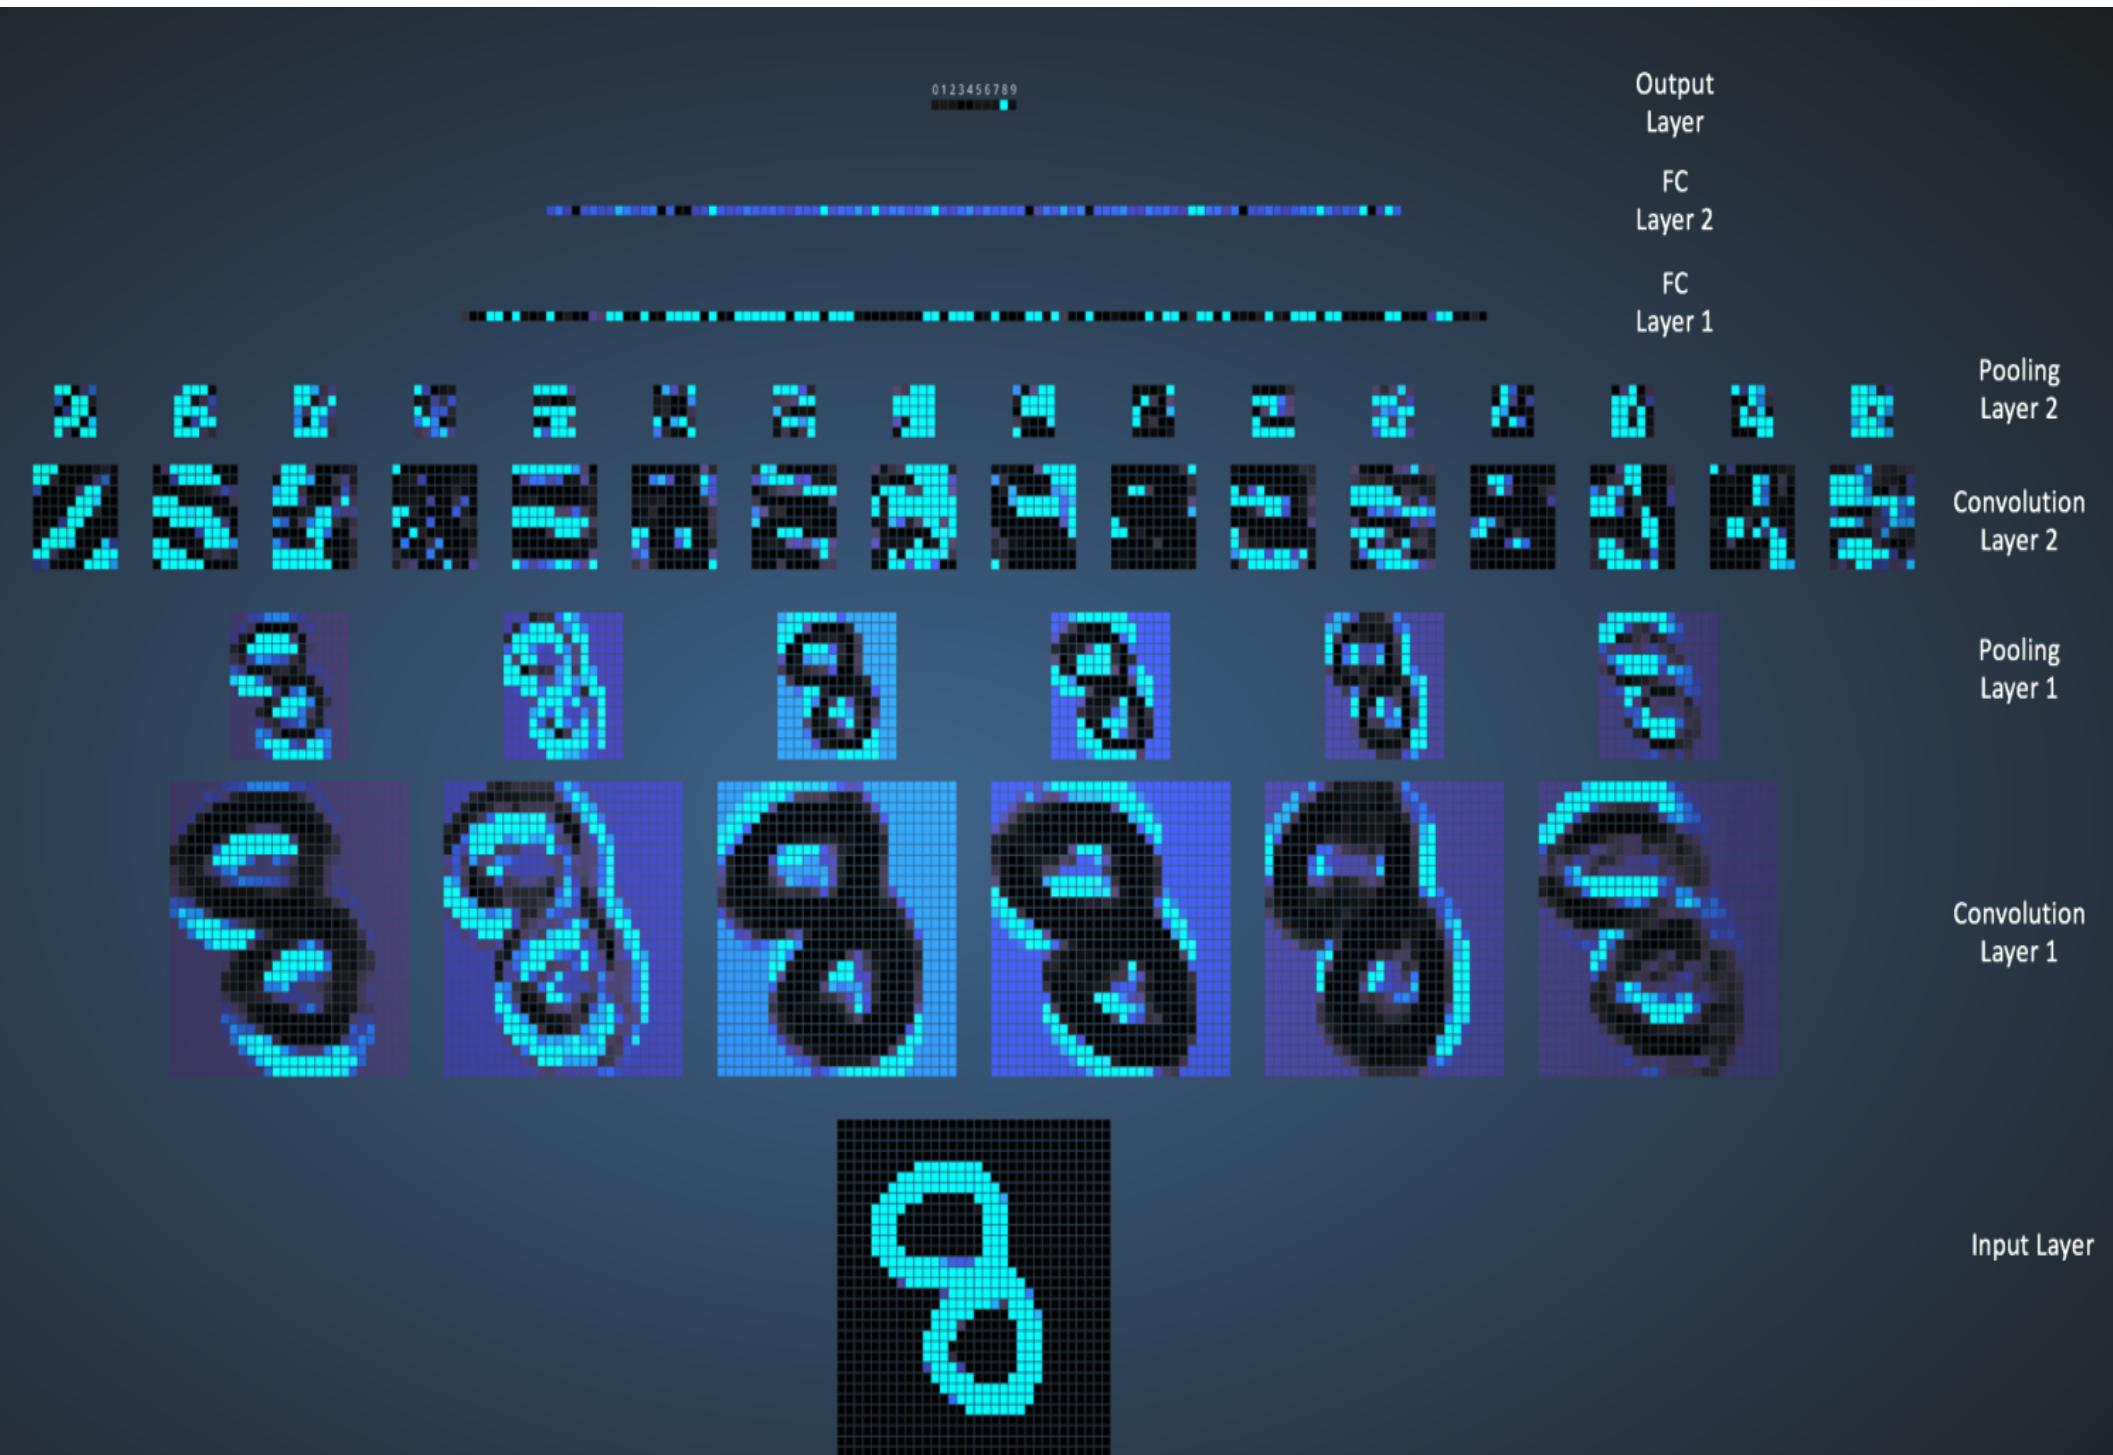


Layer 2

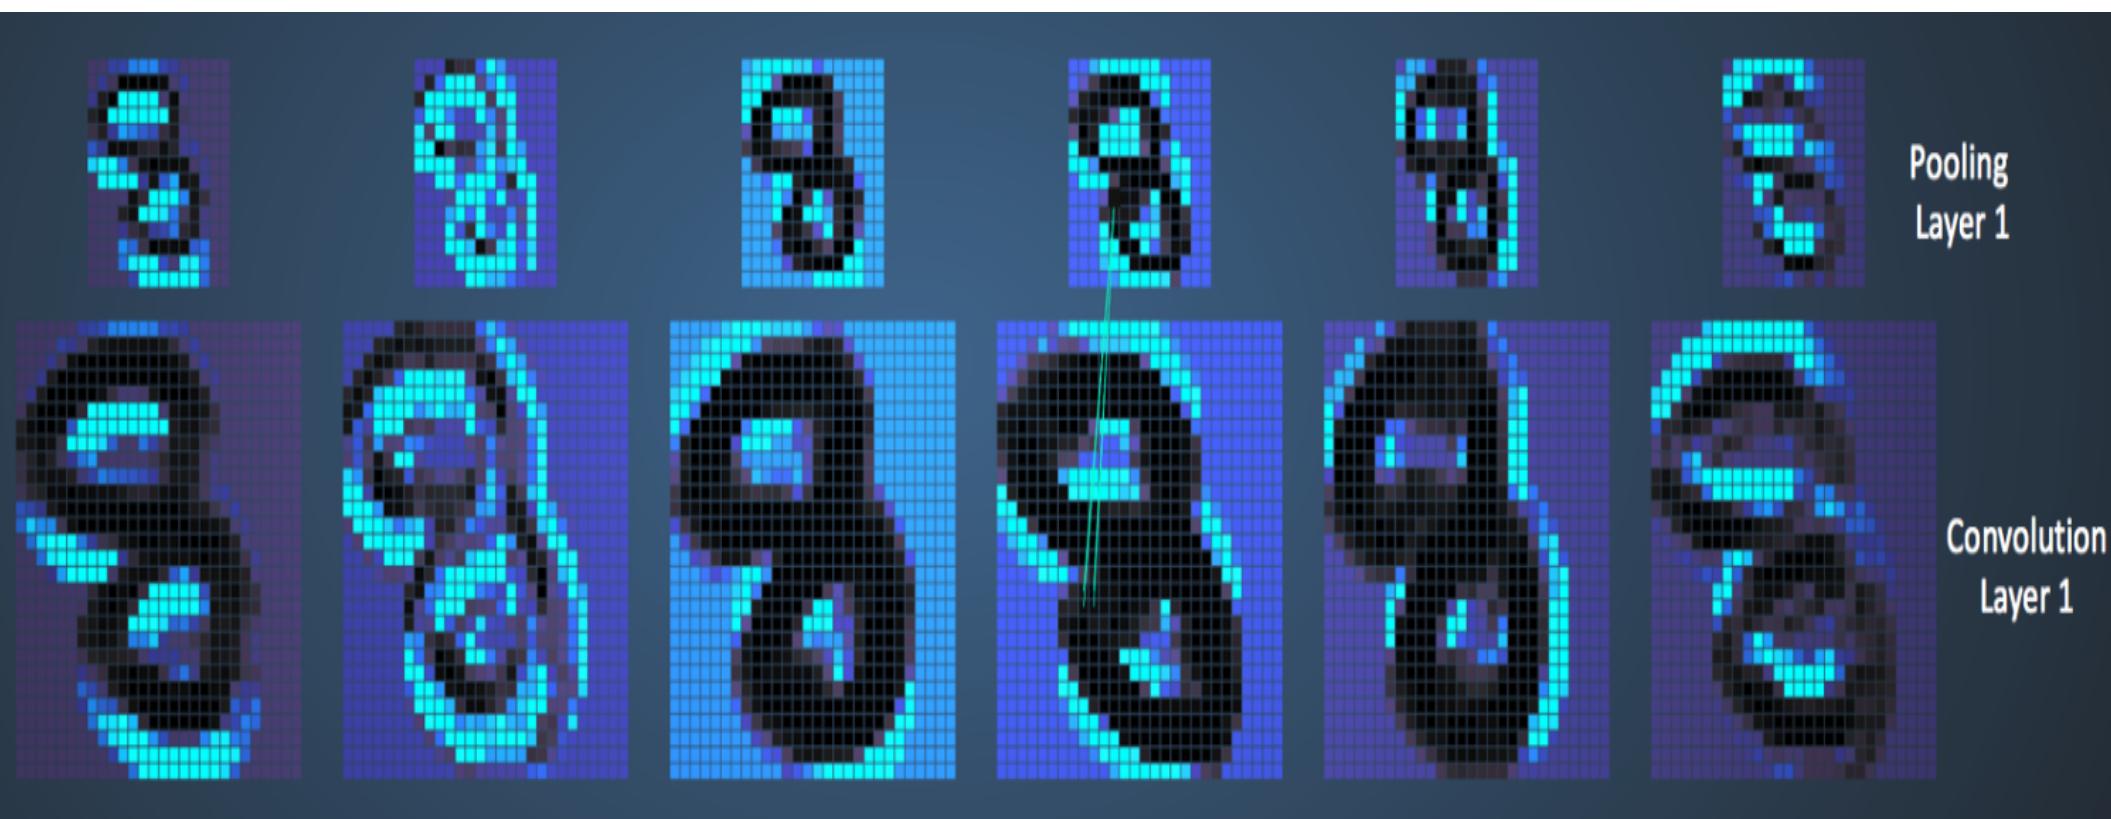


Layer 1

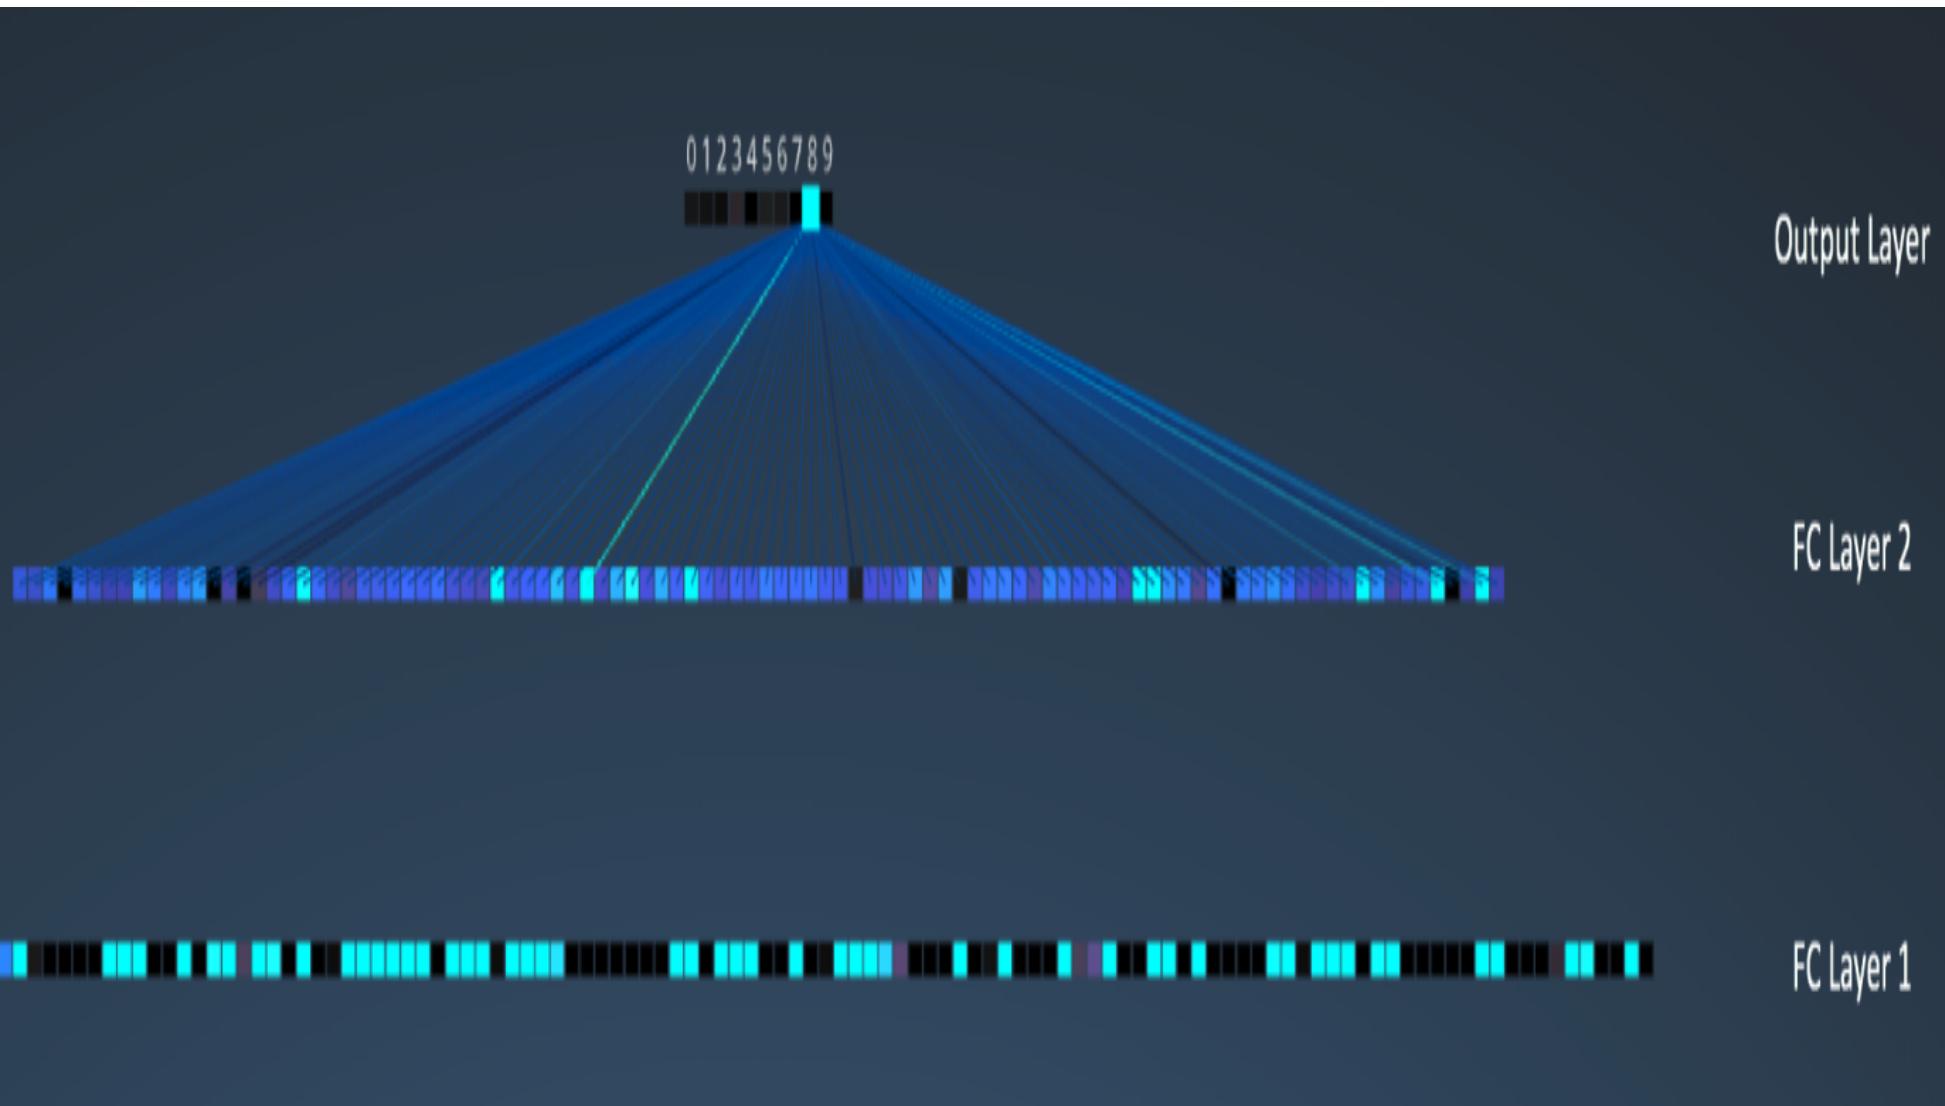
Convolved Neural Network: Visualization



Convulated Neural Network: Visualization

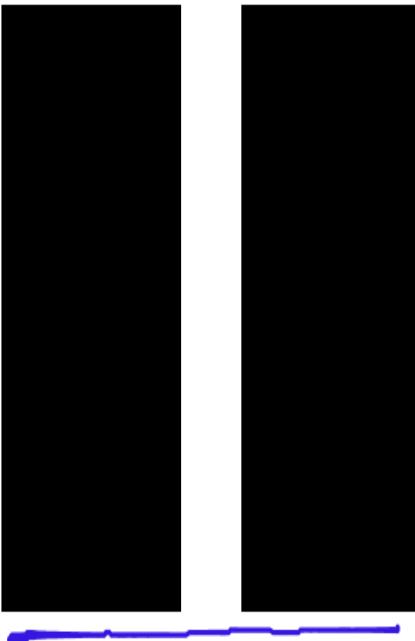


Convulated Neural Network: Visualization

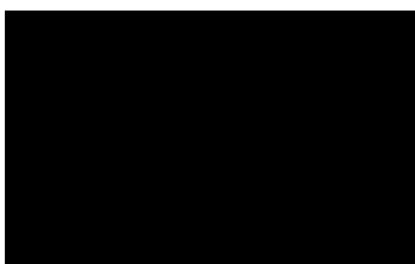


Convolved Neural Network: Filter

```
from sklearn.datasets import load_sample_image  
  
def plot_image(image):  
    plt.imshow(image, cmap="gray", interpolation="nearest")  
    plt.axis("off")  
  
def plot_color_image(image):  
    plt.imshow(image.astype(np.uint8), interpolation="nearest")  
    plt.axis("off")  
  
fmap = np.zeros(shape=(7, 7, 1, 2), dtype=np.float32)  
fmap[:, 3, 0, 0] = 1  
fmap[3, :, 0, 1] = 1  
fmap[:, :, 0, 0]  
plot image(fmap[:, :, 0, 0])  
plt.show()  
plot image(fmap[:, :, 0, 1])  
plt.show()
```



- Filter to detect vertical edges



- Filter to detect horizontal edges.



Convolved Neural Network: Convolution

```
china = load_sample_image("china.jpg")
flower = load_sample_image("flower.jpg")
image = china[150:220, 130:250]
height, width, channels = image.shape
image_grayscale = image.mean(axis=2).astype(np.float32)
images = image_grayscale.reshape(1, height, width, 1)

fmap = np.zeros(shape=(7, 7, 1, 2), dtype=np.float32)
fmap[:, 3, 0, 0] = 1
fmap[3, :, 0, 1] = 1
fmap[:, :, 0, 0]

X = tf.placeholder(tf.float32, shape=(None, height, width, 1))
feature_maps = tf.constant(fmap)
convolution = tf.nn.conv2d(X, feature_maps, strides=[1,1,1,1], padding="SAME", use_cudnn_on_gpu=False)

with tf.Session() as sess:
    output = convolution.eval(feed_dict={X: images})

plot_image(images[0, :, :, 0])
fl.savefig("china_original", tight_layout=False)
plt.show()

plot_image(output[0, :, :, 0])
fl.savefig("china_vertical", tight_layout=False)
plt.show()

plot_image(output[0, :, :, 1])
fl.savefig("china_horizontal", tight_layout=False)
plt.show()
```

- Convolutional layer
- Filters as params.



Convolved Neural Network:Convolution

```
from sklearn.datasets import load_sample_image

def plot_image(image):
    plt.imshow(image, cmap="gray", interpolation="nearest")
    plt.axis("off")

def plot_color_image(image):
    plt.imshow(image.astype(np.uint8), interpolation="nearest")
    plt.axis("off")

china = load_sample_image("china.jpg")
flower = load_sample_image("flower.jpg")
dataset = np.array([china, flower], dtype=np.float32)
batch_size, height, width, channels = dataset.shape

filters = np.zeros(shape=(7, 7, channels, 2), dtype=np.float32)
filters[:, 3, :, 0] = 1 # vertical line
filters[3, :, :, 1] = 1 # horizontal line

X = tf.placeholder(tf.float32, shape=(None, height, width, channels))
convolution = tf.nn.conv2d(X, filters, strides=[1,2,2,1], padding="SAME")

with tf.Session() as sess:
    output = sess.run(convolution, feed_dict={X: dataset})

for image_index in (0, 1):
    for feature_map_index in (0, 1):#every image second feature map
        plot_image(output[image_index, :, :, feature_map_index])
        plt.show()
```

Convolved Neural Network: Convolution: Padding

```
def plot_image(image):
    plt.imshow(image, cmap="gray", interpolation="nearest")
    plt.axis("off")

def plot_color_image(image):
    plt.imshow(image.astype(np.uint8), interpolation="nearest")
    plt.axis("off")

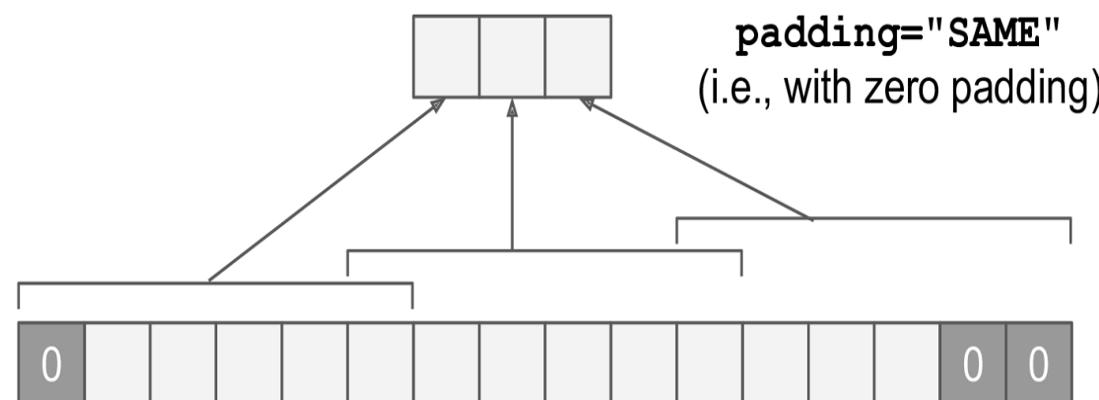
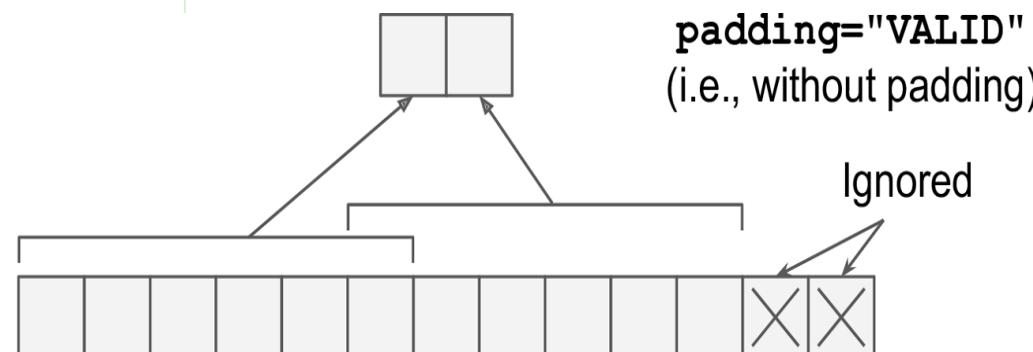
tf.reset_default_graph()

filter_primes = np.array([2., 3., 5., 7., 11., 13.], dtype=np.float32)
x = tf.constant(np.arange(1, 13+1, dtype=np.float32).reshape([1, 1, 13, 1]))
filters = tf.constant(filter_primes.reshape(1, 6, 1, 1))

valid_conv = tf.nn.conv2d(x, filters, strides=[1, 1, 5, 1], padding='VALID')
same_conv = tf.nn.conv2d(x, filters, strides=[1, 1, 5, 1], padding='SAME')

with tf.Session() as sess:
    print("VALID:\n", valid_conv.eval())
    print("SAME:\n", same_conv.eval())

print(filter_primes)
print("VALID:")
print(np.array([1,2,3,4,5,6]).T.dot(filter_primes))
print(np.array([6,7,8,9,10,11]).T.dot(filter_primes))
print("SAME:")
print(np.array([0,1,2,3,4,5]).T.dot(filter_primes))
print(np.array([5,6,7,8,9,10]).T.dot(filter_primes))
print(np.array([10,11,12,13,0,0]).T.dot(filter_primes))
```



Convolved Neural Network: Convolution: Padding

```
def plot_image(image):
    plt.imshow(image, cmap="gray", interpolation="nearest")
    plt.axis("off")

def plot_color_image(image):
    plt.imshow(image.astype(np.uint8), interpolation="nearest")
    plt.axis("off")

tf.reset_default_graph()

china = load_sample_image("china.jpg")
flower = load_sample_image("flower.jpg")
dataset = np.array([china, flower], dtype=np.float32)
batch_size, height, width, channels = dataset.shape

batch_size, height, width, channels = dataset.shape

filters = np.zeros(shape=(7, 7, channels, 2), dtype=np.float32)
filters[:, 3, :, 0] = 1 # vertical line
filters[3, :, :, 1] = 1 # horizontal line

X = tf.placeholder(tf.float32, shape=(None, height, width, channels))
max_pool = tf.nn.max_pool(X, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="VALID")

with tf.Session() as sess:
    output = sess.run(max_pool, feed_dict={X: dataset})

plot_color_image(dataset[0])
fl.save_fig("china_original")
plt.show()

plot_color_image(output[0])
fl.save_fig("china_max_pool")
plt.show()
```



original



pooled.

Convolved Neural Network:Tensorflow:CNN-1

```
graph = tf.Graph()
with graph.as_default():
    with tf.name_scope("inputs"):
        X = tf.placeholder(tf.float32, shape=[None, n_inputs], name="X")
        X_reshaped = tf.reshape(X, shape=[-1, height, width, channels])
        y = tf.placeholder(tf.int32, shape=[None], name="y")
        is_training = tf.placeholder_with_default(False, shape=[], name='is_training')
    conv1 = tf.layers.conv2d(X_reshaped, filters=conv1_fmaps, kernel_size=conv1_ksize, strides=conv1_stride,
                           padding=conv1_pad, activation=tf.nn.relu, name="conv1")
    conv2 = tf.layers.conv2d(conv1, filters=conv2_fmaps, kernel_size=conv2_ksize, strides=conv2_stride,
                           padding=conv2_pad, activation=tf.nn.relu, name="conv2")
    with tf.name_scope("pool3"):
        pool3 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="VALID")
        pool3_flat = tf.reshape(pool3, shape=[-1, pool3_fmaps * 14 * 14])
        pool3_flat_drop = tf.layers.dropout(pool3_flat, conv2_dropout_rate, training=is_training)
    with tf.name_scope("fc1"):
        fc1 = tf.layers.dense(pool3_flat_drop, n_fc1, activation=tf.nn.relu, name="fc1")
        fc1_drop = tf.layers.dropout(fc1, fc1_dropout_rate, training=is_training)

    with tf.name_scope("output"):
        logits = tf.layers.dense(fc1, n_outputs, name="output")
        Y_proba = tf.nn.softmax(logits, name="Y_proba")

    with tf.name_scope("train"):
        xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=logits, labels=y)
        loss = tf.reduce_mean(xentropy)
        optimizer = tf.train.AdamOptimizer()
        training_op = optimizer.minimize(loss)

    with tf.name_scope("eval"):
        correct = tf.nn.in_top_k(logits, y, 1)
        accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))

    with tf.name_scope("init_and_save"):
        init = tf.global_variables_initializer()
        saver = tf.train.Saver()

mse_summary = tf.summary.scalar('ACCURACY', accuracy)
```

Convolved Neural Network:Tensorflow:CNN-2

```
mse_summary = tf.summary.scalar('ACCURACY', accuracy)
summary_writer = tf.summary.FileWriter(logdir, tf.get_default_graph())

from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("/tmp/data/")

n_epochs = 25
batch_size = 50

best_acc_val = 0
check_interval = 100
checks_since_last_progress = 0
max_checks_without_progress = 100
best_model_params = None

with tf.Session(graph=graph) as sess:
    init.run()
    for epoch in range(n_epochs):
        for iteration in range(mnist.train.num_examples // batch_size):
            X_batch, y_batch = mnist.train.next_batch(batch_size)
            sess.run(training_op, feed_dict={X: X_batch, y: y_batch, is_training: True})
            if iteration % check_interval == 0:
                acc_val = accuracy.eval(feed_dict={X: mnist.test.images[:2000], y: mnist.test.labels[:2000]})
                if acc_val > best_acc_val:
                    best_acc_val = acc_val
                    checks_since_last_progress = 0
                    best_model_params = get_model_params()
                else:
                    checks_since_last_progress += 1
            acc_train = accuracy.eval(feed_dict={X: X_batch, y: y_batch})
            acc_test = accuracy.eval(feed_dict={X: mnist.test.images[2000:], y: mnist.test.labels[2000:]})
            print(epoch, "Train accuracy:", acc_train, "Test accuracy:", acc_test, "Best validation accuracy:", best_acc_val)
        if checks_since_last_progress > max_checks_without_progress:
            print("Early stopping!")
            break

    if best_model_params:
        restore_model_params(best_model_params)
```

Convolved Neural Network:Tensorflow:CNN-3

```
(ml_home) mohit@nomind:~/Work/ArtificialIntelligence$ ./DL/CNN/cnnmnist.py
Extracting /tmp/data/train-images-idx3-ubyte.gz
Extracting /tmp/data/train-labels-idx1-ubyte.gz
Extracting /tmp/data/t10k-images-idx3-ubyte.gz
Extracting /tmp/data/t10k-labels-idx1-ubyte.gz
2017-07-02 06:11:10.925714: W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow
p CPU computations.
2017-07-02 06:11:10.925742: W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow
p CPU computations.
2017-07-02 06:11:10.925757: W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow
PU computations.
2017-07-02 06:11:10.925763: W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow
CPU computations.
2017-07-02 06:11:10.925778: W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow
PU computations.
0 Train accuracy: 0.98 Test accuracy: 0.9875 Best validation accuracy: 0.978
1 Train accuracy: 1.0 Test accuracy: 0.989 Best validation accuracy: 0.9835
2 Train accuracy: 1.0 Test accuracy: 0.98925 Best validation accuracy: 0.985
3 Train accuracy: 1.0 Test accuracy: 0.9925 Best validation accuracy: 0.987
4 Train accuracy: 1.0 Test accuracy: 0.991 Best validation accuracy: 0.987
5 Train accuracy: 1.0 Test accuracy: 0.992125 Best validation accuracy: 0.987
6 Train accuracy: 1.0 Test accuracy: 0.992 Best validation accuracy: 0.989
7 Train accuracy: 1.0 Test accuracy: 0.99 Best validation accuracy: 0.989
8 Train accuracy: 1.0 Test accuracy: 0.99225 Best validation accuracy: 0.989
9 Train accuracy: 1.0 Test accuracy: 0.990375 Best validation accuracy: 0.989
10 Train accuracy: 1.0 Test accuracy: 0.991375 Best validation accuracy: 0.989
11 Train accuracy: 1.0 Test accuracy: 0.991625 Best validation accuracy: 0.989
12 Train accuracy: 1.0 Test accuracy: 0.99025 Best validation accuracy: 0.989
13 Train accuracy: 1.0 Test accuracy: 0.99175 Best validation accuracy: 0.989
14 Train accuracy: 1.0 Test accuracy: 0.991375 Best validation accuracy: 0.989
15 Train accuracy: 1.0 Test accuracy: 0.99175 Best validation accuracy: 0.989
16 Train accuracy: 1.0 Test accuracy: 0.991875 Best validation accuracy: 0.989
Early stopping!
Final accuracy on test set: 0.991625
```

Convolved Neural Network:Tensorflow:Pretrained Models

```
import matplotlib.image as mpimg
test_image = mpimg.imread(os.path.join("input/images","cnn","test_image.png"))[:, :, :channels]
plt.imshow(test_image)
plt.axis("off")
plt.show()

from tensorflow.contrib.slim.nets import inception
import tensorflow.contrib.slim as slim

tf.reset_default_graph()

X = tf.placeholder(tf.float32, shape=[None, height, width, channels], name="X")
with slim.arg_scope(inception.inception_v3_arg_scope()):
    logits, end_points = inception.inception_v3(X, num_classes=1001, is_training=False)
predictions = end_points["Predictions"]
saver = tf.train.Saver()

X_test = test_image.reshape(-1, height, width, channels)

with tf.Session() as sess:
    saver.restore(sess, INCEPTION_V3_CHECKPOINT_PATH)
    predictions_val = predictions.eval(feed_dict={X: X_test})

print("Predicted Class:", class_names[np.argmax(predictions_val[0])])

print(np.argmax(predictions_val, axis=1))

top_5 = np.argpartition(predictions_val[0], -5)[-5:]
top_5 = top_5[np.argsort(predictions_val[0][top_5])]
for i in top_5:
    print("{0}: {1:.2f}%".format(class_names[i], 100*predictions_val[0][i]))
```



```
swing: 0.04%
beer bottle: 0.05%
common newt, Triturus vulgaris: 0.05%
red fox, Vulpes vulpes: 2.36%
hyena, hyaena: 93.85%
```