

Supervised:Logistic Regression

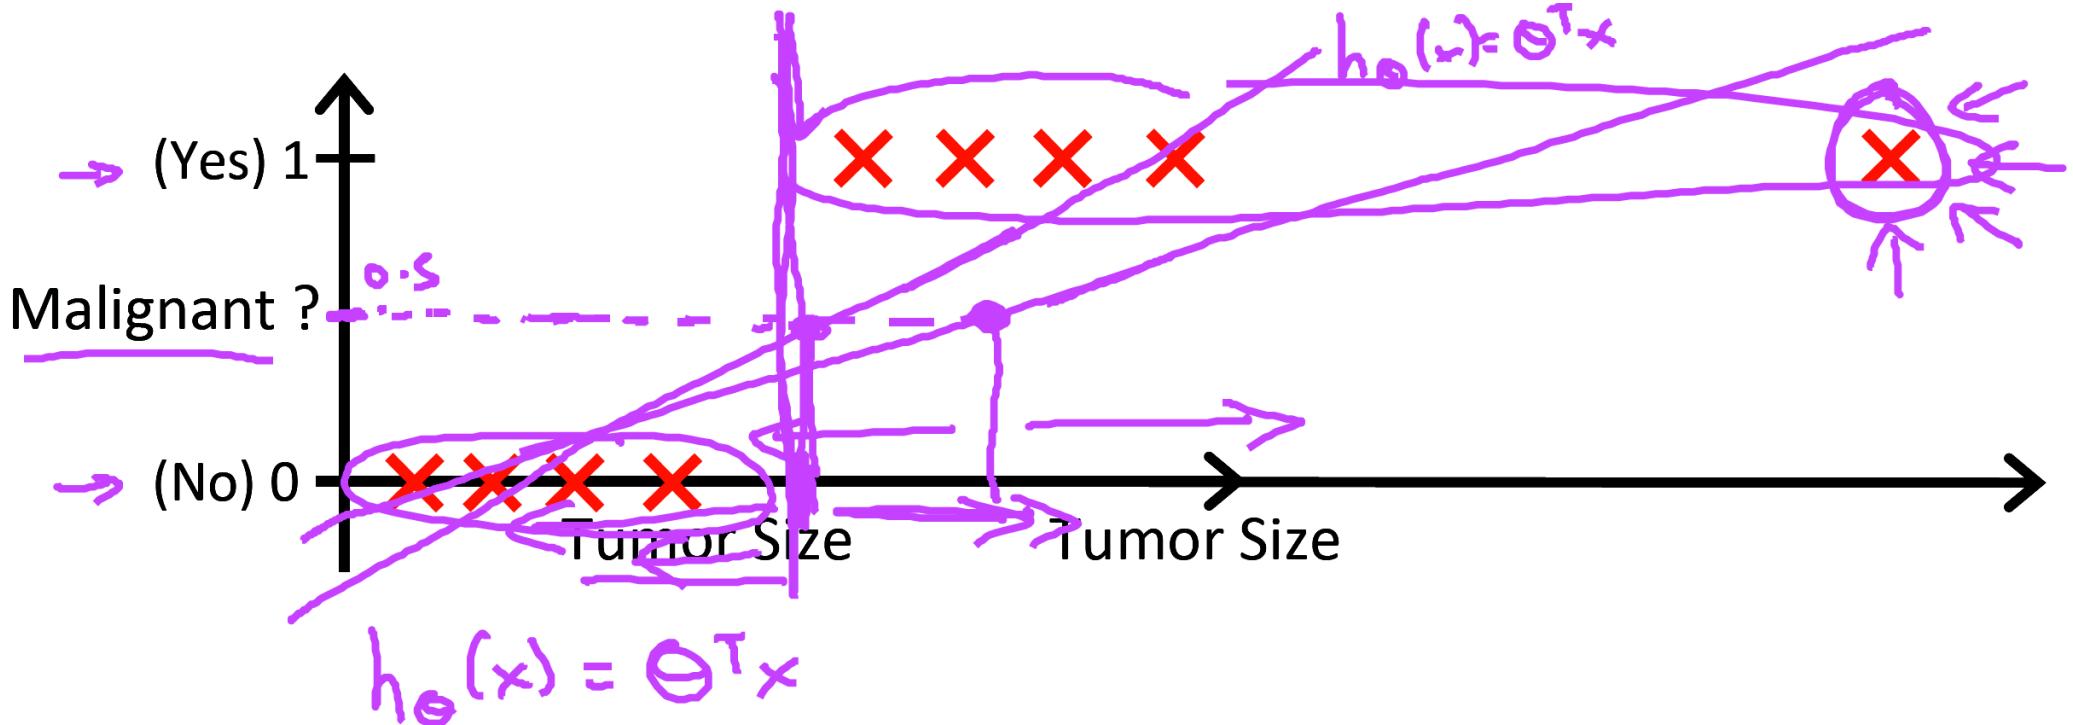
Classification

- Email: Spam / Not Spam?
- Online Transactions: Fraudulent (Yes / No)?
- Tumor: Malignant / Benign ?

- $y \in \{0, 1\}$
 - 0: “Negative Class” (e.g., benign tumor)
 - 1: “Positive Class” (e.g., malignant tumor)

$$\rightarrow y \in \{0, 1, 2, 3\}$$

Classification



→ Threshold classifier output $h_{\theta}(x)$ at 0.5:

→ If $h_{\theta}(x) \geq 0.5$, predict "y = 1"

If $h_{\theta}(x) < 0.5$, predict "y = 0"

Classification

Classification: $y = 0 \text{ or } 1$

$h_\theta(x)$ can be $\underline{> 1}$ or $\underline{< 0}$

Logistic Regression: $0 \leq h_\theta(x) \leq 1$



Classification

Classification

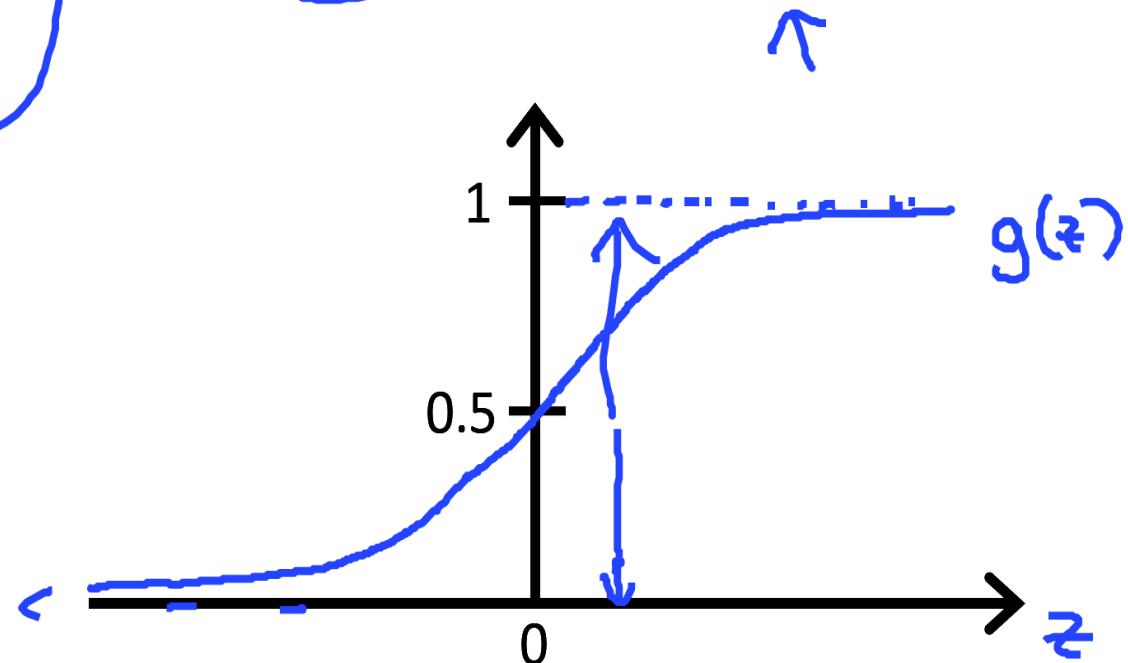
Logistic Regression Model

Want $0 \leq h_\theta(x) \leq 1$

$$h_\theta(x) = g(\theta^T x)$$

$$\rightarrow g(z) = \frac{1}{1+e^{-z}}$$

$$h_\theta(x) = \frac{1}{1+e^{-\theta^T x}}$$



↳ Sigmoid function

↳ Logistic function

Parameters θ .

Classification

Interpretation of Hypothesis Output

$h_{\theta}(x)$ = estimated probability that $y = 1$ on input x

Example: If $\underline{x} = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 & \leftarrow \\ \underline{\text{tumorSize}} & \leftarrow \end{bmatrix}$

$$\underline{h_{\theta}(x) = 0.7} \quad \underline{y=1}$$

Tell patient that 70% chance of tumor being malignant

$h_{\theta}(x) = P(y=1|x; \theta)$

“probability that $y = 1$, given x , parameterized by θ ”

$y = 0 \text{ or } 1$

$$\begin{aligned} \rightarrow P(y = 0|x; \theta) + P(y = 1|x; \theta) &= 1 \\ \rightarrow P(y = 0|x; \theta) &= 1 - P(y = 1|x; \theta) \end{aligned}$$

Classification

Logistic regression

$$h_{\theta}(x) = g(\theta^T x)$$

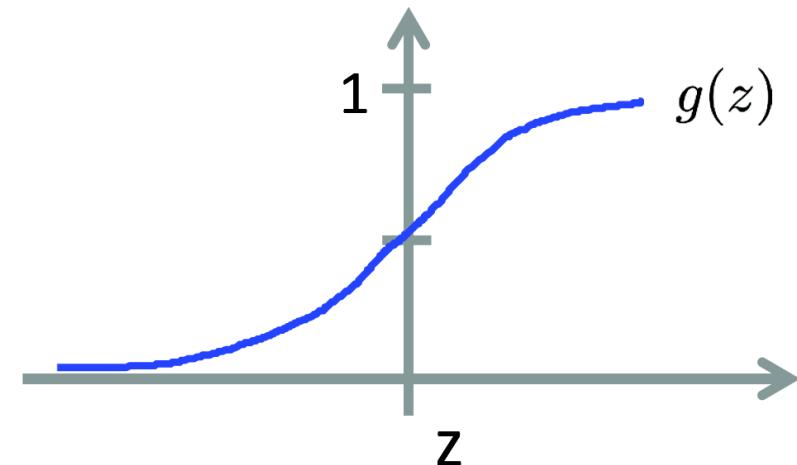
$$g(z) = \frac{1}{1+e^{-z}}$$

Suppose predict " $y = 1$ " if $h_{\theta}(x) \geq 0.5$

$$\theta^T x \geq 0$$

predict " $y = 0$ " if $h_{\theta}(x) < 0.5$

$$\theta^T x < 0$$



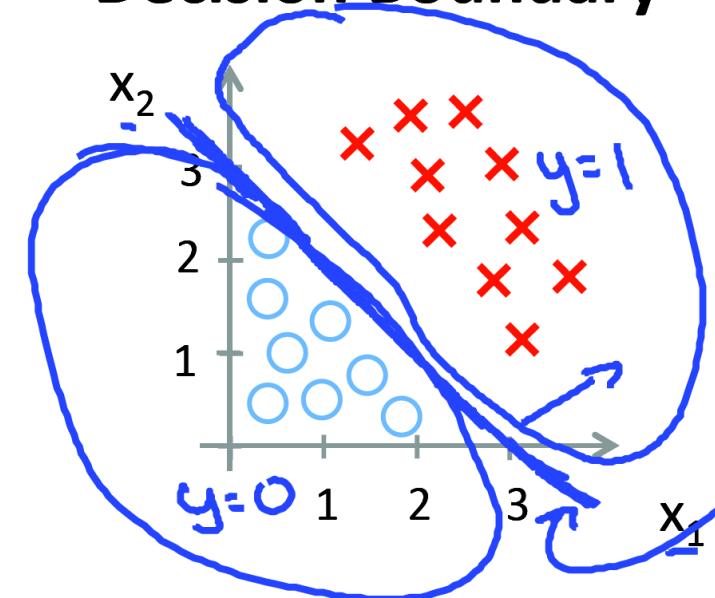
$$g(z) \geq 0.5 \\ \text{when } z \geq 0$$

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) < 0.5 \\ \text{when } z < 0$$

Classification

Decision Boundary



Predict " $y = 1$ " if

$$-3 + x_1 + x_2 \geq 0$$

$$\theta^T x$$

$$x_1 + x_2 \geq 3$$

$$\theta = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$$

$$h_{\theta}(x) = g(\theta_0 + \underline{\theta_1} x_1 + \underline{\theta_2} x_2)$$

Decision boundary

$$x_1, x_2$$

$$h_{\theta}(x) = 0.5$$

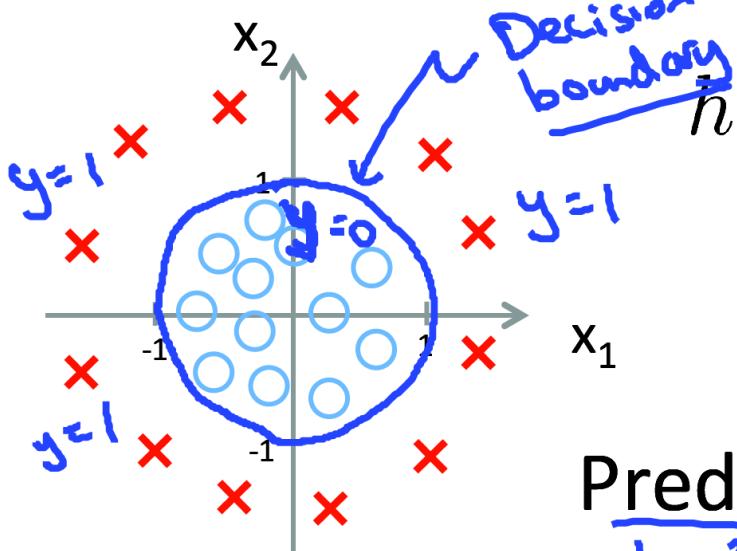
$$x_1 + x_2 = 3$$

$$x_1 + x_2 < 3$$

$$y = 0$$

Classification

Non-linear decision boundaries



Predict " $y = 1$ " if

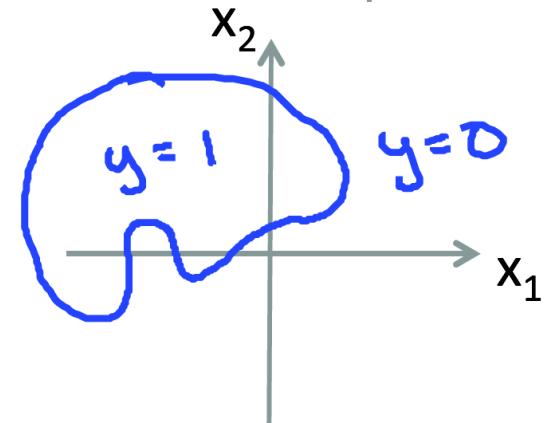
$$x_1^2 + x_2^2 \geq 1$$

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

$$\theta = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

$$-1 + x_1^2 + x_2^2 \geq 0$$

$$x_1^2 + x_2^2 \geq 1$$



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 \underline{x_1^2} + \theta_4 \underline{x_1^2 x_2} + \theta_5 \underline{x_1^2 x_2^2} + \theta_6 \underline{x_1^3 x_2} + \dots)$$

Classification

Training set:

m examples

$$h_{\theta}(x) = \frac{1}{1 + e^{-\underline{\theta}^T x}}$$

$$x \in \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix}$$

$$x_0 = 1, y \in \{0, 1\}$$

How to choose parameters θ ?

Classification: Cost function

Cost function

→ Linear regression:

Logistic

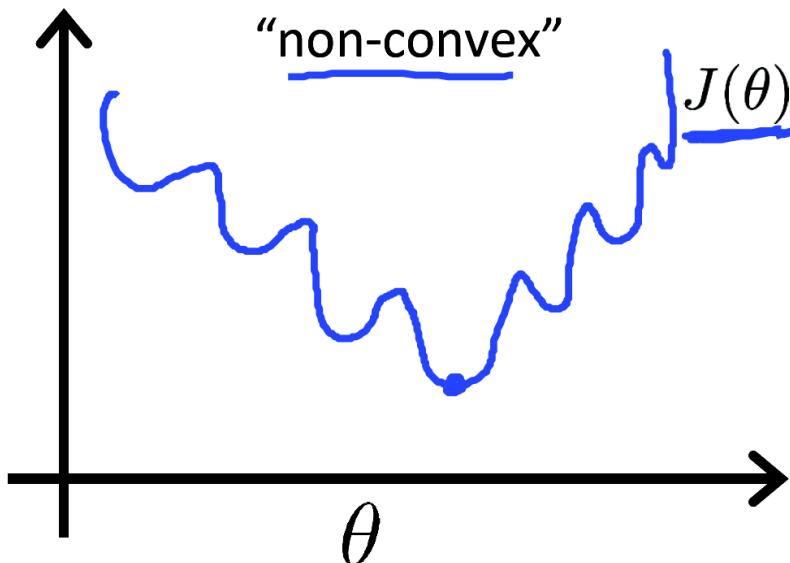
$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_\theta(x^{(i)}) - y^{(i)})^2$$

cost($h_\theta(x^{(i)})$, $y^{(i)}$)

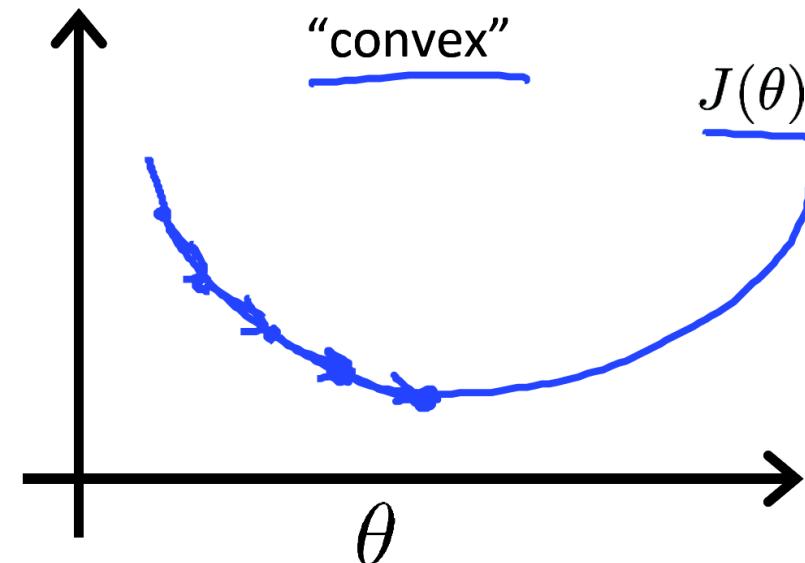
Cost($h_\theta(x^{(i)})$, $y^{(i)}$)

$$\frac{1}{2} (h_\theta(x^{(i)}) - y^{(i)})^2$$

$$h_\theta(x^{(i)}) = e^{-\theta^T x^{(i)}}$$



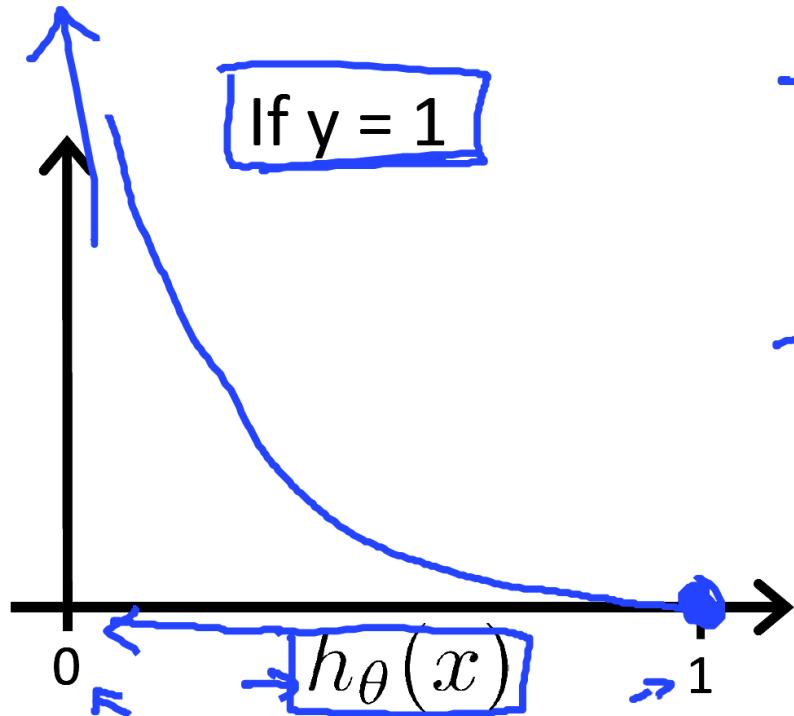
"non-convex"



"convex"

Classification: Logistic Regression: Cost function

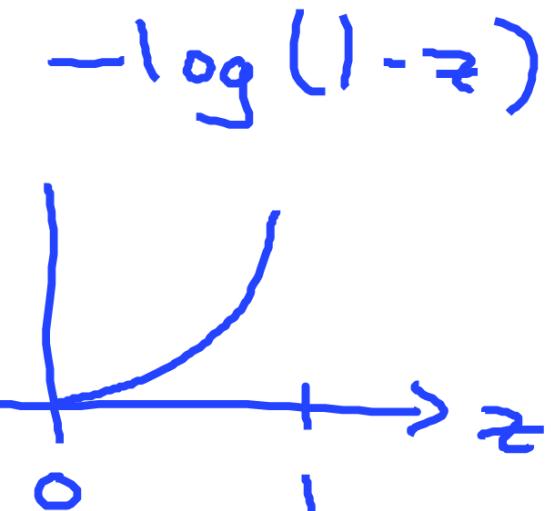
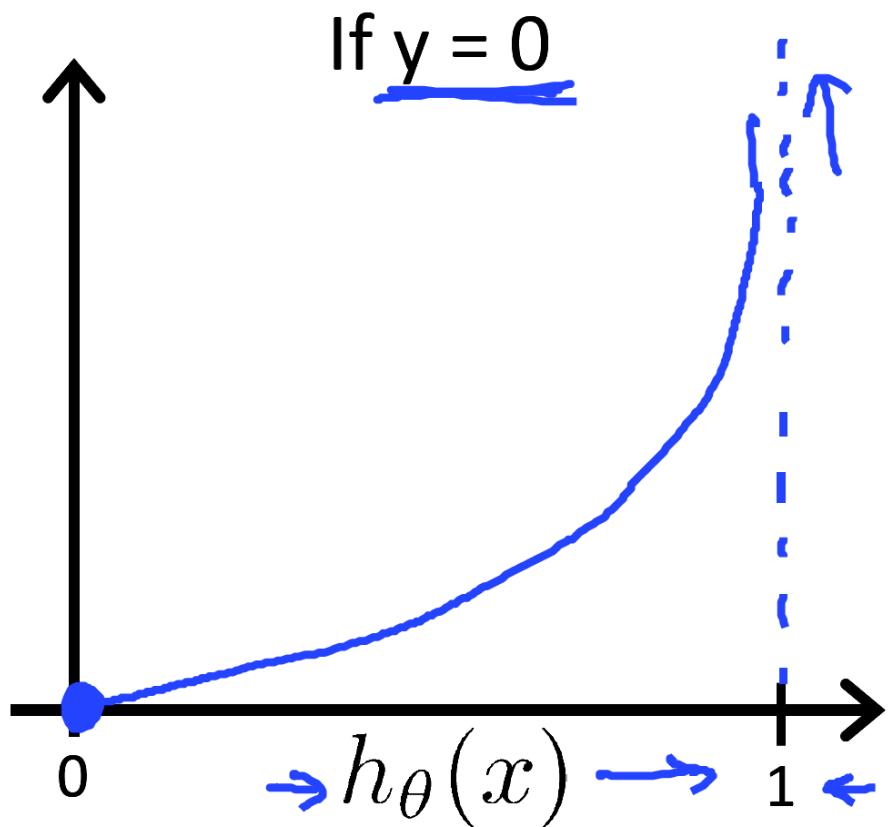
$$\text{Cost}(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y = 1 \\ -\log(1 - h_\theta(x)) & \text{if } y = 0 \end{cases}$$



- Cost = 0 if $y = 1, h_\theta(x) = 1$
But as $h_\theta(x) \rightarrow 0$
 $\text{Cost} \rightarrow \infty$
- Captures intuition that if $h_\theta(x) = 0$,
(predict $P(y = 1|x; \theta) = 0$), but $y = 1$,
we'll penalize learning algorithm by a very
large cost.

Classification: Logistic Regression: Cost function

$$\text{Cost}(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y = 1 \\ -\log(1 - h_\theta(x)) & \text{if } y = 0 \end{cases}$$



Classification: Logistic Regression: Cost function

Logistic regression cost function

$$\rightarrow J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_\theta(x^{(i)}), y^{(i)})$$

$$\rightarrow \text{Cost}(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y = 1 \\ -\log(1 - h_\theta(x)) & \text{if } y = 0 \end{cases}$$

Note: $y = 0$ or 1 always

$$\rightarrow \text{Cost}(h_\theta(x), y) = -y \log(h_\theta(x)) - (1-y) \log(1-h_\theta(x))$$

If $y=1$: $\text{Cost}(h_\theta(x), y) = -\log h_\theta(x)$

If $y=0$: $\text{Cost}(h_\theta(x), y) = -\log(1-h_\theta(x))$

Classification: Logistic Regression: Cost function

Logistic regression cost function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_\theta(x^{(i)}), y^{(i)})$$

$$= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_\theta(x^{(i)})) \right]$$

To fit parameters θ :

$$\min_{\theta} J(\theta)$$

Gret $\underline{\theta}$

To make a prediction given new \underline{x} :

$$\text{Output } \underline{h_\theta(x)} = \frac{1}{1+e^{-\theta^T x}}$$

$$\underline{\text{p}(\text{y}=1 | \underline{x}; \theta)}$$

Classification: Logistic Regression: Cost function

$$\rightarrow J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_\theta(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

(simultaneously update all θ_j)

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Classification: Logistic Regression: Cost function

Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_\theta(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\rightarrow \theta_j := \theta_j - \alpha \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

(simultaneously update all θ_j)

$$\Theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \quad \text{for } i = 0 \dots n$$

$$h_\theta(x) = \Theta^\top x$$

$$h_\theta(x) = \frac{1}{1 + e^{-\Theta^\top x}}$$

Algorithm looks identical to linear regression!

Classification: Logistic Regression: Cost function

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1+e^{-z}}$$

$$\text{Cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1-y) \log(1-h_{\theta}(x))$$

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m y \log h_{\theta}(x_i) + (1-y) \log(1-h_{\theta}(x_i))$$

$$1. \frac{d}{dz} \left(\frac{1}{1+e^{-z}} \right) \quad \text{substitute } 1+e^z = u$$

$$\frac{d}{du} \left(\frac{1}{u} \right) = \frac{d}{du} (u^{-1}) = -u^{-2}$$

$$2. \frac{d}{dz} \left(\frac{1}{1+e^{-z}} \right) = -e^{-z}$$

$$3. \frac{d}{dz} \left(\frac{1}{1+e^{-z}} \right) = \frac{dg}{du} \cdot \frac{du}{dz}$$

$$= -u^{-2} x - e^{-z}$$

$$= \frac{1}{u^2} x e^{-z}$$

$$= \frac{e^{-z}}{(1+e^{-z})^2}$$

Derivatives
cheat sheet

$$\frac{d}{dx} x^3 = 3x^2$$

$$\frac{d}{dx} \sin(x) = \cos(x)$$

$$\frac{d}{dx} \cos(x) = -\sin(x)$$

$$\frac{d}{dx} \tan(x) = \frac{1}{\cos^2(x)}$$

$$\frac{d}{dx} (e^x) = e^x$$

$$\frac{d}{dx} e^{-x} = -e^{-x}$$

$$\frac{d}{dx} \ln(x) = \frac{1}{x} = x^{-1}$$

$$F(x) = f(g(x))$$

$$F'(x) = f'(g(x)) \cdot g'(x)$$

Classification: Logistic Regression: Cost function

4. $e^{-z}/(1+e^{-z})^2 = \frac{1+e^{-z}-1}{(1+e^{-z})^2} = \frac{1}{(1+e^{-z})} - \frac{1}{(1+e^{-z})^2}$

$$= \frac{1}{1+e^{-z}} \left(1 - \frac{1}{1+e^{-z}} \right)$$

$$= \boxed{g(z)(1-g(z))}$$

5. $\frac{\partial}{\partial \theta} J_{\theta} = -\frac{1}{m} \left[\frac{\partial}{\partial \theta} \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)})) \right]$

$$= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \frac{1}{h_{\theta}(x^{(i)})} \frac{\partial}{\partial \theta} h_{\theta}(x^{(i)}) + (1-y^{(i)}) \frac{1}{1-h_{\theta}(x^{(i)})} \left(-\frac{\partial}{\partial \theta} h_{\theta}(x^{(i)}) \right) \right]$$

$$= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \frac{x_j^{(i)}}{h_{\theta}(x^{(i)})} (h_{\theta}(x^{(i)})(1-h_{\theta}(x^{(i)})) - (1-y^{(i)}) \frac{x_j^{(i)}}{1-h_{\theta}(x^{(i)})} (h_{\theta}(x^{(i)})(1-h_{\theta}(x^{(i)}))) \right]$$

Classification: Logistic Regression: Cost function

$$\begin{aligned}
 4. \quad e^{-z}/(1+e^{-z})^2 &= \frac{1+e^{-z}-1}{(1+e^{-z})^2} = \frac{1}{(1+e^{-z})} - \frac{1}{(1+e^{-z})^2} \\
 &= \frac{1}{1+e^{-z}} \left(1 - \frac{1}{1+e^{-z}} \right) \\
 &= \boxed{g(z)(1-g(z))}
 \end{aligned}$$

$$\begin{aligned}
 5. \quad \frac{\partial}{\partial \theta} J_{\theta} &= -1/m \left[\frac{\partial}{\partial \theta} \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)})) \right] \\
 &= -1/m \left[\sum_{i=1}^m y^{(i)} \frac{1}{h_{\theta}(x^{(i)})} \frac{\partial}{\partial \theta} h_{\theta}(x^{(i)}) + (1-y^{(i)}) \frac{1}{1-h_{\theta}(x^{(i)})} \left(-\frac{\partial}{\partial \theta} h_{\theta}(x^{(i)}) \right) \right] \\
 &= -1/m \left[\sum_{i=1}^m y^{(i)} \frac{x_j^{(i)}}{h_{\theta}(x^{(i)})} (h_{\theta}(x^{(i)}))(1-h_{\theta}(x^{(i)})) - (1-y^{(i)}) \frac{x_j^{(i)}}{1-h_{\theta}(x^{(i)})} (h_{\theta}(x^{(i)}))(1-h_{\theta}(x^{(i)})) \right]
 \end{aligned}$$

Classification: Logistic Regression: Cost function

$$= -\frac{1}{m} \left[\sum y^{(i)} x_j^{(i)} (1 - h_{\theta} x^{(i)}) - (1 - y^{(i)}) x_j^{(i)} h_{\theta} x^{(i)} \right]$$

$$= -\frac{1}{m} \left[\sum y^{(i)} x_j^{(i)} - y^{(i)} x_j^{(i)} h_{\theta} x^{(i)} - x_j^{(i)} h_{\theta} x^{(i)} + y^{(i)} x_j^{(i)} h_{\theta} x^{(i)} \right]$$

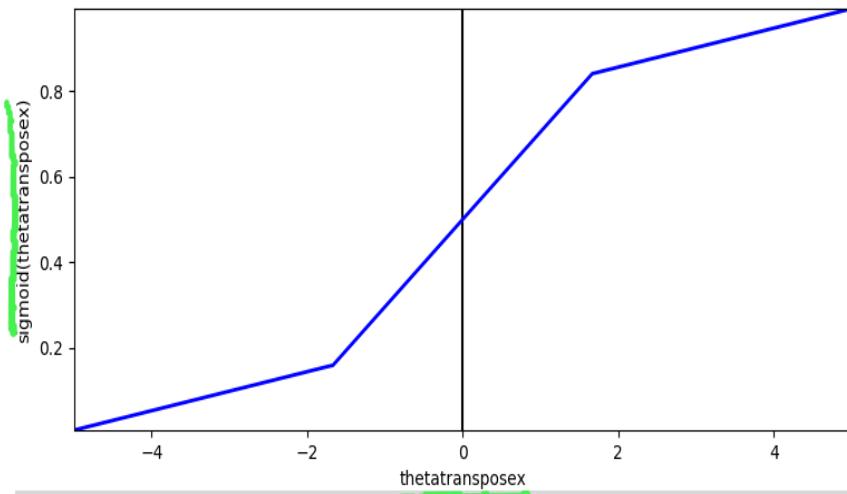
$$= -\frac{1}{m} \left[\sum y^{(i)} x_j^{(i)} - x_j^{(i)} h_{\theta} x^{(i)} \right]$$

$$= \frac{1}{m} \sum_{i=1}^m (h_{\theta} x^{(i)} - y^{(i)}) x_j^{(i)}$$

$$= \boxed{\frac{1}{m} \sum_{i=1}^m (\sigma(\theta^T x^{(i)}) - y^{(i)}) x_j^{(i)}}$$

Classification: Logistic Regression

```
def compute_cost(thetas,X,y):  
    m = X.shape[0]  
    print("training set size:\n",X.shape[0])  
    costval=[]  
    sigval=[]  
    ttxvals=[]  
    for theta in thetas:  
        theta = np.reshape(theta,(len(theta),1))  
        ttx=X.dot(theta)  
        sig=sigmoid(X.dot(theta))  
        sigval.append(sig)  
        lg = -log(sig)  
        oneminuslog=-log(1-sigmoid(X.dot(theta)))  
        J=cost(X,y,theta)  
        costval.append(J)  
    if theta==1:  
        print("X:\n",X)  
        print("X.dot(theta):\n",ttx)  
        print("sig:",sig)  
        plot(ttx,sig,"thet transpose","sigmoid(thet transpose)", "")  
        print("log:",lg)  
        plot(sig,lg,"sigmoid(thet transpose)", "-log(sigmoid(thet transpose))", "")  
        plot(ttx,lg,"thet transpose", "-log(sigmoid(thet transpose))", "")  
        print("oneminuslog:",oneminuslog)  
        plot(sig,oneminuslog,"sigmoid(thet transpose)", "-log(1-sigmoid(thet transpose))", "")  
        plot(ttx,oneminuslog,"thet transpose", "-log(1-sigmoid(thet transpose))", "")  
        print("Cost:",J)  
    plot(thetas,costval,"theta","cost", "")
```



```
def sigmoid(X):  
    return 1 / (1 + np.exp(-X))
```

• also called the squashing function.

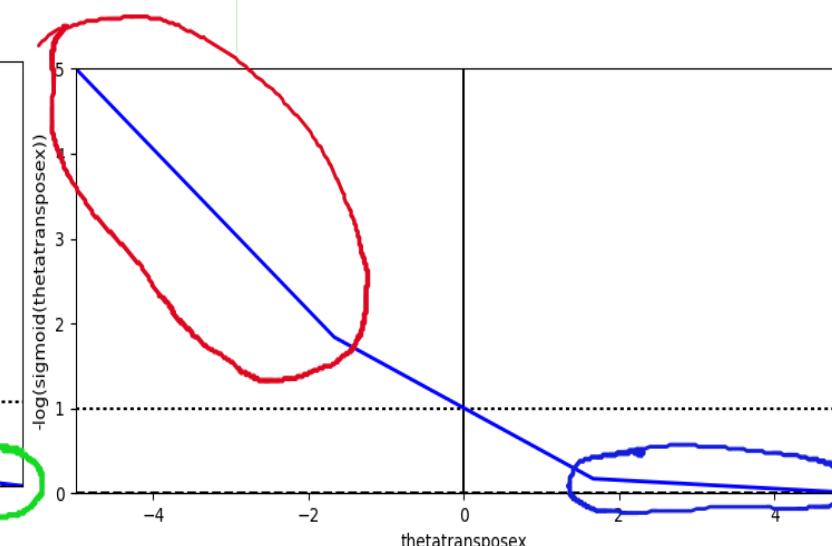
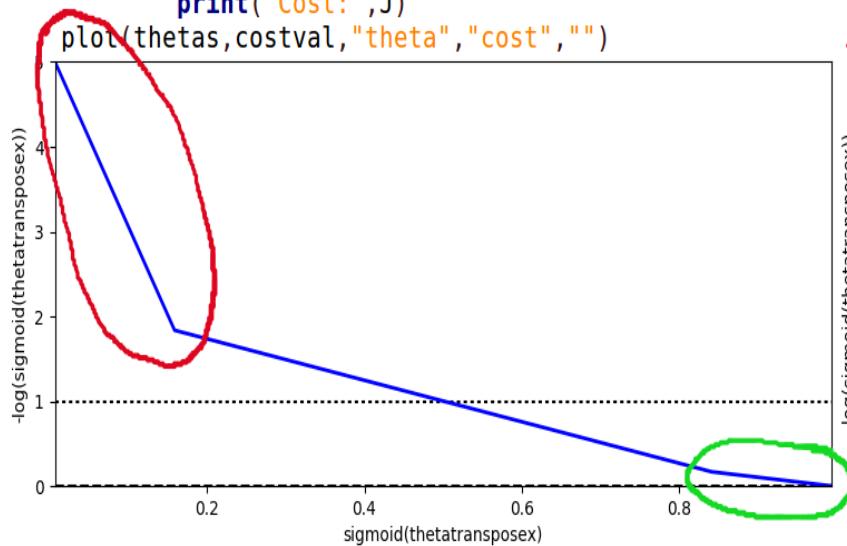
X: [0. 3.33333333 6.66666667 10.]
y: [False False True True]
X:
[[-5.]
[-1.66666667]
[1.66666667]
[5.]]
| scaled input
X.dot(theta):
[[-5.]
[-1.66666667]
[1.66666667]
[5.]]
| $\Theta^T X$ for $\Theta \rightarrow 1$
sig: [[0.00669285]
[0.1588691]
[0.8411309]
[0.99330715]]
| Sigmoid or
log: [[5.00671535]
[1.83967466]
[0.17300799]
[0.00671535]]
|
oneminuslog: [[0.00671535]
[0.17300799]
[1.83967466]
[5.00671535]]
|
Cost: [0.08986167]

Classification: Logistic Regression

```

def compute_cost(thetas,X,y):
    m = X.shape[0]
    print("training set size:\n",X.shape[0])
    costval=[]
    sigval=[]
    ttxvals=[]
    for theta in thetas:
        theta = np.reshape(theta,(len(theta),1))
        ttx=X.dot(theta)
        sig=sigmoid(X.dot(theta))
        sigval.append(sig)
        lg = -log(sig)
        oneminuslog=-log(1-sigmoid(X.dot(theta)))
        J=cost(X,y,theta)
        costval.append(J)
        if theta==1:
            print("X:\n",X)
            print("X.dot(theta):\n",ttx)
            print("sig:",sig)
            plot(ttx,sig,"thet transpose","sigmoid(thet transpose)", "")
            print("log:",lg)
            plot(sig,lg,"sigmoid(thet transpose)","-log(sigmoid(thet transpose))", "")
            plot(ttx,lg,"thet transpose","-log(sigmoid(thet transpose))", "")
            print("oneminuslog:",oneminuslog)
            plot(sig,oneminuslog,"sigmoid(thet transpose)","-log(1-sigmoid(thet transpose))", "")
            plot(ttx,oneminuslog,"thet transpose","-log(1-sigmoid(thet transpose))", "")
            print("Cost:",J)
            plot(thetas,costval,"theta","cost", "")

```



$$Cost(h_\theta x, y) = \boxed{-y \log(h_\theta x)} - (1-y) \log(1-h_\theta x)$$

- when ($y=1$), and if $h_\theta x$ is also 1 then the cost tends to 0, because the algo gets it right.
- As $\theta^T x \gg 0$, $h_\theta x \rightarrow 1$ and, $-\log(h_\theta x) \rightarrow 0$
- when ($y=1$) but if $h_\theta x \rightarrow 0$, the cost goes up or we penalize the algo

```

X: [ 0. 3.33333333 6.66666667 10. ]
y: [ False False True True ]
X:
[[ -5.      ]
 [-1.66666667]
 [ 1.66666667]
 [ 5.      ]]
X.dot(theta):
[[ -5.      ]
 [-1.66666667]
 [ 1.66666667]
 [ 5.      ]]
sig: [[ 0.00669285]
 [ 0.1588691 ]
 [ 0.8411309 ]
 [ 0.99330715]]
log: [[ 5.00671535]
 [ 1.83967466]
 [ 0.17300799]
 [ 0.00671535]]
oneminuslog: [[ 0.00671535]
 [ 0.17300799]
 [ 1.83967466]
 [ 5.00671535]]
Cost: [ 0.08986167]

```

Classification: Logistic Regression

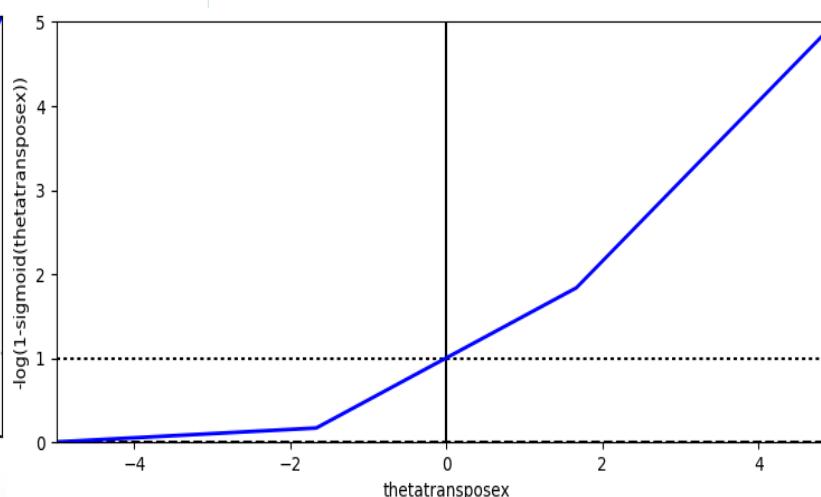
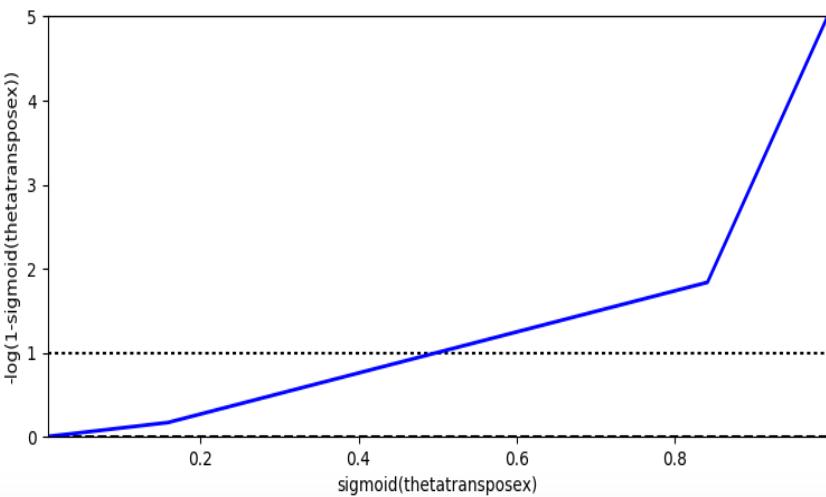
```

def compute_cost(thetas,X,y):
    m = X.shape[0]
    print("training set size:\n",X.shape[0])
    costval=[]
    sigval=[]
    ttxvals=[]
    for theta in thetas:
        theta = np.reshape(theta,(len(theta),1))
        ttx=X.dot(theta)
        sig=sigmoid(X.dot(theta))
        sigval.append(sig)
        lg = -log(sig)
        oneminuslog=-log(1-sigmoid(X.dot(theta)))
        J=cost(X,y,theta)
        costval.append(J)
    if theta==1:
        print("X:\n",X)
        print("X.dot(theta):\n",ttx)
        print("sig:",sig)
        plot(ttx,sig,"thet transpose","sigmoid(thet transpose)", "")
        print("log:",lg)
        plot(sig,lg,"sigmoid(thet transpose)", "-log(sigmoid(thet transpose))", "")
        plot(ttx,lg,"thet transpose", "-log(sigmoid(thet transpose))", "")
        print("oneminuslog:",oneminuslog)
        plot(sig,oneminuslog,"sigmoid(thet transpose)", "-log(1-sigmoid(thet transpose))", "")
        plot(ttx,oneminuslog,"thet transpose", "-log(1-sigmoid(thet transpose))", "")
        print("Cost:",J)
    plot(thetas,costval,"theta","cost", "")

```

$$Cost(h_\theta x, y) = -y \log(h_\theta x) - (1-y) \log(1-h_\theta x)$$

• when $y=0$...



```

X: [ 0. 3.33333333 6.66666667 10
y: [ False False True True]
X:
[[ -5. ]
[-1.66666667]
[ 1.66666667]
[ 5. ]]
X.dot(theta):
[[ -5. ]
[-1.66666667]
[ 1.66666667]
[ 5. ]]
sig: [[ 0.00669285]
[ 0.1588691 ]
[ 0.8411309 ]
[ 0.99330715]]
log: [[ 5.00671535]
[ 1.83967466]
[ 0.17300799]
[ 0.00671535]]
oneminuslog: [[ 0.00671535]
[ 0.17300799]
[ 1.83967466]
[ 5.00671535]]
Cost: [ 0.08986167]

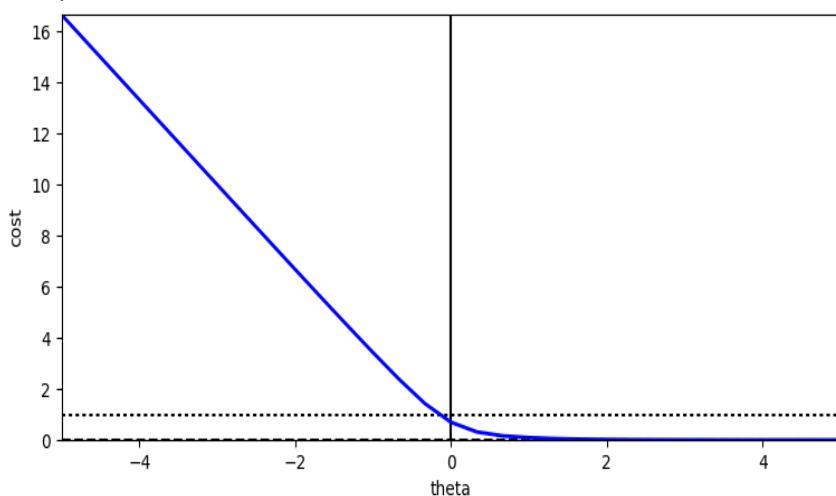
```

Classification: Logistic Regression

```

def compute_cost(thetas,X,y):
    m = X.shape[0]
    print("training set size:\n",X.shape[0])
    costval=[]
    sigval=[]
    ttxvals=[]
    for theta in thetas:
        theta = np.reshape(theta,(len(theta),1))
        ttx=X.dot(theta)
        sig=sigmoid(X.dot(theta))
        sigval.append(sig)
        lg = -log(sig)
        oneminuslog=-log(1-sigmoid(X.dot(theta)))
        J=cost(X,y,theta)
        costval.append(J)
    if theta==1:
        print("X:\n",X)
        print("X.dot(theta):\n",ttx)
        print("sig:",sig)
        plot(ttx,sig,"thet transpose","sigmoid(thet transpose)", "")
        print("log:",lg)
        plot(sig,lg,"sigmoid(thet transpose)","-log(sigmoid(thet transpose))", "")
        plot(ttx,lg,"thet transpose","-log(sigmoid(thet transpose))", "")
        print("oneminuslog:",oneminuslog)
        plot(sig,oneminuslog,"sigmoid(thet transpose)","-log(1-sigmoid(thet transpose))", "")
        plot(ttx,oneminuslog,"thet transpose","-log(1-sigmoid(thet transpose))", "")
        print("Cost:",J)
    plot(thetas,costval,"theta","cost", "")

```



$$J(\theta) = \frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)}) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)})) \right]$$

```

def cost(X,y,theta):
    m = X.shape[0]
    return (1./m) * (-transpose(y).dot(log(sigmoid(X.dot(theta)))))
            - transpose(1-y).dot(log(1-sigmoid(X.dot(theta)))))

```

X: [0. 3.33333333 6.66666667 10.]
 y: [False False True True]
 X:
 [[-5.]
 [-1.66666667]
 [1.66666667]
 [5.]]
 X.dot(theta):
 [[-5.]
 [-1.66666667]
 [1.66666667]
 [5.]]
 sig: [[0.0069285]
 [0.1588691]
 [0.8411309]
 [0.99330715]]
 log: [[5.00671535]
 [1.83967466]
 [0.17300799]
 [0.00671535]]
 oneminuslog: [[0.00671535]
 [0.17300799]
 [1.83967466]
 [5.00671535]]
 Cost: [0.08986167]

1 + } / m
 ↙

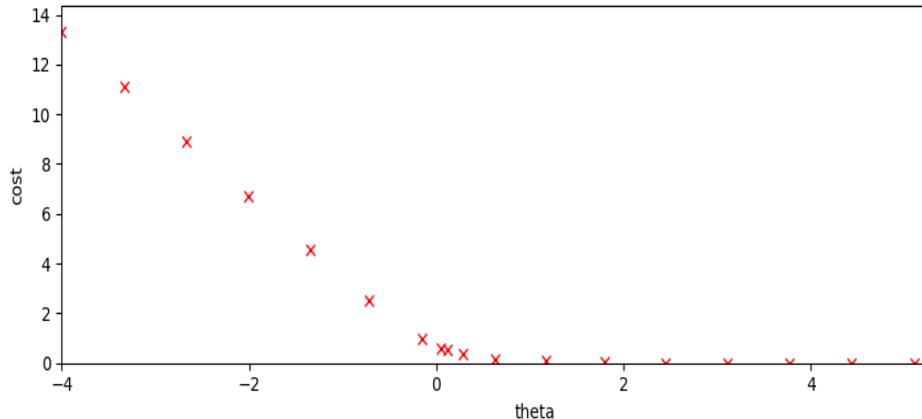
Classification: Logistic Regression: BGD

```

def gradient(X,y,theta):
    m = len(X)
    return 1/m * X.T.dot(np.sum(sigmoid(X.dot(theta))-y,axis=1).reshape(-1,1))

def plot_gradient_descent(theta, eta, theta_path=None):
    m = len(tumor_size_scaled)
    n_iterations = 20
    costvalprev=0.0
    change=0.0
    for iteration in range(n_iterations):
        y_predict=sigmoid(thetaTransposeX(tumor_size_scaled,theta))
        costval=cost(tumor_size_scaled,mal,theta)
        if iteration < iterations_to_track:
            if theta_path_bgd_10 is not None:
                theta_path_bgd_10.append(theta)
            if cost_path_bgd_10 is not None:
                cost_path_bgd_10.append(costval)
            if y_predict_bgd_10 is not None:
                y_predict_bgd_10.append(y_predict)
        gradients = gradient(tumor_size_scaled,mal,theta)
        if gradients > 0: #going up so reverse direction
            theta = theta + eta * gradients
        else:
            theta = theta - eta * gradients
        if iteration !=0:
            if (costvalprev-costval < .0004):
                print("breaking:")
                break
        costvalprev=costval
    print("gradient:",gradients," theta:",theta," costval:",costval)

```



$$\text{gradient} \Rightarrow \frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (\sigma(\theta^T x^i) - y^i) x^i_j$$

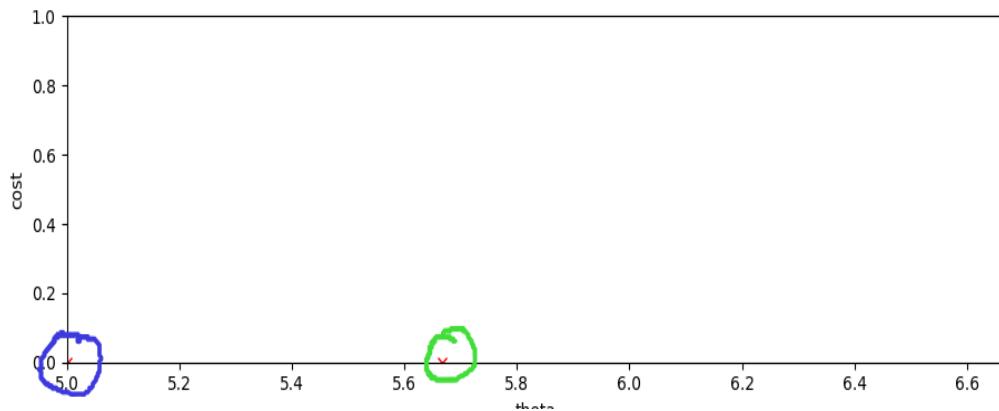
- convergence rule
- going up or down the slope.

```
(ml_home) mohit@nomind:~/Work/ArtificialIntelligence$ ./ML/logisticregression/logisticregressiongraph2.py
theta init: [[-4.]]
gradient: [[-6.66242993]] theta: [[-3.33375701]] costval: [ 13.33396924]
gradient: [[-6.65383834]] theta: [[-2.66837317]] costval: [ 11.11445126]
gradient: [[-6.62806817]] theta: [[-2.00556636]] costval: [ 8.90039917]
gradient: [[-6.5524319]] theta: [[-1.35032317]] costval: [ 6.70261041]
gradient: [[-6.33731371]] theta: [[-0.7165918]] costval: [ 4.55173898]
gradient: [[-5.62129158]] theta: [[-0.15446264]] costval: [ 2.53464648]
gradient: [[-2.05356512]] theta: [[ 0.05089387]] costval: [ 0.99110364]
gradient: [[ 0.70340617]] theta: [[ 0.12123449]] costval: [ 0.6128099]
gradient: [[ 1.63848263]] theta: [[ 0.28508275]] costval: [ 0.51625916]
gradient: [[ 3.45052092]] theta: [[ 0.63013485]] costval: [ 0.34948867]
gradient: [[ 5.39207964]] theta: [[ 1.16934281]] costval: [ 0.17096596]
gradient: [[ 6.22228]] theta: [[ 1.79157081]] costval: [ 0.06802139]
gradient: [[ 6.50516429]] theta: [[ 2.44208724]] costval: [ 0.02469341]
gradient: [[ 6.61665006]] theta: [[ 3.10315284]] costval: [ 0.00846787]
gradient: [[ 6.64785935]] theta: [[ 3.76793878]] costval: [ 0.0028289]
gradient: [[ 6.66043297]] theta: [[ 4.43398208]] costval: [ 0.00093592]
gradient: [[ 6.66460989]] theta: [[ 5.10044307]] costval: [ 0.00030861]
breaking:
theta best: [[ 5.76704197]]
```

Classification:Logistic Regression:BGD

```
def gradient(X,y,theta):
    m = len(X)
    return 1/m * X.T.dot(np.sum(sigmoid(X.dot(theta))-y,axis=1).reshape(-1,1))

def plot_gradient_descent(theta, eta, theta_path=None):
    m = len(tumor_size_scaled)
    n_iterations = 20
    costvalprev=0.0
    change=0.0
    for iteration in range(n_iterations):
        y_predict=sigmoid(thetaTransposeX(tumor_size_scaled,theta))
        costval=cost(tumor_size_scaled,mal,theta)
        if iteration < iterations_to_track:
            if theta_path_bgd_10 is not None:
                theta_path_bgd_10.append(theta)
            if cost_path_bgd_10 is not None:
                cost_path_bgd_10.append(costval)
            if y_predict_bgd_10 is not None:
                y_predict_bgd_10.append(y_predict)
        gradients = gradient(tumor_size_scaled,mal,theta)
        if iteration !=0:
            if (costvalprev-costval < .0004):
                print("breaking:")
                break
        if gradients > 0: #going up so reverse direction
            theta = theta + eta * gradients
        else:
            theta = theta - eta * gradients
        costvalprev=costval
    print("gradient:",gradients," theta:",theta," costval:",costval)
```



• theta init = 5
• theta best = 5.66

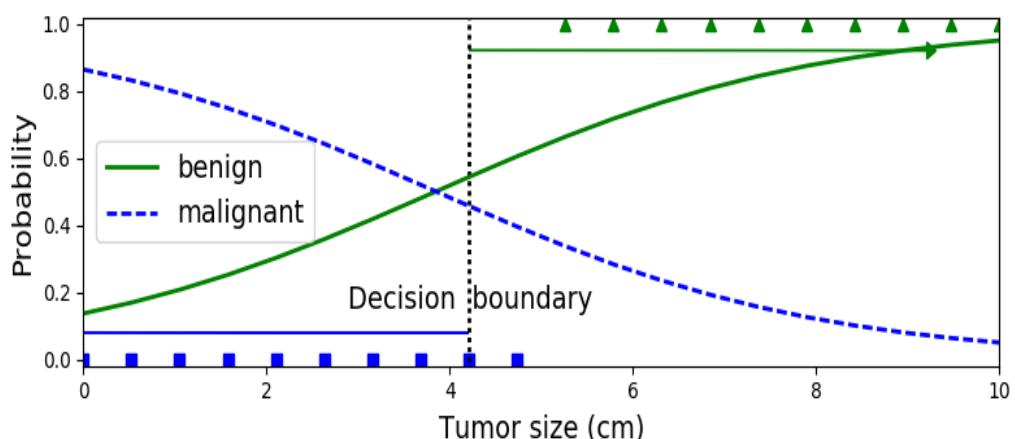
```
(ml_home) mohit@nomind:~/Work/ArtificialIntelligence$ ./ML/logisticregression/logisticregressiongraph2.py
thetainit: [[ 5.1]]
gradient: [[ 6.66586563]]  theta: [[ 5.66658656]]  costval: [ 0.00012017]
breaking:
theta best: [[ 5.66658656]]
```

Classification: Logistic Regression(scikit):descision boundary

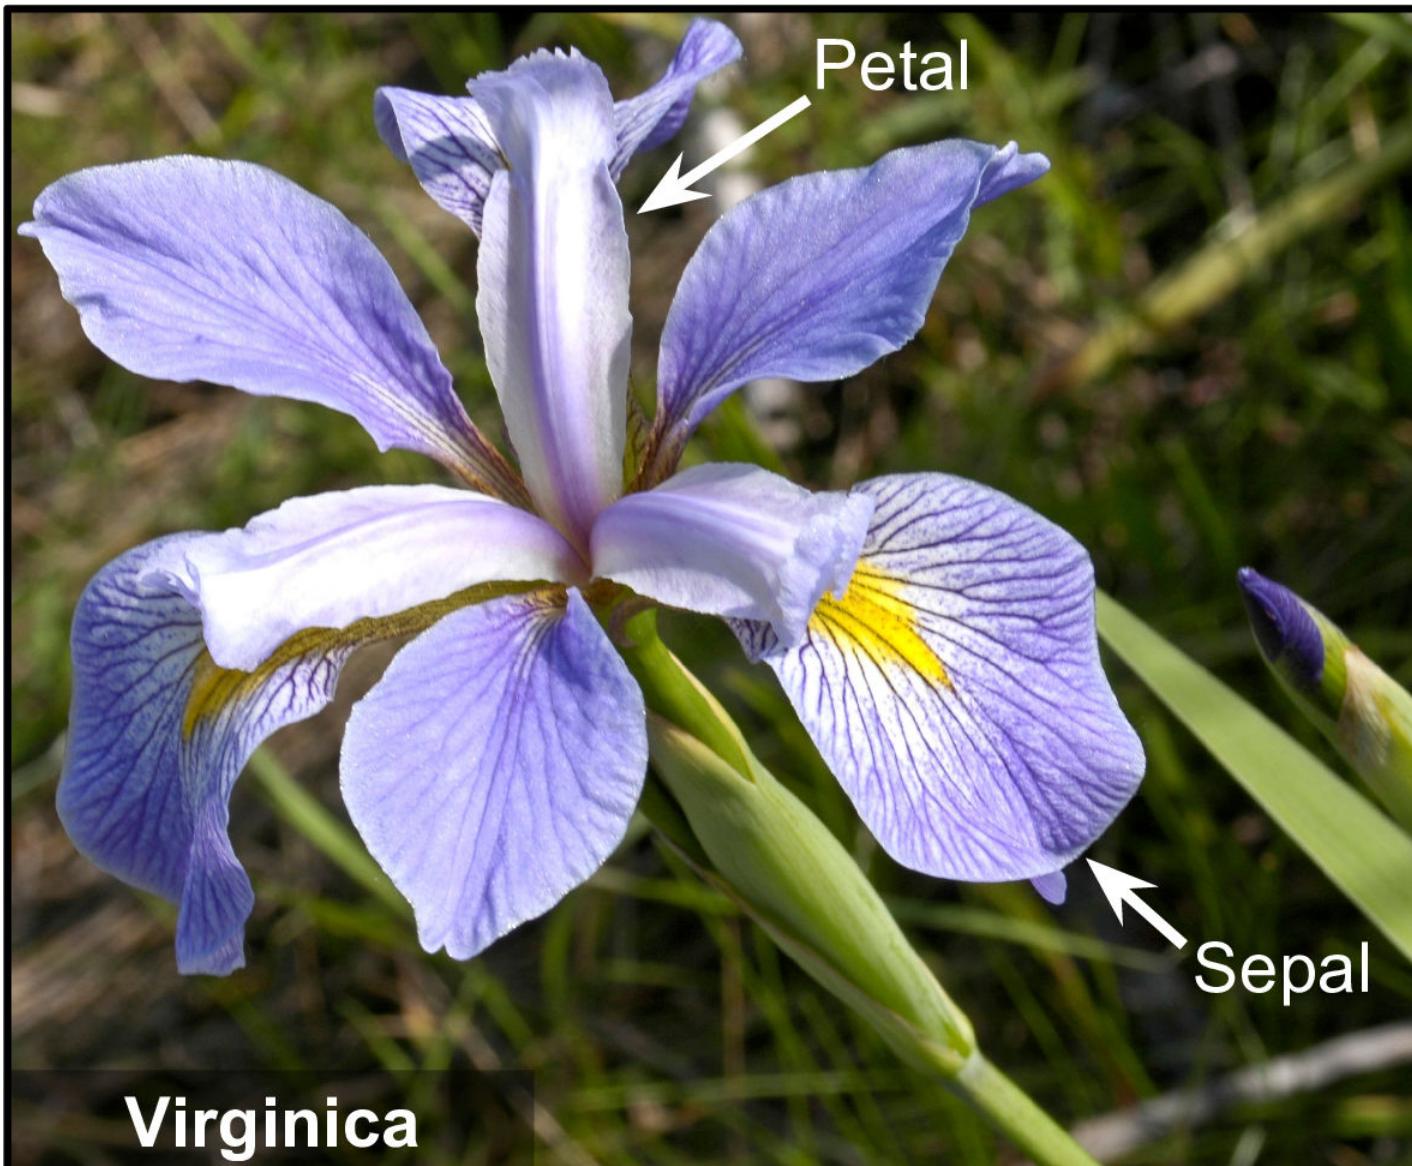
```
tumor_size=np.linspace(0.0, 10.0, num=20)
print("\nX:",tumor_size)
tumor_size_new=tumor_size.reshape(-1,1)
print("\nx_new:",tumor_size_new)

mal=(tumor_size>5)
print("\nmal:",mal)

log_reg = LogisticRegression()
log_reg.fit(tumor_size_new, mal)
print(log_reg)
print(log_reg.coef_)
print(log_reg.intercept_)
X_test = np.linspace(0.0, 10.0, 20).reshape(-1, 1)
mal_proba = log_reg.predict_proba(X_test)
decision_boundary = X_test[mal_proba[:, 1] >= 0.5][0]
print("decision_boundary:",decision_boundary)
```



Classification: Logistic Regression(scikit):descision boundary



Classification: Logistic Regression(scikit):descision boundary

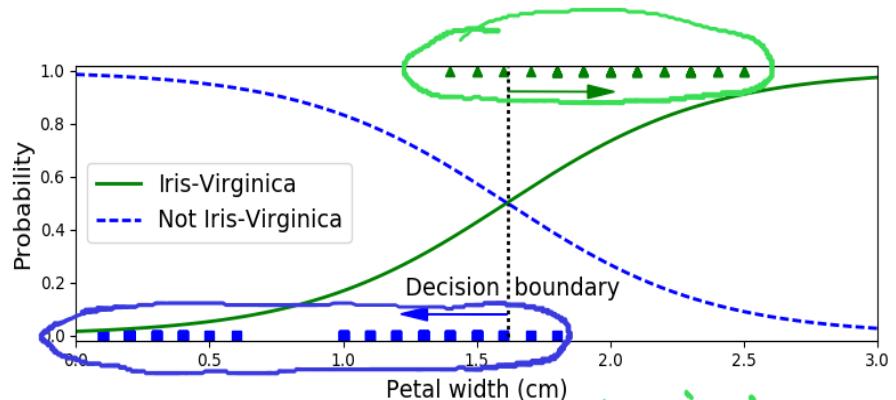
```
iris = datasets.load_iris()
print(list(iris.keys()))
print(iris.DESCR)

X = iris["data"][:, 3:] # petal width
y = (iris["target"] == 2).astype(np.int) # 1 if Iris-Virginica, else 0

log_reg = LogisticRegression()
log_reg.fit(X, y)
X1=np.linspace(0, 3, 1000);

X_new = np.linspace(0, 3, 1000).reshape(-1, 1)

y_proba = log_reg.predict_proba(X_new)
decision_boundary = X_new[y_proba[:, 1] >= 0.5][0]
print("decision_boundary:",decision_boundary)
```



- 1.4 - 2.5 cm Iris Virginica
- 0.1 - 1.8 cm other specie
- Decision Boundary at ~ 1.6 and the algorithm is not so sure around that mark.
- may be adding more features will help.

Classification: Logistic Regression(scikit): linear descision boundary

```
iris = datasets.load_iris()

X = iris["data"][:, (2, 3)] # petal length, petal width
y = (iris["target"] == 2).astype(np.int)

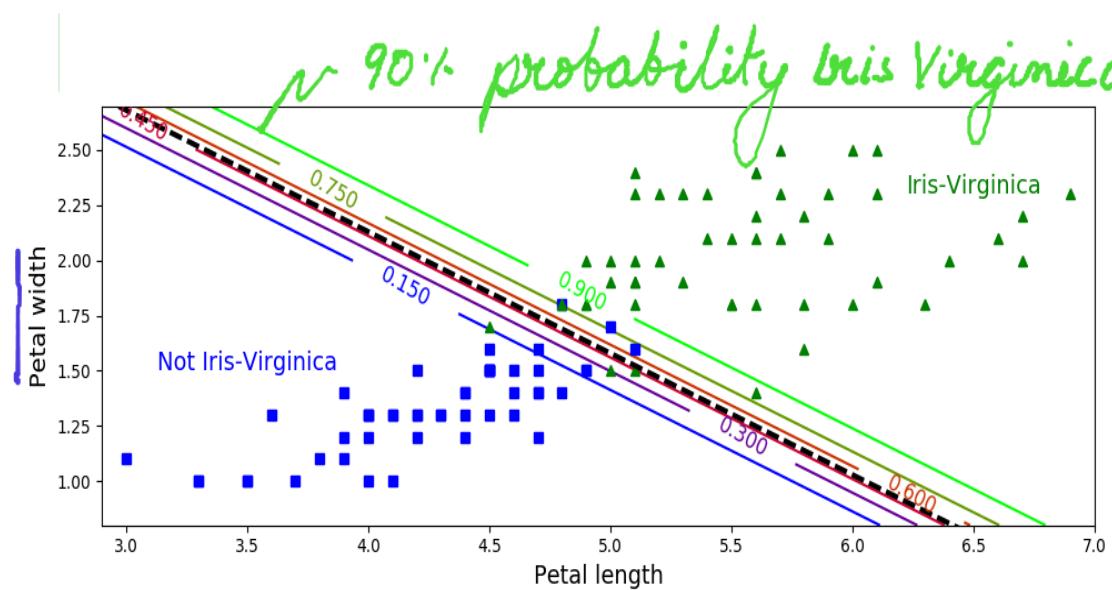
log_reg = LogisticRegression(C=10**10)
log_reg.fit(X, y)

x0, x1 = np.meshgrid(
    np.linspace(2.9, 7, 500).reshape(-1, 1),
    np.linspace(0.8, 2.7, 200).reshape(-1, 1),
)
X_new = np.c_[x0.ravel(), x1.ravel()]
y_proba = log_reg.predict_proba(X_new)
plt.figure(figsize=(10, 4))
plt.plot(X[y==0, 0], X[y==0, 1], "bs")
plt.plot(X[y==1, 0], X[y==1, 1], "g^")
zz = y_proba[:, 1].reshape(x0.shape)
contour = plt.contour(x0, x1, zz, cmap=plt.cm.brg)

left_right = np.array([2.9, 7])
boundary = -(log_reg.coef_[0][0] * left_right + log_reg.intercept_[0]) / log_reg.coef_[0][1]

plt.clabel(contour, inline=1, fontsize=12)
plt.plot(left_right, boundary, "k--", linewidth=3)
plt.text(3.5, 1.5, "Not Iris-Virginica", fontsize=14, color="b", ha="center")
plt.text(6.5, 2.3, "Iris-Virginica", fontsize=14, color="g", ha="center")
plt.xlabel("Petal length", fontsize=14)
plt.ylabel("Petal width", fontsize=14)
plt.axis([2.9, 7, 0.8, 2.7])
```

- $C = \frac{1}{\lambda}$ controls regularization
- LR adds L_2 penalty by default.

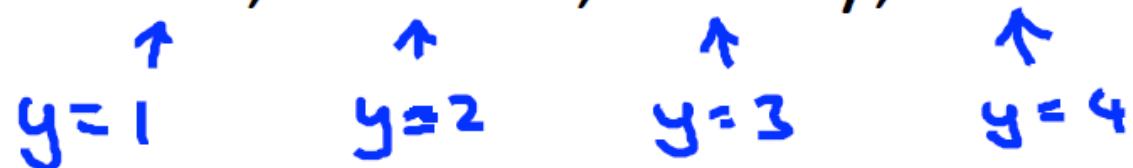


- 2 features on a contour plot.
- dashed line is the DB.
- each parallel line represents a specific probability. From 15% - 90%.

Multiclass Classification

Multiclass classification

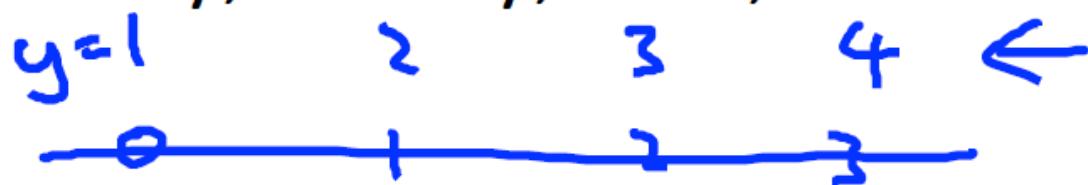
Email foldering/tagging: Work, Friends, Family, Hobby



Medical diagrams: Not ill, Cold, Flu

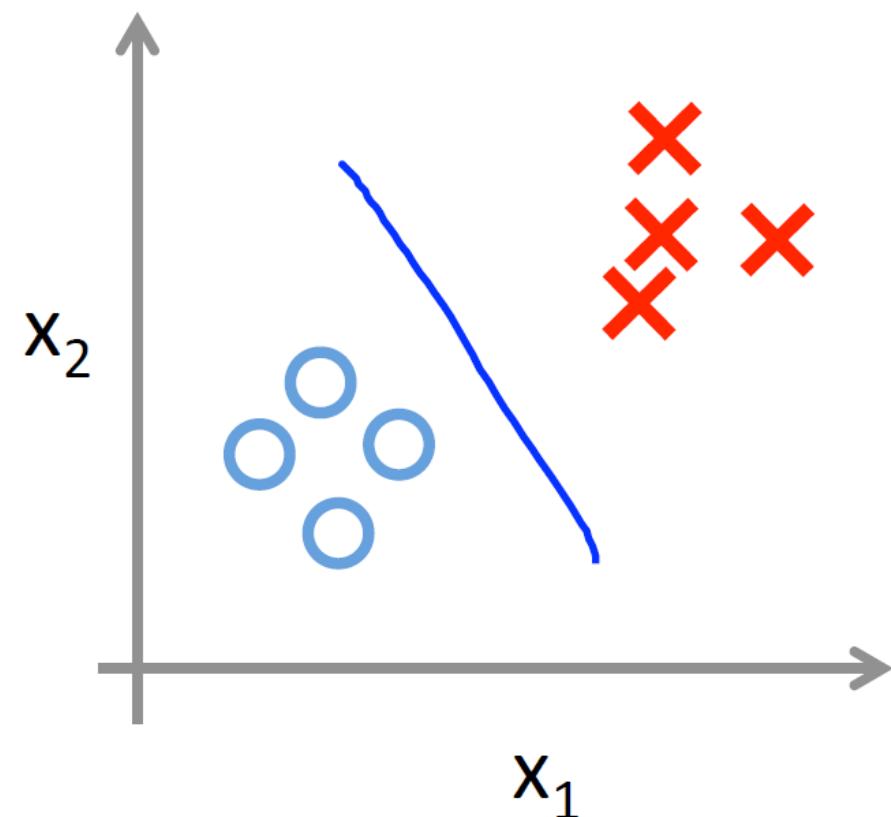


Weather: Sunny, Cloudy, Rain, Snow

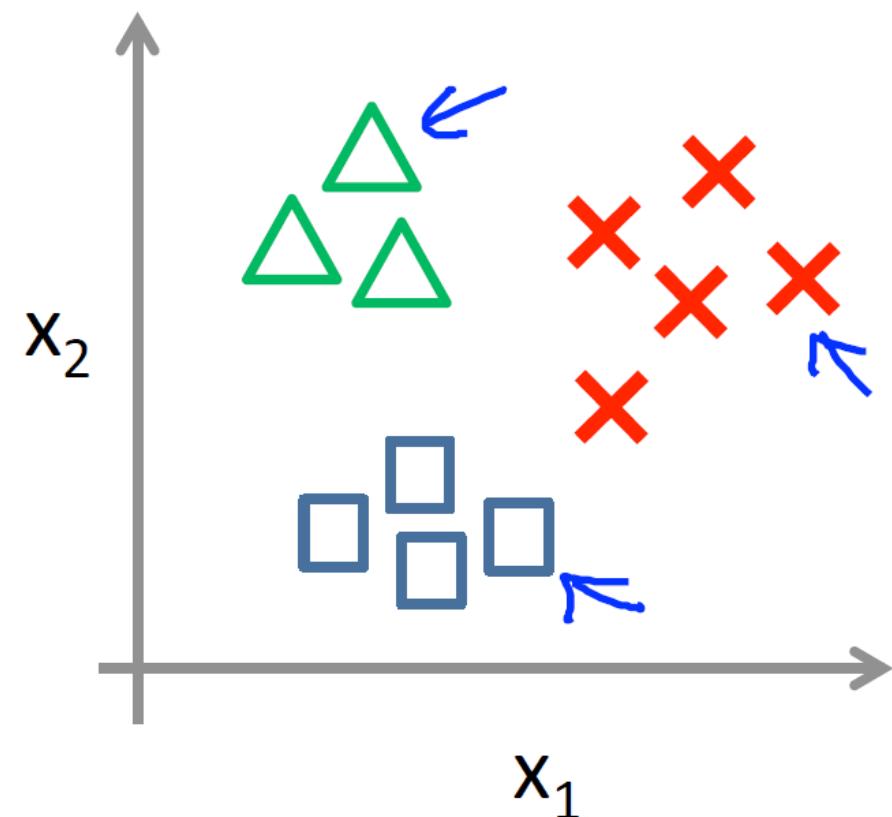


Multiclass Classification

Binary classification:

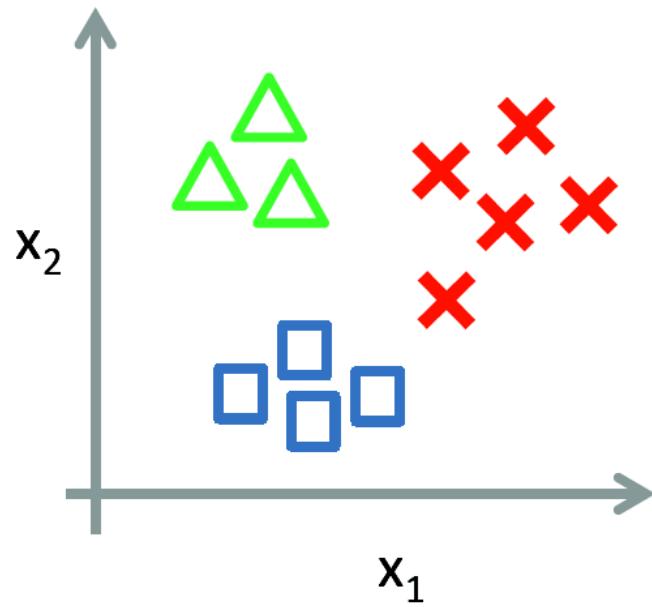


Multi-class classification:



Multiclass Classification

One-vs-all (one-vs-rest):

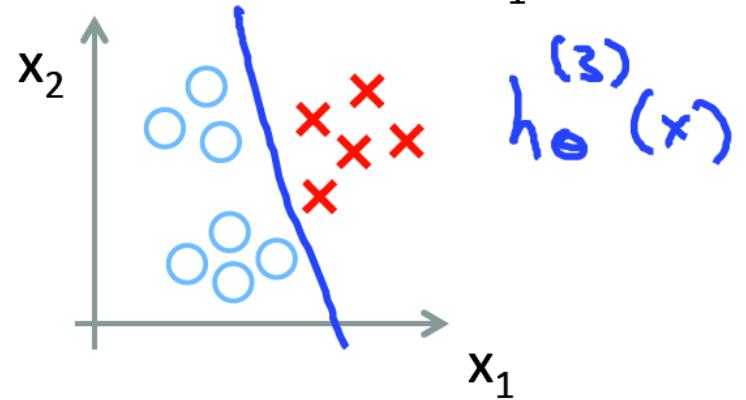
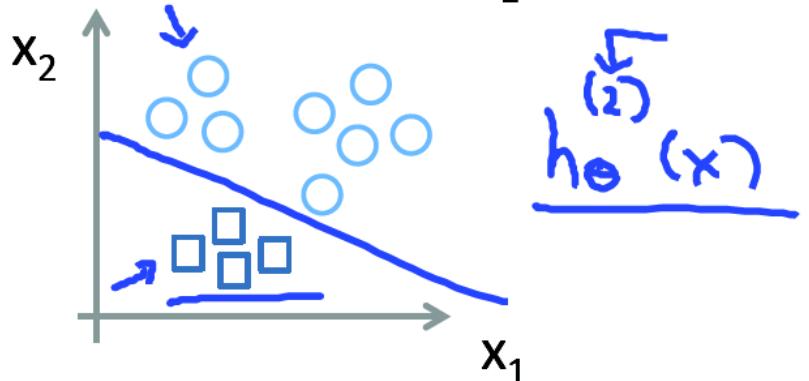
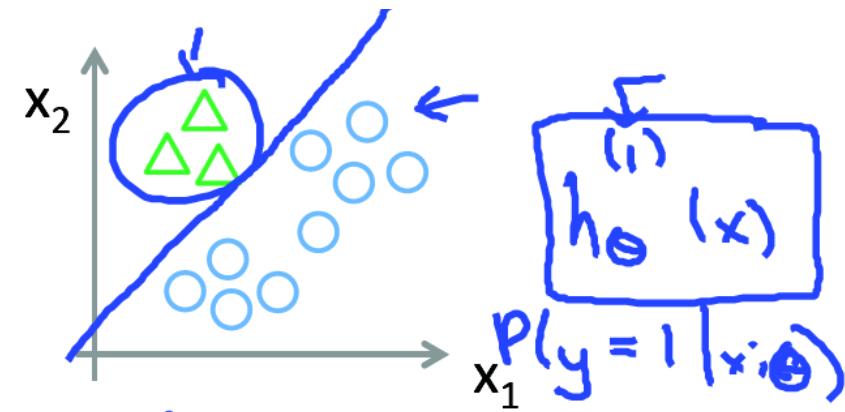
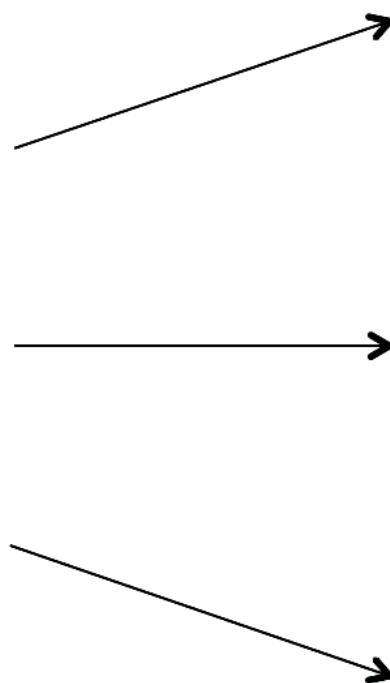


Class 1: ←

Class 2: ←

Class 3: ←

$$h_{\theta}^{(i)}(x) = P(y = i|x; \theta) \quad (i = 1, 2, 3)$$



Multiclass Classification:Softmax Regression

```
fl=fileloader("input")
iris = datasets.load_iris()
X = iris["data"][:, (2, 3)] # petal length, petal width
y = iris["target"]

softmax_reg = LogisticRegression(multi_class="multinomial", solver="lbfgs", C=10)
softmax_reg.fit(X, y)
x0, x1 = np.meshgrid(
    np.linspace(0, 8, 500).reshape(-1, 1),
    np.linspace(0, 3.5, 200).reshape(-1, 1),
)
X_new = np.c_[x0.ravel(), x1.ravel()]
y_proba = softmax_reg.predict_proba(X_new)
y_predict = softmax_reg.predict(X_new)

zz1 = y_proba[:, 1].reshape(x0.shape)
zz = y_predict.reshape(x0.shape)

plt.figure(figsize=(10, 4))
plt.plot(X[y==2, 0], X[y==2, 1], "g^", label="Iris-Virginica")
plt.plot(X[y==1, 0], X[y==1, 1], "bs", label="Iris-Versicolor")
plt.plot(X[y==0, 0], X[y==0, 1], "yo", label="Iris-Setosa")

from matplotlib.colors import ListedColormap
custom_cmap = ListedColormap(['#fafab0','#9898ff','#a0faa0'])
```

