

Support Vector Machine


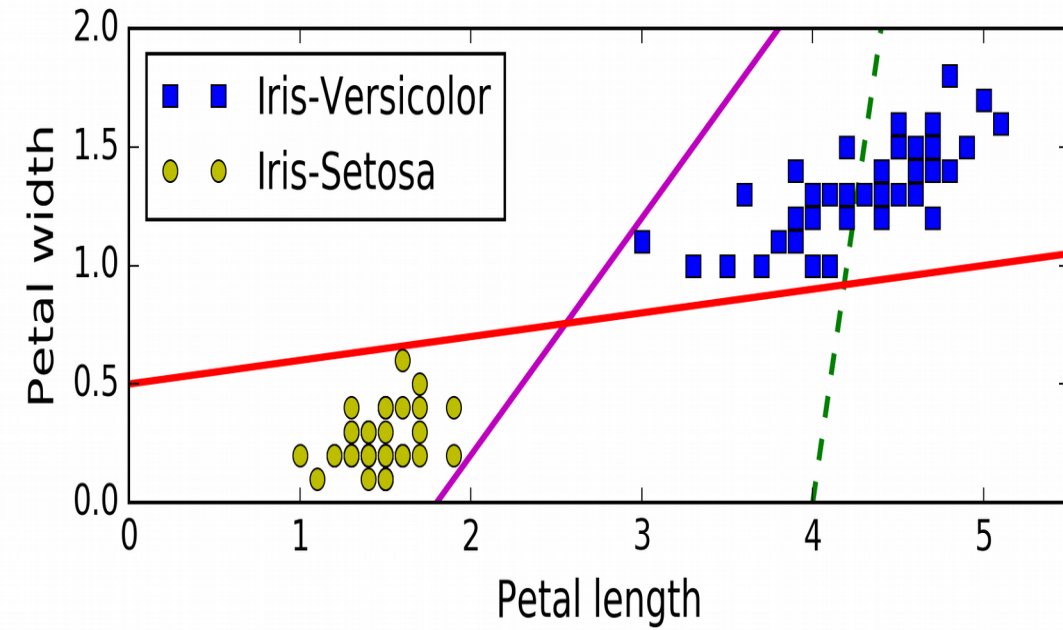
Support Vector Machine

```
fl=fileloader("input")
iris = datasets.load_iris()
X = iris["data"][:, (2, 3)] # petal length, petal width
y = iris["target"]

setosa_or_versicolor = (y == 0) | (y == 1)
X = X[setosa_or_versicolor]
y = y[setosa_or_versicolor]

# SVM Classifier model
svm_clf = SVC(kernel="linear", C=float("inf"))
svm_clf.fit(X, y)
print("svm_clf:",svm_clf)
```

Support Vector Machine



Support Vector Machine


```
fl=fileloader("input")
Xs = np.array([[1, 50], [5, 20], [3, 80], [5, 60]]).astype(np.float64)
ys = np.array([0, 0, 1, 1])
svm_clf = SVC(kernel="linear", C=100)
svm_clf.fit(Xs, ys)
```

```
plt.figure(figsize=(12,3.2))
plt.subplot(121)
plt.plot(Xs[:, 0][ys==1], Xs[:, 1][ys==1], "bo")
plt.plot(Xs[:, 0][ys==0], Xs[:, 1][ys==0], "ms")
plot_svc_decision_boundary(svm_clf, 0, 6)
plt.xlabel("$x_0$", fontsize=20)
plt.ylabel("$x_1$", fontsize=20, rotation=0)
plt.title("Unscaled", fontsize=16)
plt.axis([0, 6, 0, 90])
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(Xs)
svm_clf.fit(X_scaled, ys)
```


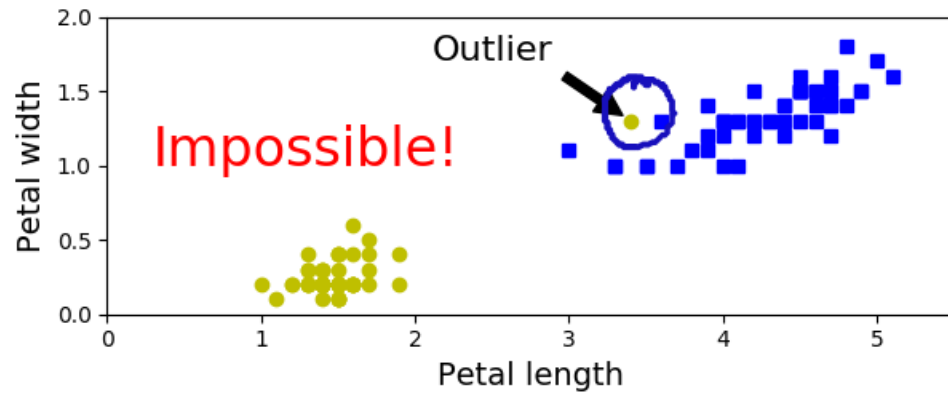
```
plt.subplot(122)
plt.plot(X_scaled[:, 0][ys==1], X_scaled[:, 1][ys==1], "bo")
plt.plot(X_scaled[:, 0][ys==0], X_scaled[:, 1][ys==0], "ms")
plot_svc_decision_boundary(svm_clf, -2, 2)
plt.xlabel("$x_0$", fontsize=20)
plt.title("Scaled", fontsize=16)
plt.axis([-2, 2, -2, 2])
```

```
fl.save_fig("sensitivity_to_feature_scales_plot")
plt.show()
```



- features in different scales
- SVM is quite sensitive to feature scales.
- The decision boundary looks much better with scaling.

Support Vector Machine: Hard Margin



- Hard margin classification requires data to be linearly separable.
- It is also quite sensitive to outliers.

Support Vector Machine: Soft Margin



```
scaler = StandardScaler()
svm_clf1 = LinearSVC(C=100, loss="hinge")
svm_clf2 = LinearSVC(C=1, loss="hinge")
```

```
scaled_svm_clf1 = Pipeline((
    ("scaler", scaler),
    ("linear_svc", svm_clf1),
))
scaled_svm_clf2 = Pipeline((
    ("scaler", scaler),
    ("linear_svc", svm_clf2),
))
```

```
scaled_svm_clf1.fit(X, y)
scaled_svm_clf2.fit(X, y)
```

- The objective is to find a good balance between keeping the street as wide as possible and limiting margin violations.


- Recall that C is like $1/\lambda$. So lower values have the effect of regularization



Support Vector Machine: Soft Margin

- Alternatively, you could use the SVC class, using `SVC(kernel="linear", C=1)`, but it is much slower, especially with large training sets, so it is not recommended.
- Another option is to use the SGDClassifier class, with `SGDClassifier(loss="hinge", alpha=1/(m*C))`. This applies regular Stochastic Gradient Descent to train a linear SVM classifier.
 - It does not converge as fast as the LinearSVC class, but it can be useful to handle huge datasets that do not fit in memory (out-of-core training), or to handle online classification tasks.

Support Vector Machine: Non Linear



- Linear SVM classifiers are good but many dataset are not even close to being linearly separable.
- Often adding more features make it so.

Support Vector Machine: Non Linear


```
X, y = make_moons(n_samples=100, noise=0.15, random_state=42)

polynomial_svm_clf = Pipeline((
    ("poly_features", PolynomialFeatures(degree=3)),
    ("scaler", StandardScaler()),
    ("svm_clf", LinearSVC(C=10, loss="hinge"))
))

polynomial_svm_clf.fit(X, y)

def plot_predictions(clf, axes):
    x0s = np.linspace(axes[0], axes[1], 100)
    x1s = np.linspace(axes[2], axes[3], 100)
    x0, x1 = np.meshgrid(x0s, x1s)
    X = np.c_[x0.ravel(), x1.ravel()]
    y_pred = clf.predict(X).reshape(x0.shape)
    y_decision = clf.decision_function(X).reshape(x0.shape)
    plt.contourf(x0, x1, y_pred, cmap=plt.cm.brg, alpha=0.2)
    plt.contourf(x0, x1, y_decision, cmap=plt.cm.brg, alpha=0.1)



plot_predictions(polynomial_svm_clf, [-1.5, 2.5, -1, 1.5])
plot_dataset(X, y, [-1.5, 2.5, -1, 1.5])
```



- Adding polynomial works great with most ML algorithms.
- Low degrees cannot deal with complex data set, higher degree creates huge number of features making the model too slow.

Support Vector Machine: Non Linear: Kernel

```
poly_kernel_svm_clf = Pipeline((  
    ("scaler", StandardScaler()),  
    ("svm_clf", SVC(kernel="poly", degree=3, coef0=1, C=5))  
))  
poly100_kernel_svm_clf = Pipeline((  
    ("scaler", StandardScaler()),  
    ("svm_clf", SVC(kernel="poly", degree=10, coef0=100, C=5))  
))  
  
poly_kernel_svm_clf.fit(X, y)  
poly100_kernel_svm_clf.fit(X, y)
```



- Increase or decrease depending on whether the model is underfitting or overfitting.
- Grid search to find the right balance.

• SVMs however can use a great mathematical trick called Kernel.

• It has the effect of high degree polynomial without the combinatorial explosion.

• γ controls how much the model is effected by low and high degree polynomials.

Support Vector Machine:Non Linear:Kernel

```



fl=fileloader("input")
X1D = np.linspace(-4, 4, 9).reshape(-1, 1)
X2D = np.c_[X1D, X1D**2]
y = np.array([0, 0, 1, 1, 1, 1, 1, 0, 0])
def gaussian_rbf(x, landmark, gamma):
    return np.exp(-gamma * np.linalg.norm(x - landmark, axis=1)**2)

gamma = 0.3

x1s = np.linspace(-4.5, 4.5, 200).reshape(-1, 1)
x2s = gaussian_rbf(x1s, -2, gamma)
x3s = gaussian_rbf(x1s, 1, gamma)

XK = np.c_[gaussian_rbf(X1D, -2, gamma), gaussian_rbf(X1D, 1, gamma)]
yk = np.array([0, 0, 1, 1, 1, 1, 1, 0, 0])
    
```

- $\text{sim}(x, l) = \exp(-\lambda \|x - l\|^2)$
- if $\lambda = 0.3$
 $x_2 = e^{-(-0.3 \times 1^2)} = 0.74$
 $x_3 = e^{-(-0.3 \times 2^2)} = 0.3$



Support Vector Machine:Non Linear:Kernel

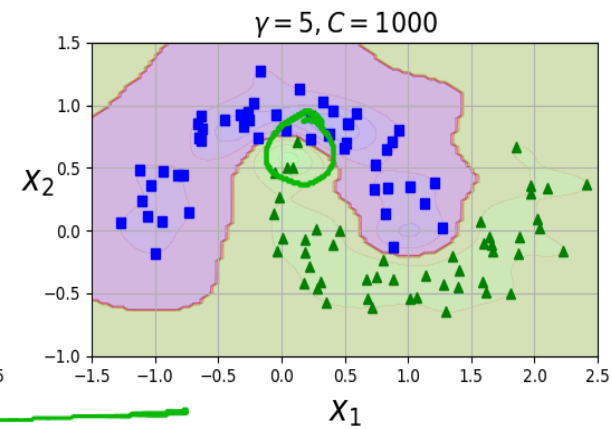
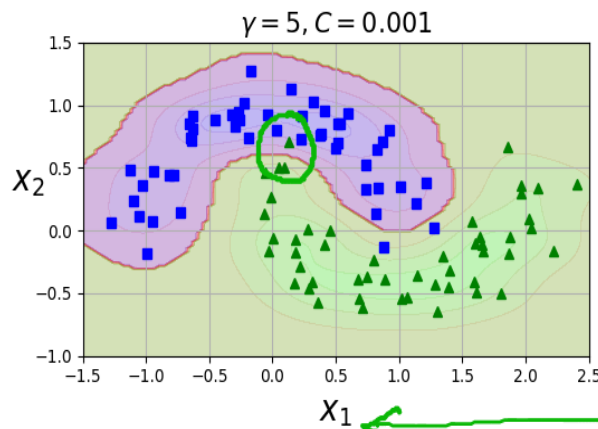


```
gamma1, gamma2 = 0.1, 5  
C1, C2 = 0.001, 1000  
hyperparams = (gamma1, C1), (gamma1, C2), (gamma2, C1), (gamma2, C2)
```

```
svm_clfs = []  
for gamma, C in hyperparams:  
    rbf_kernel_svm_clf = Pipeline((  
        ("scaler", StandardScaler()),  
        ("svm_clf", SVC(kernel="rbf", gamma=gamma, C=C))  
    ))  
    rbf_kernel_svm_clf.fit(X, y)  
    svm_clfs.append(rbf_kernel_svm_clf)
```

```
plt.figure(figsize=(11, 7))
```

```
for i, svm_clf in enumerate(svm_clfs):  
    plt.subplot(221 + i)  
    plot_predictions(svm_clf, [-1.5, 2.5, -1, 1.5])  
    plot_dataset(X, y, [-1.5, 2.5, -1, 1.5])  
    gamma, C = hyperparams[i]  
    plt.title(r"$\gamma = {}, C = {}".format(gamma, C), font
```

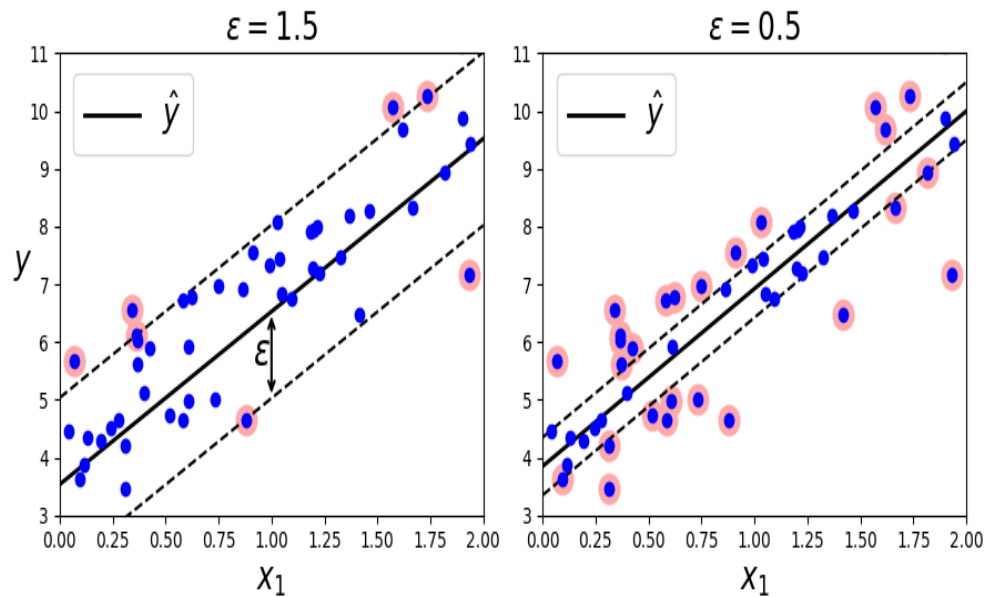
- higher γ would make the bell curve steeper.
- lower C values would have wider streets at the cost of margin violations



Support Vector Machine:Regression

```
rnd.seed(42)
m = 50
X = 2 * rnd.rand(m, 1)
y = (4 + 3 * X + rnd.randn(m, 1)).ravel()
```

```
svm_reg1 = LinearSVR(epsilon=1.5)
svm_reg2 = LinearSVR(epsilon=0.5)
svm_reg1.fit(X, y)
svm_reg2.fit(X, y)
```



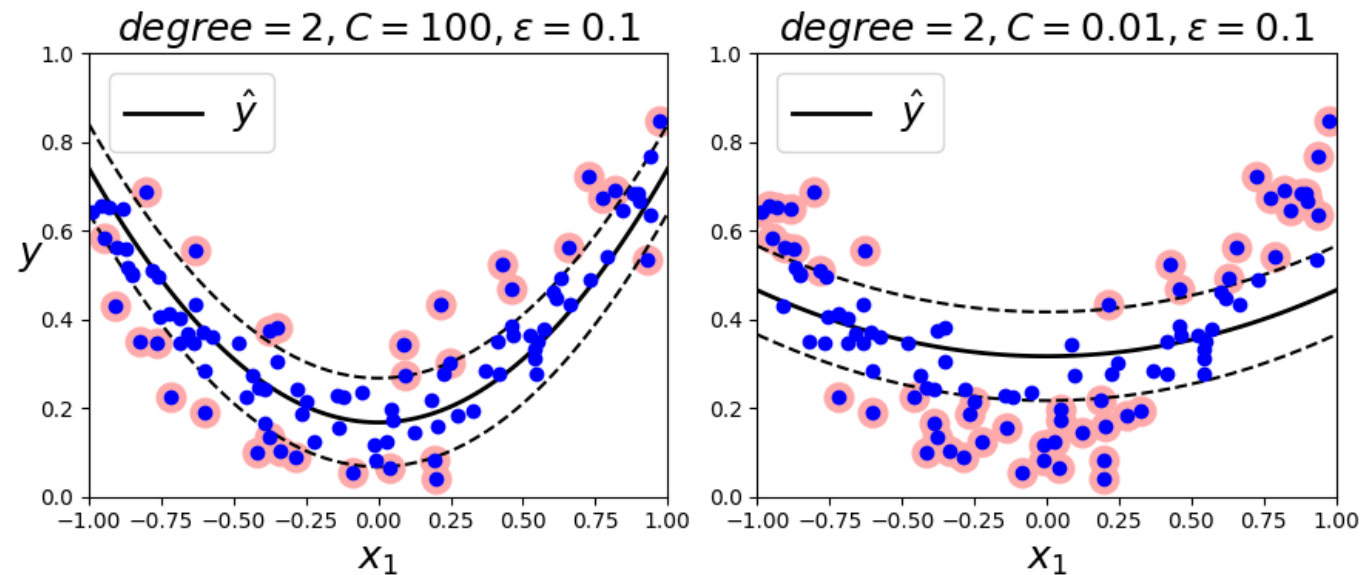
- Objective is reversed. Instead of fitting the widest possible street, while limiting margin violations, SVM Regression tries to fit as many instances on the street with as few margin violations.
- ϵ controls the width of the street.

Support Vector Machine:Regression

```
rnd.seed(42)
m = 100
X = 2 * rnd.rand(m, 1) - 1
y = (0.2 + 0.1 * X + 0.5 * X**2 + rnd.randn(m, 1)/10).ravel()
```

```
svm_poly_reg1 = SVR(kernel="poly", degree=2, C=100, epsilon=0.1)
svm_poly_reg2 = SVR(kernel="poly", degree=2, C=0.01, epsilon=0.1)
svm_poly_reg1.fit(X, y)
svm_poly_reg2.fit(X, y)
```

• SVR kernel.



Support Vector Machine:Regression

```
X, y = make_moons(n_samples=1000, noise=0.4)
plt.plot(X[:, 0][y==0], X[:, 1][y==0], "bs")
plt.plot(X[:, 0][y==1], X[:, 1][y==1], "g^")
plt.show()

tol = 0.1
tols = []
times = []
for i in range(10):
    svm_clf = SVC(kernel="poly", gamma=3, C=10, tol=tol, verbose=1)
    t1 = time.time()
    svm_clf.fit(X, y)
    t2 = time.time()
    times.append(t2-t1)
    tols.append(tol)
    #print(i, tol, t2-t1)
    tol /= 10
plt.semilogx(tols, times)
plt.show()
```

